

Multi-Class Annotation Errors in Causal Inference: Topology-Dependent Bias, Sensitivity, and Correction

Anonymous ACL submission

Abstract

Large language models are increasingly used as annotators in computational social science, yet annotation errors propagate into downstream causal estimates in ways that depend on both the confusion matrix structure and the causal DAG topology. We derive bias expressions for seven canonical topologies under K -class non-differential misclassification in linear structural equation models and show that $K \geq 3$ introduces a structural possibility of sign reversal absent in the binary case. We prove sign preservation for confounding and collider, and verify it numerically for mediation, M-bias, and front-door; exposure and IV topologies can reverse the treatment effect sign. Under realistic LLM accuracy, sign reversal of non-dominant coefficients occurs under exposure; under empirical LLM matrices (diagonal 0.45–0.81), sign reversal is confined to exposure; IV is fragile under uniform Dirichlet priors but exhibits zero reversals under empirical LLM matrices. We introduce Annotation Sensitivity Values (ASVs), with proved ($K=2$) and numerically verified ($K \geq 3$) monotonicity, yielding a vulnerability ranking from IV (maximally fragile) to M-bias (immune), and show that matrix-based corrections are provably ineffective under linear regression for four topologies due to Frisch–Waugh–Lovell invariance. Topology-aware analysis is necessary for valid causal inference with LLM-generated annotations.

1 Introduction

Large language models are rapidly replacing human annotators in computational social science, with researchers using LLM-generated labels for tasks ranging from sentiment classification to political stance detection (Gillardi et al., 2023; Ziems et al., 2024; Egami

et al., 2023; Baumann et al., 2025). Crucially, most of these tasks are inherently multi-class: stance detection uses 3–4 categories, sentiment analysis 5, and disaster response taxonomies 10 or more. This shift promises dramatic cost reduction, but introduces a critical validity gap: annotation errors (equivalently, misclassification) from imperfect classifiers propagate into downstream causal estimates in ways that depend on both the error structure and the causal design. Baumann et al. (2025) demonstrate that across 37 annotation tasks, 31% of hypotheses yield incorrect conclusions with state-of-the-art LLMs, rising to 50% for smaller models; these errors are highly correlated across models (60% agreement when both err) (Kim et al., 2025), undermining ensemble-based mitigation. Analyzing the *structure* of annotation error, rather than simply reducing its magnitude, is essential for valid downstream inference.

The transition at $K \geq 3$ is not gradual but structural. For $K=2$, the stochastic constraint forces any confusion matrix with positive diagonal to preserve coefficient signs—an algebraic property of the 2×2 matrix dimension (Brenner, 1993). This guarantee *vanishes* at $K \geq 3$: the higher-dimensional confusion matrix admits negative effective weights on categories with opposite-sign coefficients, enabling sign reversal even under high accuracy (Veierød and Laake, 2001). Under a uniform Dirichlet prior, 83.4% of $K=3$ exposure confusion matrices produce sign reversal of at least one secondary coefficient; under realistic accuracy (diagonal ≥ 0.7), 2.1% still do. The primary treatment coefficient is never reversed under the tested β configurations where a single category dominates ($\beta_1 \gg \beta_2$)—an empirical regularity we discuss in Section 3.

Standard classification metrics do not pre-

dict downstream causal risk: within a narrow F_1 range of 0.614–0.690 across three frontier LLMs, relative bias varies by an order of magnitude across topologies (from 4.1% to 42.1%). F_1 shows no significant rank correlation with causal bias across 56 HumAID scenarios (Spearman $\rho = -0.10$, $p = 0.45$), confirming that *topology*, not annotator quality, determines downstream risk. This fundamental asymmetry has received limited systematic treatment. Dosemeci et al. (1990) and Veierød and Laake (2001) noted the possibility of polytomous sign reversal; Brenner (1993) identified amplification for polytomous confounders; Lee and Wood-Doughty (2024) extended classifier error analysis to categorical variables and LLM classifiers, but covered only one DAG topology. We present the first systematic analysis that jointly addresses multi-class ($K \geq 3$) misclassification across multiple DAG topologies under structural equation models, extending prior binary or single-topology analyses (Brenner, 1993; TeBlunthuis et al., 2024). Our analytical predictions match Monte Carlo estimates within 1.6% MAPE across 161 scenarios.

We address these gaps with three contributions:

1. **Multi-class bias taxonomy.** We derive bias expressions for seven canonical DAG topologies¹ (Figure 1) under K -class non-differential misclassification, establishing a sign-preservation/reversal dichotomy: five topologies preserve the treatment effect sign, while exposure and IV can reverse it (binarization amplifies this gap $>150\times$, from $\bar{\Delta}=0.006$ for M-bias to 0.926 for IV; Appendix B.1).
2. **Annotation Sensitivity Values.** Building on the bias taxonomy, we define three ASV diagnostics—magnitude, significance, and sign-flip—with proved envelope monotonicity for $K=2$ (Theorem 1) and $K=3$ symmetric confusion matrices (Theorem 2), with ε -approximate monotonicity empirically validated for general $K \geq 3$ (Conjecture 1,

$\varepsilon < 0.02$; adversarial structures in Appendix A.11). These reveal a vulnerability spectrum from IV (maximally fragile) to M-bias (immune, $ASV = \infty$), enabling topology-specific error tolerance assessment.

3. **Topology-aware correction routing.** We establish that FWL invariance renders matrix-based corrections provably ineffective for four covariate topologies, evaluate five methods across 210 configurations—PPI++ harms M-bias (−88%), MC-SIMEX degrades front-door (−62.4%)—and provide a topology-aware decision procedure (Algorithm 1, Appendix).

2 Related Work

Measurement error in causal inference. Foundations appear in Carroll et al. (2006); Gustafson (2004); Buonaccorsi (2010); Bound et al. (2001), with DAG-based formalization by Hernán and Cole (2009) and Kuroki and Pearl (2014). For binary variables, non-differential misclassification biases toward the null (Bross, 1954; Brenner, 1993; Ogburn and VanderWeele, 2012), but this fails for polytomous exposures (Dosemeci et al., 1990; Jurek et al., 2005; Veierød and Laake, 2001). Extensions address differential measurement error (Edwards et al., 2021; Imai and Yamamoto, 2010), endogenous misreporting (Nguimkeu et al., 2019), mismeasured covariates under IPW (Shu and Yi, 2019b,a), mediation with mismeasured mediators (Valeri and VanderWeele, 2013), nonclassical error (Schennach, 2020), and multiply robust estimation (Wei et al., 2025). All analyze a single DAG topology; ours characterizes how the *same* confusion matrix produces qualitatively different bias across seven topologies.

Text classifiers and LLM annotation. Grimmer et al. (2022) and Feder et al. (2022) provide foundations for text-as-data and causal inference in NLP. Wood-Doughty et al. (2018) first applied DAG reasoning to classifier error; TeBlunthuis et al. (2024) gave the most complete binary treatment ($K=2$, four DAGs); Lee and Wood-Doughty (2024) extended to categorical LLM classifiers but covered only confounding. On annotation

¹These span the four structural roles of the annotated variable: treatment (exposure), outcome-side covariate (confounding, mediation, M-bias, front-door), selection variable (collider), and instrument (IV).




Figure 1: Seven canonical DAG topologies with starred-node annotation (Hernán and Cole, 2009). A (latent) $\rightarrow A^*$ (observed) via a dashed measurement edge; the causal role of A varies: treatment (exposure), covariate (confounding, mediation, M-bias, front-door), selection (collider), instrument (IV).

quality, Gilardi et al. (2023) and Ziems et al. (2024) document rapid LLM annotator adoption (surveyed by Tan et al. 2024), Baumann et al. (2025) found 31% incorrect conclusions, and Kim et al. (2025) showed 60% error correlation across models (Pangakis et al., 2023; Cheng et al., 2024; Camuffo et al., 2026; Messing, 2026). Audinet de Pieuchon et al. (2025) benchmark debiasing; Keith et al. (2020) survey the landscape. These works either restrict to binary labels, cover a single topology, or diagnose quality without connecting it to causal validity.

Sensitivity analysis and correction methods. Our ASV framework draws on sensitivity analysis without assumptions (Ding and VanderWeele, 2016), E-values (VanderWeele and Ding, 2017; Mathur et al., 2018), robustness values (Cinelli and Hazlett, 2020), and mediation sensitivity (Imai et al., 2010; Pryzant et al., 2021), but parametrizes annotation error severity rather than omitted variable strength. Lash et al. (2009) and Blackwell et al. (2017) provide complementary frameworks; Mahajan (2006) establishes identification under misclassified regressors. For correction, PPI++ (Angelopoulos et al., 2023) and DSL (Egami et al., 2023) are topology-agnostic, with extensions by Zrnić and Candès (2024); Demirel et al. (2024); Cadei et al. (2025); Gligorić et al. (2025). We show this agnosticism is costly: correction effectiveness is itself topology-dependent, with matrix-based methods provably ineffective for four topologies.

3 Method

3.1 Setup and Notation

Let A denote a true K -class categorical variable (e.g., sentiment with $K=3$ levels) and A^* its observed, potentially misclassified version produced by an LLM annotator. The misclassification process is characterized by a $K \times K$ confusion matrix C with entries $C_{jk} = P(A^*=j \mid A=k)$, where each column sums to one (column-stochastic: each column gives the distribution of predicted labels for a given true label). Following Hernán and Cole (2009), we represent this measurement process using *starred-node* DAGs: the true variable A is latent, and the observed A^* is connected to A via a dashed measurement edge.

Notation conventions. Our analysis characterizes probability limits ($\text{plim}, N \rightarrow \infty$) rather than finite-sample variance. The FWL theorem (Frisch and Waugh, 1933) guarantees that each regression coefficient equals the coefficient from a simple regression after residualizing all other variables—this is why matrix-based corrections fail for covariate topologies (Section 4.5). Table 27 (Appendix C) collects all notation.

Definition 1 (Non-differential misclassification). *Misclassification of A is non-differential with respect to outcome Y and covariates $\mathbf{Z}$ if $P(A^*=j \mid A=k, Y, \mathbf{Z}) = P(A^*=j \mid A=k) = C_{jk}$ for all j, k .*

In batch annotation pipelines, the LLM processes each text independently without access to Y , making non-differentiality an architectural invariant (Ziems et al., 2024); Y-stratified permutation tests confirm this (1/12 rejection at $\alpha=0.05$; Appendix B.26). When

annotation error is driven by label-intrinsic ambiguity (e.g., vague category boundaries) rather than content features correlated with Y , non-differentiality approximately holds; when error depends on features that also predict Y (e.g., stance toward a policy affecting the outcome), the assumption may be violated and we recommend the Y -stratified test (Appendix B.26). Under stress testing, $\varepsilon=0.01$ structured departure produces sign reversal in 14/16 $K \geq 3$ exposure configurations (Appendix B.6).

Throughout, τ denotes the causal estimand of interest, whose definition varies by topology: the confounder’s coefficient, the treatment effect, the indirect effect, or the Wald ratio (Appendix A.3ff.).

We study seven canonical DAG topologies (Figure 1) (Pearl, 2009; Hernán and Robins, 2020) in which A occupies structurally distinct causal roles: confounder, exposure, mediator, collider, M-bias variable, front-door mediator, and instrument. Throughout, we derive probability limits under *linear* structural equation models (Carroll et al., 2006), the standard framework for exact characterization of topology-dependent bias; robustness under logistic regression (Stefanski and Carroll, 1985) is verified in Appendix B.7 and B.14.

3.2 Multi-Class Bias Expressions

For each topology, we derive the plim^2 of the naive estimator that substitutes A^* for A . Let τ denote the target causal parameter—its precise definition is topology-specific (e.g., the coefficient on X for confounding, the category-specific β_j for exposure, or the direct effect τ_{direct} for mediation)—and $\hat{\tau}(A^*)$ the estimator using misclassified annotations.

Confounding ($A = \text{confounder}$). When A confounds the $X \rightarrow Y$ relationship, the OLS estimator controlling for A^* has probability limit

$$\text{plim } \hat{\tau}(A^*) = (\mathbb{E}[Z'Z])^{-1} \mathbb{E}[Z'X] \beta, \quad (1)$$

where $Z = [X, \text{dummies}(A^*)]$, β contains the true coefficients, and the target parameter τ

²The probability limit (plim) is the value to which an estimator converges as sample size $N \rightarrow \infty$; it characterizes systematic bias separately from sampling variability.

is the coefficient on X . Misclassification produces residual confounding that monotonically approaches the omitted variable bias, but the sign of $\hat{\tau}$ is preserved: $\tau_{\text{sign_flip}} = 0$ structurally.³

Exposure ($A = \text{treatment variable}$). When A is the exposure of interest, the biased coefficient for category j becomes a weighted combination of true coefficients:

$$\text{plim } \hat{\beta}_j(A^*) = \sum_{k=1}^K w_{jk} \beta_k, \quad (2)$$

where w_{jk} depends on the confusion matrix C and the population marginal $\pi_k = P(A=k)$. Unlike the binary case, where non-differential misclassification guarantees attenuation toward the null (Brenner, 1993), the $K \geq 3$ case permits w_{jk} to assign negative weight to categories with opposite-sign coefficients, producing *sign reversal*. A near-permutation confusion matrix effectively swaps category labels, making sign reversal the dominant outcome (83.4% under uniform Dirichlet prior). This failure mode is structurally impossible at $K=2$ because a 2×2 stochastic matrix with positive diagonal cannot reverse the sign of a single coefficient.

Numerical Verification 2: Primary coefficient preservation. Across 23 empirical LLM configurations, the primary coefficient (the category with the largest true effect $|\beta_k|$) never reverses sign, while 61% reverse a secondary (non-dominant) coefficient—confirmed under 100,000 Dirichlet matrices with diagonal ≥ 0.7 across six beta-ratio conditions (1.0 to 50.0) with 0% primary sign-flip (Appendix A.9).

Mediation ($A = \text{mediator}$). When A mediates the $X \rightarrow Y$ relationship (VanderWeele, 2015), the direct effect estimator controlling for A^* has probability limit

$$\text{plim } \hat{\tau}_{\text{direct}}(A^*) = \tau_{\text{direct}} + (1 - \lambda(C)) \tau_{\text{indirect}}, \quad (3)$$

where $\lambda(C) \in [0, 1]$ under proportional effects (Appendix A, Proposition 4), with $\lambda(I) = 1$

³Consistent with Brenner (1993): misclassification degrades confounding control, producing amplification ($|\hat{\tau}| > |\tau|$, universal in this topology) but never sign reversal ($\text{sign}(\hat{\tau}) \neq \text{sign}(\tau)$).




Figure 2: Bias distributions under 100,000 random $K=3$ confusion matrices (uniform Dirichlet; realistic-accuracy results in Section 4.2). Red = sign reversal. Exposure: 83.4% sign reversal; IV: 52.6%.

under perfect classification. Misclassification ($\lambda < 1$) causes the residual indirect effect to be absorbed into the direct path estimate, producing upward bias in $|\hat{\tau}_{\text{direct}}|$ but preserving the sign ($\tau_{\text{sign_flip}} = 0$).

Collider, M-bias, and front-door. Three topologies exhibit *protective* or *immune* behavior. **Collider:** conditioning on A^* instead of A attenuates the spurious path, so annotation error reduces collider bias ($\tau_{\text{sign_flip}} = 0$). **M-bias:** in the M-shaped DAG (Greenland, 2003), misclassifying A weakens the M-bias path; this topology is immune ($\text{ASV} = \infty$). **Front-door:** the front-door estimator (Belle-mare et al., 2024) is attenuated in both $X \rightarrow A$ and $A \rightarrow Y$ stages, producing moderate bias with $\tau_{\text{sign_flip}} = 0$.

Instrumental variable ($A = \text{instrument}$). The Wald estimator (Angrist and Imbens, 1995) $\hat{\tau}_{\text{IV}} = \text{cov}(A^*, Y) / \text{cov}(A^*, X)$ has C in the denominator. Misclassification can drive first-stage coefficients toward zero or reverse their sign; under random Dirichlet matrices, 52.6% produce sign reversal, making IV the most fragile topology (Appendix B.4).

Structural sign-flip guarantee.

Proposition 1 (Sign preservation—analytic results). *Under non-differential misclassification with $C_{kk} > 0$ for all k : (a) Confounding preserves sign when $\text{sign}(\tau) = \text{sign}(\tau + B_{\text{OV}})$, where $B_{\text{OV}} = \text{Cov}(g(A), f(A)) / \text{Var}(T)$ is the omitted-variable bias (Cauchy-Schwarz; Appendix A.3). (b) Collider preserves sign when $|\tau| > \sqrt{\text{Var}(g) \text{Var}(h) / \mathbb{E}[\text{Var}(T|A)]}$ (Berkson attenuation; Appendix A.7); empirically verified across 100,000 matrices with zero violations. (c) Exposure can produce sign reversal for $K \geq 3$, even with $C_{kk} > 0.5$ (counterexample*

in Appendix A.9).

Sign preservation for non-exposure topologies holds under topology-specific conditions (Table 2, Appendix A.10), verified across 100,000 random matrices with zero violations; analytic proofs for general K remain open.

3.3 Annotation Sensitivity Values

We adapt the sensitivity analysis paradigm—E-values (VanderWeele and Ding, 2017), partial- R^2 bounds (Cinelli and Hazlett, 2020), coefficient stability (Oster, 2019)—to $K \times K$ annotation error, defining Annotation Sensitivity Values (ASVs): we prove ε -approximate monotonicity along $C(\delta)$ (Theorem 1), reducing K^2 dimensions to a scalar; define three ASV diagnostics for topology-dependent failure modes; and combine ASVs with FWL invariance to route corrections.

Parameterization. We set $C(\delta) = (1-\delta)I + \delta C_0$, where C_0 is a reference confusion matrix and $\delta \in [0, 1]$ controls severity ($\delta=0$: perfect; $\delta=1$: matches C_0). This linear path is one of many in the column-stochastic polytope from I to C_0 ; ASV thresholds are path-dependent, but our bootstrap procedure (Section 4.3) captures estimation uncertainty along the empirical path. We recommend bootstrapping C_0 and reporting ASV confidence intervals (Appendix B.8).

Three ASV types. We define three complementary measures: (1) **Magnitude ASV** $\delta_{\text{mag}}^* = \min\{\delta : |\text{plim } \hat{\tau}(C(\delta)) - \tau| > \epsilon\}$ —analogous to E-values; (2) **Significance ASV** $\delta_{\text{sig}}^* = \min\{\delta : |\text{plim } \hat{\tau}(C(\delta))| < z_{\alpha/2} \cdot \text{SE}(N)\}$ —relevant for attenuating topologies; (3) **Sign-flip ASV** $\delta_{\text{sf}}^* = \min\{\delta : \text{sign}(\hat{\tau}) \neq$

sign(τ)—exposure and IV only. The significance ASV depends on N through $SE(N)$, exhibiting a sharp phase transition at a critical N^* below which all rejections are blocked.

Theorem 1 (Envelope monotonicity— $K=2$). *For $K=2$ confounding and mediation topologies with proportional effects ($f = \lambda g$ for some scalar λ) under $C(\delta) = (1-\delta)I + \delta C_0$ with $(C_0)_{jj} \geq \frac{1}{2}$, $|B(C(\delta))|$ is monotonically non-decreasing in $\delta \in [0, 1]$.*

The proof (Appendix A.4, A.5) exploits perspective function convexity; $(C_0)_{jj} \geq \frac{1}{2}$ excludes degenerate matrices. Proportional effects ($f = \lambda g$) is plausible when overall annotation competence drives both paths (e.g., a single quality dimension); it may fail under multi-dimensional asymmetric confusion where different error modes affect f and g differentially. All 23 empirical configurations preserve sign under this assumption, with 97.3% preserving sign under non-proportional effects (Appendix B.13; Tables 14, 18).

Theorem 2 ($K=3$ symmetric monotonicity). *For $K=3$ with symmetric confusion matrix $C_{jk} = c$ for $j \neq k$ and $C_{kk} = 1-2c$, $c \in (0, 1/3)$, and proportional effects ($f = \lambda g$), envelope monotonicity holds exactly: $|B(C(\delta))|$ is strictly monotone in δ for confounding, mediation, and collider topologies whenever $\text{Var}(g) > 0$.*

Theorem 2 serves as a theoretical anchor: exact monotonicity is achievable at $K=3$ under symmetric C_0 . For general asymmetric matrices (empirical $K=3$ median $d_{\text{rel}}=0.33$; Remark 4), Conjecture 1 provides numerical coverage ($\varepsilon < 0.017$). The proof (Appendix A.6) exploits the one-parameter reduction $s = c\delta$ to show $dV_B/ds < 0$ pointwise; under non-proportional effects, approximate monotonicity holds ($\varepsilon < 0.005$; Appendix A.6, Remark 2).

From proofs to numerical evidence. Theorems 1–2 establish exact monotonicity for $K=2$ and $K=3$ symmetric matrices. The factorization in Appendix A.6 exploits off-diagonal symmetry to reduce to univariate analysis—a strategy that does not extend to general $K \geq 3$ column-stochastic structures. We therefore state a conjecture supported by numerical evidence:

Conjecture 1 (ε -approximate envelope monotonicity). *For general $K \geq 3$ non-IV topologies under $C(\delta) = (1-\delta)I + \delta C_0$, $|B(C(\delta))|$ is ε -approximately monotonic with $\varepsilon < 0.02$ for $|\beta| \leq 1$, scaling with effect magnitude (Table 11, 13; Appendix B.12). The safe/fragile partition—five sign-preserving topologies versus exposure and IV—is robust across all tested parameterizations.*

Conjecture 1 is supported by 100,000 Dirichlet-sampled matrices and 545,945 adversarial curves across four structural families (Appendix A.11), with maximum observed violation $\varepsilon < 0.017$.

Remark 1 (ASV and E-values as complementary tools). *ASV diagnostics and E-values (VanderWeele and Ding, 2017) address orthogonal bias sources—annotation error vs. unmeasured confounding—and are most valuable jointly (Table 22, Appendix). Key distinctions: ASV is topology-dependent (same C has different consequences across topologies); sign-flip ASV targets a failure mode absent from confounding analysis; ASV searches the column-stochastic polytope while E-values optimize over $[1, \infty)$.*

Under alternating-sign effects ($\beta = [0, 1.0, -0.5]$), magnitude ASV ranges from 0.09 (exposure) to ∞ (M-bias), a $4\times$ difference determined by topology (Table 10). Sign-flip ASV < 1 flags sign-reversal-prone pairs; magnitude ASV $< \delta_{\text{observed}}$ indicates bias exceeds tolerance (Algorithm 1); the safe/fragile partition is stable across parameterizations ($\varepsilon < 0.02$; Appendix B.11).

4 Experiments

We validate through Monte Carlo simulation ($K=3$), cross- K robustness checks ($K=5, 7$), and empirical validation with real LLM confusion matrices from three NLP tasks ($K=2, 3, 10$).

4.1 Setup

For each topology, we sample 100,000 confusion matrices from a Dirichlet(1, 1, 1) prior, compute analytical probability limits, and validate against Monte Carlo simulation ($N=10,000$). ASV experiments use $C(\delta) = (1-\delta)I + \delta C_0$.

4.2 Bias Taxonomy

Theory-MC agreement. Analytical probability limits match Monte Carlo estimates with MAPE of 1.6%, averaged across 161 scenarios (23 LLM configurations $\times$ 7 topologies; per-topology range: 0.3%–3.7%).

Exposure DAG: sign reversal under multi-class misclassification. Under realistic accuracy (diagonal ≥ 0.7), the violation rate is 13.7% with 2.1% sign reversal; at diagonal ≥ 0.9 , violation persists at 11.9% (Appendix B.4). Under the full Dirichlet prior, 85% of matrices violate the binary “toward-null” guarantee and 83.4% cause sign reversal; $\min_k C_{kk} < 0.55$ predicts sign reversal with 91% accuracy (Appendix B.25). The $K=2$ control confirms classical attenuation (Brenner, 1993), establishing the multi-class failure as structurally distinct.

Structural guarantee for non-exposure topologies. Five topologies exhibit zero sign-flip rate (Table 1): confounding, mediation, and front-door show 100% amplification (Brenner, 1993; Dosemeci et al., 1990) but never reverse the sign; collider is protective ($|B(C)|$ non-increasing; Conjecture 1); M-bias shows weakest sensitivity (53.2% violation). All patterns persist at $K=5, 7$ (Appendix B.9).

IV: Wald estimator instability. The IV topology is the sole non-exposure exception: 52.6% of random matrices cause sign reversal, stable across $K=3, 5, 7$. Diagonal dominance ($C_{kk} > 0.5$) eliminates sign reversal but preserves amplification (Appendix B.4); however, all 23 empirical LLM configurations show zero IV sign reversal.

The bias distributions are visualized in Figure 2.

4.3 ASV Framework Validation

ASV results. The ASV envelope is ε -approximately monotonic for non-IV topologies ($\varepsilon < 0.02$ for non-exposure under moderate effects, larger for exposure; Conjecture 1). Figure 3 shows the vulnerability ranking: IV is maximally fragile, M-bias is immune (ASV = ∞). Sign-flip ASV: 83.4% of exposure matrices yield $\delta_{sf}^* < 1$; IV ranges 1.1–6.7 (Table 4). Significance ASV exhibits a phase transition at $N^* \approx 6,000$ –10,000; bootstrap CIs ($B=1,000$;

Table 1: Bias under $K=3$ non-differential misclassification (100,000 Dirichlet(1, 1, 1) matrices). Violation: $|\text{bias}| > 0$; Sign Flip: any coefficient reverses sign; Amplif.: $|\hat{\tau}| > |\tau|$. $\tau_{sf}=0?$: sign-flip guarantee holds unconditionally (Yes), conditionally (Cond.; Proposition 1), or not (No). Collider 100% amplification reflects Berkson bias—misclassification is protective.

Topology	Violation	Sign Flip	Amplif.	$\tau_{sf}=0?$
Confounding	100%	0%	100%	Cond.
Exposure	85%	83.4%	14.1%	No
Mediation	100%	0%	100%	Cond.
Front-door	100%	0%	100%	Cond.
Collider	100%	0%	100%	Cond.
M-bias	53.2%	0%	10.4%	Yes
IV	100%	52.6%	100%	No




Figure 3: ASV vulnerability ranking across topologies. Three ASV types (magnitude, significance, sign-flip) reveal a topology-dependent spectrum from IV (maximally fragile) to M-bias (immune, ASV = ∞).

Appendix B.8) confirm stable topology ranking despite wide IV intervals.

4.4 Empirical Validation

We validate using real LLM confusion matrices from HumAID (Alam et al., 2021) ($K=10$), VAST ($K=3$), and CivilComments ($K=2$): four LLMs, 23 configurations, 161 bias scenarios.

Empirical bias ranking. HumAID ($K=10$) matches predictions: IV (12.4–42.1%) and exposure (9.0–17.4%) dominate; remaining topologies $<15\%$. F_1 shows no rank correlation with bias ($\rho = -0.10$, $p = 0.45$; Figure 4), but per-topology patterns diverge: five topologies show $F_1 \uparrow \rightarrow \text{bias} \downarrow$ ($\rho \leq -0.53$), while M-bias reverses ($\rho = 0.87$; Appendix B.23). Of exposure coefficients, 14/23 (61%) exhibit sign reversal—all secondary; zero across 138 non-exposure scenarios.




Figure 4: Topology determines downstream causal bias, not F_1 . HumAID $K=10$ (56 points): within F_1 range 0.614–0.690, bias varies 10 $\times$ (IV 42.1% vs. collider <5.6%).

4.5 Correction Guidance

We evaluate five correction methods (naive plug-in, MLA, DSL (Egami et al., 2023), MC-SIMEX (Küchenhoff et al., 2006), PPI++ (Angelopoulos et al., 2023)) across all topologies (210 configurations; Table 5, Appendix B).

FWL invariance (observation). FWL invariance implies matrix-based methods (naive plug-in, MLA) achieve 0% bias reduction for covariate topologies (Appendix B.2); under logistic regression this breaks but ranking persists ($|\Delta| < 0.02$; Appendix B.14).

Method ranking. PPI++ achieves the highest RMSE reduction (47.0% at $K=3$) but requires gold labels and harms M-bias (−88%); MC-SIMEX is gold-free and K -robust (27–70%), though it harms front-door (−62.4%); rankings are stable across K ($\rho \geq 0.80$; Appendix B.21).

Correction decision procedure. Given topology $\mathcal{T}$ and gold-label availability: (1) **Risk:** M-bias $\rightarrow$ low; confounding/mediation/collider/front-door $\rightarrow$ medium; exposure/IV $\rightarrow$ high. (2) **Method:** with gold, PPI++ except M-bias; without gold, DSL for confounding/mediation, MC-SIMEX for IV/collider. (3) **FWL:** skip plug-in/MLA for covariate topologies (0% reduction). If $\mathcal{T}$ uncertain, compute ASV under candidate topologies and report worst-case sensitivity (Algorithm 1, Appendix B); thresholds may require recalibration for

different K or DGP families.

Case studies and sensitivity. ClinicalBERT and HumAID ($K=10$) confirm all patterns—FWL invariance, PPI++ harm for M-bias/IV, partition recoverability—under linear and logistic models ($|\Delta| < 0.02$; Appendix B.3, B.21). An end-to-end demo ($K=3$, $N=10,000$, 100 replicates; Table 19) confirms: confounding achieves 0% improvement (FWL), PPI++ reduces exposure bias from 42.9% to 0.2%, while topology-agnostic IV correction amplifies bias (−139%). A semi-real case study using 10 production VAST confusion matrices ($d_{\text{rel}} = 0.30\text{--}0.38$) confirms both patterns: PPI++ achieves 93.6% median exposure bias reduction while confounding correction yields 0% improvement (Appendix B.20).

5 Conclusion

We characterized how multi-class annotation errors interact with causal DAG topology: at $K \geq 3$, exposure and IV can reverse treatment effect signs while five topologies preserve them. ASVs yield a topology-dependent vulnerability ranking; matrix-based corrections are provably ineffective for four covariate topologies (FWL invariance). Practitioners should identify their topology and compute ASV before selecting corrections.

Limitations

We assume non-differential misclassification (Appendix B.6). Closed-form results (Theorems 1–2) assume linear SEMs with proportional effects ($f=\lambda g$); FWL invariance extends to logistic models for covariate topologies (Appendix B.2), but for exposure, logistic specification can increase bias (Appendix B.14), and practitioners should verify linearity assumptions. Sign preservation holds at 97.3% under non-proportional effects (Table 18). Proving general $K \geq 3$ envelope monotonicity beyond symmetric matrices remains open (Schennach, 2020).

References

Firoj Alam, Ferda Ofli, and Muhammad Imran. 2021. HumAID: Human-annotated disaster incidents data from Twitter. In *Proceedings of the*

661	<i>International AAAI Conference on Web and Social Media</i> , volume 15, pages 933–942.	714
662		715
663	Anastasios N. Angelopoulos, Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnić.	716
664	2023. Prediction-powered inference . <i>Science</i> ,	717
665	382(6671):669–674.	718
666		719
667	Joshua D. Angrist and Guido W. Imbens. 1995.	720
668	Two-stage least squares estimation of average	721
669	causal effects in models with variable treatment	722
670	intensity . <i>Journal of the American Statistical</i>	723
671	<i>Association</i> , 90(430):431–442.	724
672	Nicolas Audinet de Pieuchon, Adel Daoud, Con-	
673	nor Thomas Jerzak, Moa Johansson, and	
674	Richard Johansson. 2025. Benchmarking debi-	
675	asing methods for LLM-based parameter esti-	
676	mates. In <i>Proceedings of the 2025 Conference</i>	
677	<i>on Empirical Methods in Natural Language Pro-</i>	
678	<i>cessing</i> , pages 19757–19772.	
679	Maximilian Baumann, Paul Röttger, Aleksandra	
680	Urman, Sina Wendsjö, Flor Miriam Plaza-del	
681	Arco, Nora Gruber, and Dirk Hovy. 2025. Large	
682	language model hacking: Quantifying the hid-	
683	den risks of using LLMs for text annotation.	
684	<i>arXiv preprint arXiv:2509.08825</i> . Preprint.	
685	Marc F. Bellemare, Jeffrey R. Bloem, and Noah	
686	Wexler. 2024. The paper of how: Estimat-	
687	ing treatment effects using the front-door crite-	
688	rion . <i>Oxford Bulletin of Economics and Statis-</i>	
689	<i>tics</i> , 86(4):951–993.	
690	Matthew Blackwell, James Honaker, and Gary	
691	King. 2017. A unified approach to measure-	
692	ment error and missing data: Overview and	
693	applications . <i>Sociological Methods & Research</i> ,	
694	46(3):303–341.	
695	John Bound, Charles Brown, and Nancy Math-	
696	iowetz. 2001. Measurement error in survey data.	
697	In James J. Heckman and Edward E. Leamer,	
698	editors, <i>Handbook of Econometrics</i> , volume 5,	
699	chapter 59, pages 3705–3843. Elsevier.	
700	Hermann Brenner. 1993. Bias due to non-	
701	differential misclassification of polytomous con-	
702	founders . <i>Journal of Clinical Epidemiology</i> ,	
703	46(1):57–63.	
704	Irwin Bross. 1954. Misclassification in 2 x 2 tables .	
705	<i>Biometrics</i> , 10(4):478–486.	
706	John P. Buonaccorsi. 2010. <i>Measurement Error:</i>	
707	<i>Models, Methods, and Applications</i> . Chapman	
708	& Hall/CRC.	
709	Riccardo Cadei, Ilker Demirel, Piersilvio De Bar-	
710	tolomeis, Lukas Lindorfer, Sylvia Cremer,	
711	Cordelia Schmid, and Francesco Locatello. 2025.	
712	Prediction-powered causal inferences. <i>arXiv</i>	
713	<i>preprint arXiv:2502.06343</i> .	
	Arnaldo Camuffo, Alfonso Gambardella, Saeid	714
	Kazemi, Jakub Malachowski, and Abhinav	715
	Pandey. 2026. Variance-aware LLM annota-	716
	tion for strategy research: Sources, diagnostics,	717
	and a protocol for reliable measurement . <i>arXiv</i>	718
	<i>preprint arXiv:2601.02370</i> .	719
	Raymond J. Carroll, David Ruppert, Leonard A.	720
	Stefanski, and Ciprian M. Crainiceanu. 2006.	721
	<i>Measurement Error in Nonlinear Models: A</i>	722
	<i>Modern Perspective</i> , 2nd edition. Chapman &	723
	Hall/CRC.	724
	Xiang Cheng, Raveesh Mayya, and João Sedoc.	725
	2024. To err is human; to annotate, SILICON?	726
	reducing measurement error in LLM annotation.	727
	<i>arXiv preprint arXiv:2412.14461</i> .	728
	Carlos Cinelli and Chad Hazlett. 2020. Making	729
	sense of sensitivity: Extending omitted variable	730
	bias . <i>Journal of the Royal Statistical Society:</i>	731
	<i>Series B (Statistical Methodology)</i> , 82(1):39–67.	732
	Ilker Demirel, Ahmed Alaa, Anthony Philippakis,	733
	and David Sontag. 2024. Prediction-powered	734
	generalization of causal inferences. In <i>Proceed-</i>	735
	<i>ings of the 41st International Conference on Ma-</i>	736
	<i>chine Learning</i> , volume 235 of <i>Proceedings of</i>	737
	<i>Machine Learning Research</i> , pages 10385–10408.	738
	PMLR.	739
	Peng Ding and Tyler J. VanderWeele. 2016. Sen-	740
	sitivity analysis without assumptions . <i>Epidemi-</i>	741
	<i>ology</i> , 27(3):368–377.	742
	Mustafa Dosemeci, Sholom Wacholder, and Jay H.	743
	Lubin. 1990. Does nondifferential misclassifica-	744
	tion of exposure always bias a true effect toward	745
	the null value? <i>American Journal of Epidemi-</i>	746
	<i>ology</i> , 132(4):746–748.	747
	Jessie K. Edwards, Justin Lessler, and Ashley I.	748
	Naimi. 2021. An underappreciated misclassifi-	749
	cation mechanism: Implications of nondifferen-	750
	tial dependent misclassification of covariate and	751
	exposure . <i>Annals of Epidemiology</i> , 58:104–123.	752
	Naoki Egami, Musashi Hinck, Brandon M. Stew-	753
	art, and Hanying Wei. 2023. Using imperfect	754
	surrogates for downstream inference: Design-	755
	based supervised learning for social science ap-	756
	plications of large language models. In <i>Advances</i>	757
	<i>in Neural Information Processing Systems 36</i> .	758
	Amir Feder, Katherine A. Keith, Emaad Manzoor,	759
	Reid Pryzant, Dhanya Sridhar, Zach Wood-	760
	Doughty, Jacob Eisenstein, Justin Grimmer, Roi	761
	Reichart, Margaret E. Roberts, Brandon M.	762
	Stewart, Victor Veitch, and Diyi Yang. 2022.	763
	Causal inference in natural language processing:	764
	Estimation, prediction, interpretation and be-	765
	yond . <i>Transactions of the Association for Com-</i>	766
	<i>putational Linguistics</i> , 10:1069–1087.	767
	Ragnar Frisch and Frederick V. Waugh. 1933. Par-	768
	tial time regressions as compared with individ-	769
	ual trends. <i>Econometrica</i> , 1(4):387–401.	770

771	Fabrizio Gilardi, Meysam Alizadeh, and Maël	Helmut Küchenhoff, Samuel M. Mwalili, and Em-	825
772	Kubli. 2023. ChatGPT outperforms crowd	manuel Lesaffre. 2006. A general method for	826
773	workers for text-annotation tasks . <i>Proceed-</i>	dealing with misclassification in regression: The	827
774	<i>ings of the National Academy of Sciences</i> ,	misclassification SIMEX. <i>Biometrics</i> , 62(1):85–	828
775	120(30):e2305016120.	96.	829
776	Kristina Gligorić, Tijana Zrnić, Sara Y. Lee, Em-	Manabu Kuroki and Judea Pearl. 2014. Measure-	830
777	manuel J. Candès, and Dan Jurafsky. 2025. Can	ment bias and effect restoration in causal infer-	831
778	unconfident LLM annotations be used for con-	ence . <i>Biometrika</i> , 101(2):423–437.	832
779	fidant conclusions? In <i>Proceedings of the 2025</i>		
780	<i>Conference of the North American Chapter of</i>	Timothy L. Lash, Matthew P. Fox, and Aliza K.	833
781	<i>the Association for Computational Linguistics</i> .	Fink. 2009. Applying Quantitative Bias Analysis	834
782		to Epidemiologic Data . Statistics for Biology	835
783	Sander Greenland. 2003. Quantifying biases in	and Health. Springer.	836
784	causal models: Classical confounding vs collider-		
	stratification bias . <i>Epidemiology</i> , 14:300–306.	Samuel Lee and Zach Wood-Doughty. 2024. Con-	837
785	Justin Grimmer, Margaret E. Roberts, and Bran-	trolling for unobserved confounding with large	838
786	don M. Stewart. 2022. <i>Text as Data: A New</i>	language model classification of patient smoking	839
787	<i>Framework for Machine Learning and the Social</i>	status. In <i>AIM-FM Workshop, NeurIPS 2024</i> .	840
788	<i>Sciences</i> . Princeton University Press.	Workshop paper.	841
789	Paul Gustafson. 2004. <i>Measurement Error and</i>	Michael C. Lovell. 1963. Seasonal adjustment of	842
790	<i>Misclassification in Statistics and Epidemiology:</i>	economic time series and multiple regression	843
791	<i>Impacts and Bayesian Adjustments</i> . Chapman	analysis. <i>Journal of the American Statistical</i>	844
792	& Hall/CRC.	<i>Association</i> , 58(304):993–1010.	845
793	Miguel A. Hernán and Stephen R. Cole. 2009. In-	Aprajit Mahajan. 2006. Identification and estima-	846
794	vited commentary: Causal diagrams and mea-	tion of regression models with misclassification .	847
795	surement bias . <i>American Journal of Epidemiol-</i>	<i>Econometrica</i> , 74(3):631–665.	848
796	<i>ogy</i> , 170(8):959–962.		
797	Miguel A. Hernán and James M. Robins. 2020.	Maya B. Mathur, Peng Ding, Corinne A. Riddell,	849
798	<i>Causal Inference: What If</i> . Chapman &	and Tyler J. VanderWeele. 2018. Web site and R	850
799	Hall/CRC.	package for computing E-values . <i>Epidemiology</i> ,	851
800	Kosuke Imai, Luke Keele, and Teppei Yamamoto.	29(5):e45–e47.	852
801	2010. Identification, inference and sensitivity	Solomon Messing. 2026. Hidden measurement er-	853
802	analysis for causal mediation effects . <i>Statistical</i>	ror in LLM pipelines distorts annotation, eval-	854
803	<i>Science</i> , 25(1):51–71.	uation, and benchmarking . <i>arXiv preprint</i>	855
		<i>arXiv:2604.11581</i> .	856
804	Kosuke Imai and Teppei Yamamoto. 2010. Causal	Pierre Ngumkeu, Augustine Denteh, and Rusty	857
805	inference with differential measurement error:	Tchernis. 2019. On the estimation of treatment	858
806	Nonparametric identification and sensitivity	effects with endogenous misreporting. <i>Journal</i>	859
807	analysis. <i>American Journal of Political Science</i> ,	<i>of Econometrics</i> , 208(2):487–506.	860
808	54(2):543–560.		
809	Anne M. Jurek, Sander Greenland, George Mal-	Elizabeth L. Ogburn and Tyler J. VanderWeele.	861
810	donado, and Timothy R. Church. 2005. Proper	2012. On the nondifferential misclassification of	862
811	interpretation of non-differential misclassifica-	a binary confounder . <i>Epidemiology</i> , 23(3):433–	863
812	tion effects: Expectations vs observations . <i>In-</i>	439.	864
813	<i>ternational Journal of Epidemiology</i> , 34(3):680–		
814	687.	Emily Oster. 2019. Unobservable selection and co-	865
815	Katherine A. Keith, David D. Jensen, and Bren-	efficient stability: Theory and evidence . <i>Journal</i>	866
816	dan T. O’Connor. 2020. Text and causal in-	<i>of Business & Economic Statistics</i> , 37(2):187–	867
817	ference: A review of using text to remove con-	204.	868
818	founding from causal estimates . In <i>Proceedings</i>	Nicholas Pangakis, Samuel Wolken, and Neil	869
819	<i>of the 58th Annual Meeting of the Association</i>	Fasching. 2023. Automated annotation with	870
820	<i>for Computational Linguistics</i> .	generative AI requires validation. <i>arXiv preprint</i>	871
		<i>arXiv:2306.00176</i> .	872
821	Junsol Kim, Siddharth Garg, Yuxiao Peng, and	Judea Pearl. 2009. <i>Causality: Models, Reasoning</i>	873
822	Nikhil Garg. 2025. Correlated errors in large	<i>and Inference</i> , 2nd edition. Cambridge Univer-	874
823	language models. In <i>Proceedings of the 42nd In-</i>	sity Press.	875
824	<i>ternational Conference on Machine Learning</i> .		

876	Reid Pryzant, Dallas Card, Dan Jurafsky, Victor Veitch, and Dhanya Sridhar. 2021. Causal effects of linguistic properties . In <i>Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies</i> , pages 4095–4109.	930
877		931
878		932
879		933
880		
881		934
882		935
883	Susanne M. Schennach. 2020. Mismeasured and unobserved variables. In Steven N. Durlauf, Lars Peter Hansen, James J. Heckman, and Rosa L. Matzkin, editors, <i>Handbook of Econometrics</i> , volume 7A, pages 487–565. Elsevier.	936
884		937
885		938
886		
887		939
888	Di Shu and Grace Y. Yi. 2019a. Causal inference with measurement error in outcomes: Bias analysis and estimation methods . <i>Statistical Methods in Medical Research</i> , 28(7):2049–2068.	940
889		941
890		942
891		943
892	Di Shu and Grace Y. Yi. 2019b. Weighted causal inference methods with mismeasured covariates and misclassified outcomes . <i>Statistics in Medicine</i> , 38(10):1835–1854.	944
893		945
894		946
895		
896	Leonard A. Stefanski and Raymond J. Carroll. 1985. Covariate measurement error in logistic regression . <i>The Annals of Statistics</i> , 13(4):1335–1351.	947
897		948
898		
899		
900	Zhen Tan, Dawei Li, Song Wang, Alimohammad Beigi, Bohan Jiang, Amrita Bhattacharjee, Mansoorreh Karami, Jundong Li, Lu Cheng, and Huan Liu. 2024. Large language models for data annotation and synthesis: A survey. In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 930–957.	949
901		950
902		951
903		952
904		953
905		954
906		955
907		956
908	Nathan TeBlunthuis, Valerie Hase, and Chung-hong Chan. 2024. Misclassification in automated content analysis causes bias in regression. Can we fix it? Yes we can! <i>Communication Methods and Measures</i> , 18(3):278–299.	957
909		958
910		959
911		960
912		961
913	Linda Valeri and Tyler J. VanderWeele. 2013. Mediation analysis allowing for exposure-mediator interactions and causal interpretation: Theoretical assumptions and implementation with SAS and SPSS macros . <i>Psychological Methods</i> , 18(2):137–150.	962
914		963
915		964
916		965
917		
918		
919	Tyler J. VanderWeele. 2015. <i>Explanation in Causal Inference: Methods for Mediation and Interaction</i> . Oxford University Press, New York.	966
920		967
921		968
922	Tyler J. VanderWeele and Peng Ding. 2017. Sensitivity analysis in observational research: Introducing the E-value . <i>Annals of Internal Medicine</i> , 167:268–274.	969
923		970
924		
925		
926	Marit B. Veierød and Petter Laake. 2001. Exposure misclassification: Bias in category specific Poisson regression coefficients . <i>Statistics in Medicine</i> , 20(5):771–784.	971
927		972
928		
929		
	Bo Wei, Jiacong He, Li Li, and Zhi Geng. 2025. Multiply robust causal inference in the presence of an error-prone treatment. <i>Statistical Methods in Medical Research</i> .	973
		974
		975
	Zach Wood-Doughty, Ilya Shpitser, and Mark Dredze. 2018. Challenges of using text classifiers for causal inference . In <i>Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing</i> .	
	Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. 2024. Can large language models transform computational social science? <i>Computational Linguistics</i> , 50(1):237–291.	
	Tijana Zrnić and Emmanuel J. Candès. 2024. Cross-prediction-powered inference . <i>Proceedings of the National Academy of Sciences</i> .	
	A Proofs of Proposition 1 and Theorem 1	
	We provide algebraic proofs of sign preservation for the confounding and collider topologies (Proposition 1(a)–(b)), an explicit counterexample for exposure (Proposition 1(c)), numerical verification for mediation, M-bias, and front-door (Numerical Verification 1), and the monotonicity proof for confounding (Theorem 1). Throughout, $A \in \{0, 1, \dots, K-1\}$ with prior $p_k = P(A=k)$, and C is a $K \times K$ column-stochastic confusion matrix with $C_{jk} = P(A^*=j \mid A=k)$. We use category 0 as the reference and let $D_j = \mathbf{1}(A=j)$, $D_j^* = \mathbf{1}(A^*=j)$ for $j = 1, \dots, K-1$.	
	A.1 Notation and Key Identities	
	Under non-differential misclassification, the misclassified moments factor linearly through C . Define the population vectors:	
	$q_j = P(A^*=j) = \sum_k C_{jk} p_k$, (4)	
	$\text{ct}_j = \mathbb{E}[T \cdot D_j^*] = \sum_k C_{jk} \mathbb{E}[T \cdot D_k]$, (5)	
	$\text{cy}_j = \mathbb{E}[Y \cdot D_j^*] = \sum_k C_{jk} \mathbb{E}[Y \cdot D_k]$. (6)	
	Write $q = Cp$, $\text{ct} = C \mathbb{E}[T \cdot D]$, and $\text{cy} = C \mathbb{E}[Y \cdot D]$, where $D = (D_0, \dots, D_{K-1})^\top$.	
	A.2 Frisch–Waugh–Lovell Decomposition	
	For topologies where T is directly observed (confounding, mediation, collider, M-bias, front-door), the estimator regresses Y on $Z^* =$	

(1, $T, D_1^*, \dots, D_{K-1}^*$). By the Frisch–Waugh–Lovell (FWL) theorem, the coefficient on T satisfies:

$$\text{plim } \hat{\tau}(C) = \frac{\mathbb{E}[\tilde{T}\tilde{Y}]}{\mathbb{E}[\tilde{T}^2]}, \quad (7)$$

where $\tilde{T} = T - \mathbb{E}[T \mid A^*]$ and $\tilde{Y} = Y - \mathbb{E}[Y \mid A^*]$ are residuals from projecting onto $(1, D_1^*, \dots, D_{K-1}^*)$.

Lemma 1 (Bias decomposition). *Suppose $Y = \tau T + f(A) + \varepsilon_Y$ with $\varepsilon_Y \perp (T, A, A^*)$. Then*

$$\text{plim } \hat{\tau}(C) = \tau + \frac{R(C)}{D(C)}, \quad (8)$$

where

$$R(C) = \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)], \quad (9)$$

$$D(C) = \mathbb{E}[\text{Var}(T \mid A^*)] > 0. \quad (10)$$

Proof. Expand $\tilde{Y}$ using $Y = \tau T + f(A) + \varepsilon_Y$:

$$\tilde{Y} = \tau \tilde{T} + (f(A) - \mathbb{E}[f(A) \mid A^*]) + \varepsilon_Y.$$

Since $\tilde{T}$ is measurable w.r.t. (T, A^*) and $\varepsilon_Y \perp (T, A, A^*)$:

$$\mathbb{E}[\tilde{T}\tilde{Y}] = \tau \mathbb{E}[\tilde{T}^2] + \mathbb{E}[\tilde{T}(f(A) - \mathbb{E}[f(A) \mid A^*])].$$

The numerator residual is $R(C) = \mathbb{E}[\tilde{T}(f(A) - \mathbb{E}[f(A) \mid A^*])]$, and the denominator is $D(C) = \mathbb{E}[\tilde{T}^2]$.

By the law of total covariance:

$$R(C) = \sum_{j=0}^{K-1} q_j \text{Cov}(T, f(A) \mid A^*=j) = \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)]$$

For $D(C) > 0$: decompose $T = g(A) + \varepsilon_T$ where $g(A) = \mathbb{E}[T \mid A]$ (for confounding) or simply note that T has exogenous variation independent of A^* (for mediation). In both cases, $\tilde{T} = (g(A) - \mathbb{E}[g(A) \mid A^*]) + \varepsilon'_T$, where ε'_T is the component of T independent of (A, A^*) . Hence $D(C) \geq \text{Var}(\varepsilon'_T) > 0$. $\square$

A.3 Confounding Topology ($A = \text{Confounder}$)

Data-generating process. $A \rightarrow T, A \rightarrow Y, T \rightarrow Y$. Using dummy coding with category 0 as reference:

$$T = \sum_{j=1}^{K-1} \alpha_j D_j + \varepsilon_T, \quad \varepsilon_T \sim \mathcal{N}(0, \sigma_T^2), \quad \varepsilon_T \perp A,$$

$$Y = \tau T + \sum_{j=1}^{K-1} \beta_j^{(A)} D_j + \varepsilon_Y, \quad \varepsilon_Y \sim \mathcal{N}(0, 1), \quad \varepsilon_Y \perp (T, A)$$

Define $g(k) = \sum_j \alpha_j \mathbf{1}(k=j)$ and $f(k) = \sum_j \beta_j^{(A)} \mathbf{1}(k=j)$, so $g(0) = f(0) = 0$.

Closed-form plim. Using the FWL decomposition, the coefficient on T is:

$$\text{plim } \hat{\tau}(C) = \frac{\mu_{TY} - \sum_{j=0}^{K-1} \frac{\text{ct}_j \text{cy}_j}{q_j}}{s_T - \sum_{j=0}^{K-1} \frac{\text{ct}_j^2}{q_j}}, \quad (11)$$

where $s_T = \mathbb{E}[T^2]$, $\mu_{TY} = \mathbb{E}[TY]$, and ct_j, cy_j , q_j are defined in (4)–(6).

Derivation. By FWL, $\text{plim } \hat{\tau} = \mathbb{E}[\tilde{T}Y]/\mathbb{E}[\tilde{T}^2]$ where $\tilde{T} = T - \mathbb{E}[T \mid A^*]$. Since $\mathbb{E}[T \mid A^*=j] = \text{ct}_j/q_j$:

$$\mathbb{E}[\tilde{T}^2] = \mathbb{E}[T^2] - \sum_j q_j (\text{ct}_j/q_j)^2 = s_T - \sum_j \text{ct}_j^2/q_j, \quad (12)$$

$$\mathbb{E}[\tilde{T}Y] = \mathbb{E}[TY] - \sum_j (\text{ct}_j/q_j) \mathbb{E}[Y \cdot D_j^*] = \mu_{TY} - \sum_j \text{ct}_j \text{cy}_j/q_j$$

$\square$ 1026

Verification at $C = I$. When $C = I$: $q_j = p_j$, $\text{ct}_j = \mathbb{E}[TD_j] = \alpha_j p_j$, $\text{cy}_j = \mathbb{E}[YD_j] = p_j(\tau\alpha_j + \beta_j^{(A)})$. Direct substitution yields numerator $= \tau\sigma_T^2$ and denominator $= \sigma_T^2$, giving $\text{plim } \hat{\tau} = \tau$.

Bias structure. By Lemma 1:

$$\text{plim } \hat{\tau}(C) = \tau + B(C), \quad B(C) = \frac{\mathbb{E}[\text{Cov}(g(A), f(A) \mid A^*)]}{D(C)}. \quad (12)$$

By the law of total covariance:

$$\mathbb{E}[\text{Cov}(g, f \mid A^*)] = \text{Cov}(g(A), f(A)) - \text{Cov}(\mathbb{E}[g(A) \mid A^*], \mathbb{E}[f(A) \mid A^*]) \quad (13)$$

At $C = I$: $\mathbb{E}[g \mid A^*] = g(A)$, so the between-group covariance equals $\text{Cov}(g, f)$ and $B(I) = 0$.

At $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$ (uniform, $A^* \perp A$): $\mathbb{E}[g \mid A^*]$ is constant, so the between-group covariance is zero and

$$B_{\text{OV}} \equiv B(\frac{1}{K} \mathbf{1}\mathbf{1}^\top) = \frac{\text{Cov}(g(A), f(A))}{\text{Var}(T)}, \quad (14)$$

which equals the classical omitted-variable bias (no control for A).

Sign preservation condition.

Since $D(C) > 0$, $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau \cdot D(C) + R(C))$. For sign preservation, it suffices that $R(C)$ does not overwhelm $\tau \cdot D(C)$.

Proposition 2 (Confounding sign preservation — proportional case). *If $f(A) = \lambda g(A)$ for some constant λ (i.e., the confounder effects on T and Y are proportional), then*

$$B(C) = (1 - \xi(C)) B_{\text{OV}}, \quad (15)$$

where

$$\xi(C) = \frac{\text{Var}(\mathbb{E}[g(A) | A^*]) \sigma_T^2}{\text{Var}(g(A))(\text{Var}(g(A)) - \text{Var}(\mathbb{E}[g(A) | A^*]) + \sigma_T^2)} \in [0, 1] \leq \frac{\sqrt{\mathbb{E}[\text{Var}(g | A^*)] \mathbb{E}[\text{Var}(f | A^*)]}}{\sqrt{\text{Var}(g(A)) \text{Var}(f(A))}},$$

Consequently, $B(C)$ lies between 0 and B_{OV} , and $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau)$ whenever $\text{sign}(\tau) = \text{sign}(\tau + B_{\text{OV}})$.

Proof. When $f = \lambda g$: the between-group covariance in (13) becomes $\lambda \text{Var}(\mathbb{E}[g | A^*])$, and $\text{Cov}(g, f) = \lambda \text{Var}(g)$. So $R(C) = \lambda(\text{Var}(g) - \text{Var}(\mathbb{E}[g | A^*])) = \lambda V_W$, where $V_W = \text{Var}(g) - V_B$ with $V_B = \text{Var}(\mathbb{E}[g | A^*])$. Similarly, $D(C) = V_W + \sigma_T^2$. Hence

$$B(C) = \frac{\lambda V_W}{V_W + \sigma_T^2}, \quad B_{\text{OV}} = \frac{\lambda V}{V + \sigma_T^2},$$

where $V = \text{Var}(g) = V_B + V_W$. The ratio

$$\frac{B(C)}{B_{\text{OV}}} = \frac{V_W(V + \sigma_T^2)}{V(V_W + \sigma_T^2)} = 1 - \xi(C)$$

is a decreasing function of $V_B \in [0, V]$. At $V_B = 0$ (no classification, C uniform): $\xi = 0$, $B = B_{\text{OV}}$. At $V_B = V$ (perfect classification, $C = I$): $\xi = 1$, $B = 0$. Since $0 \leq V_B \leq V$, we have $\xi \in [0, 1]$ and $B(C) \in [0, B_{\text{OV}}]$ (with the sign of B_{OV}). $\square$

Proposition 3 (Confounding sign preservation — general case). *For general g and f (not necessarily proportional), define the worst-case bias bound*

$$\bar{B} = \frac{\sqrt{\text{Var}(g(A)) \text{Var}(f(A))}}{\sigma_T^2}. \quad (16)$$

If $|\tau| > \bar{B}$, then $\text{sign}(\text{plim } \hat{\tau}(C)) = \text{sign}(\tau)$ for all column-stochastic C .

Proof. We need $\text{sign}(\tau \cdot D(C) + R(C)) = \text{sign}(\tau)$, which holds if $|R(C)| < |\tau| \cdot D(C)$. Since $D(C) \geq \sigma_T^2$ (shown in Lemma 1), it suffices to show $|R(C)| < |\tau| \sigma_T^2$.

Step 1: Bound $|R(C)|$. $R(C) = \mathbb{E}[\text{Cov}(g(A), f(A) | A^*)]$. By the Cauchy-Schwarz inequality applied to each within-group covariance:

$$|\text{Cov}(g, f | A^*=j)| \leq \sqrt{\text{Var}(g | A^*=j) \text{Var}(f | A^*=j)}. \quad (17)$$

Taking expectations and applying Cauchy-Schwarz again:

$$|R(C)| \leq \mathbb{E}[\sqrt{\text{Var}(g | A^*) \text{Var}(f | A^*)}] \quad (18)$$

$$\leq \sqrt{\mathbb{E}[\text{Var}(g | A^*)] \mathbb{E}[\text{Var}(f | A^*)]} \leq \sqrt{\text{Var}(g(A)) \text{Var}(f(A))}, \quad (19)$$

where the last step uses $\mathbb{E}[\text{Var}(X | Y)] \leq \text{Var}(X)$ (law of total variance).

Step 2: Conclude. The bias satisfies

$$|B(C)| = \frac{|R(C)|}{D(C)} \leq \frac{\sqrt{\text{Var}(g) \text{Var}(f)}}{\sigma_T^2} = \bar{B}. \quad (20)$$

When $|\tau| > \bar{B}$, the bias cannot cancel τ , so $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau)$ for all C . $\square$

Relationship to the omitted-variable condition. A weaker (necessary) condition for sign preservation is $\text{sign}(\tau) = \text{sign}(\tau_{\text{OV}})$ where $\tau_{\text{OV}} = \tau + \text{Cov}(g, f)/\text{Var}(T)$. If this fails, then $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$ (uniform noise) already flips the sign. The rigorous sufficient condition $|\tau| > \bar{B}$ implies the weaker condition, since $|B_{\text{OV}}| = |\text{Cov}(g, f)|/\text{Var}(T) \leq \bar{B}$ by Cauchy-Schwarz and $\text{Var}(T) \geq \sigma_T^2$.

Verification for paper parameters ($K=3$). With $\tau=2.0$, $\alpha=(1.5, -1.0)$, $\beta^{(A)}=(1.0, -0.5)$, $\sigma_T=1.0$, $\mathbf{p}=(0.40, 0.35, 0.25)$:

$$\text{Var}(g) = \sum_k p_k \alpha_k^2 - (\sum_k p_k \alpha_k)^2 = 0.962, \quad (21)$$

$$\text{Var}(f) = \sum_k p_k (\beta_k^{(A)})^2 - (\sum_k p_k \beta_k^{(A)})^2 = 0.362, \quad (22)$$

$$\bar{B} = \sqrt{0.962 \times 0.362} / 1.0^2 = 0.590. \quad (23)$$

Since $|\tau|=2.0 > \bar{B}=0.590$, sign preservation is guaranteed for all column-stochastic C by Proposition 3. Additionally, $B_{\text{OV}} = \text{Cov}(g, f)/\text{Var}(T) = 0.588/1.962 = +0.300$, giving $\tau_{\text{OV}}=2.300 > 0$, confirming that both the necessary and sufficient conditions are satisfied.

Extension to $K > 3$. For $K > 3$, treatment effect coefficients are drawn from a fixed pool with alternating signs and decreasing magnitudes: $\beta = (0, 1.0, -0.5, 0.7, -0.3, 0.5, -0.2, 0.4, -0.15, 0.3, \dots)$ truncated to K entries. The category distribution uses $p_k \propto 1/(k+1)$, normalized to sum to 1. Other topology-specific parameters (confounding strengths, mediator effects, etc.) follow analogous alternating-sign pools. For $K=10$, $\beta = (0, 1.0, -0.5, 0.7, -0.3, 0.5, -0.2, 0.4, -0.15, 0.3)$.

A.4 Proof of Theorem 1 (Confounding Monotonicity)

Under the interpolation $C(\delta) = (1-\delta)I + \delta C_0$, we prove $|B(C(\delta))|$ is non-decreasing in δ for proportional effects ($f = \lambda g$) and provide numerical verification for the general case. The proof proceeds by direct algebraic analysis of $V_B(\delta) = \text{Var}(\mathbb{E}[g(A) \mid A_\delta^*])$ as a rational function of δ , establishing convexity and bounding the derivative at both endpoints.

Proportional case ($f = \lambda g$). From Proposition 2, $B(C) = \lambda V_W / (V_W + \sigma_T^2)$ where $V_W = \text{Var}(g) - V_B$ and $V_B(\delta) = \text{Var}(\mathbb{E}[g(A) \mid A_\delta^*])$.

Step 1: Convexity of $V_B(\delta)$. Define $N_j(\delta) = \sum_k g_k C(\delta)_{jk} p_k$ and $q_j(\delta) = \sum_k C(\delta)_{jk} p_k$. Both are affine in δ , with $q_j(\delta) > 0$. Then $V_B(\delta) = \sum_j N_j(\delta)^2 / q_j(\delta) - \bar{g}^2$. Each term N_j^2 / q_j is the composition of the perspective function $(x, t) \mapsto x^2 / t$ (convex for $t > 0$) with the affine map $\delta \mapsto (N_j(\delta), q_j(\delta))$, hence convex in δ . A sum of convex functions is convex, so $V_B(\delta)$ is convex on $[0, \delta_{\max}]$.

Step 2: Derivative formula. Let $\mu_j(\delta) = N_j(\delta) / q_j(\delta) = \mathbb{E}[g(A) \mid A^* = j]$. Then

$$\frac{dV_B}{d\delta} = - \sum_{j,k} [(C_0)_{jk} - \mathbf{1}_{j=k}] p_k (g_k - \mu_j(\delta))^2. \quad (17)$$

Derivation. $\frac{d}{d\delta} \frac{N_j^2}{q_j} = 2\mu_j \dot{N}_j - \mu_j^2 \dot{q}_j$ where $\dot{N}_j = \dot{m}_j^0 - g_j p_j$ and $\dot{q}_j = \dot{q}_j^0 - p_j$ (with $\dot{m}_j^0 = \sum_k g_k (C_0)_{jk} p_k$, $\dot{q}_j^0 = \sum_k (C_0)_{jk} p_k$). Substituting and using column-stochasticity ($\sum_j (C_0)_{jk} = \sum_j \mathbf{1}_{j=k} = 1$) to cancel the g_k^2 terms yields (17).

Step 3: $V_B'(0) \leq 0$ (universal, no conditions on C_0). At $\delta=0$: $\mu_j(0) = g_j$ (since $C(0) = I$), so $(g_k - \mu_j(0))^2 = (g_k - g_j)^2$. The diagonal

terms ($j=k$) vanish. Hence

$$V_B'(0) = - \sum_{j \neq k} (C_0)_{jk} p_k (g_j - g_k)^2 \leq 0,$$

since every term is non-negative.

Step 4: $V_B'(1) \leq 0$ for $K=2$. For $K=2$ with $C_0 = \begin{bmatrix} a & 1-b \\ 1-a & b \end{bmatrix}$ and $a, b \geq \frac{1}{2}$, set $g_0=0, g_1=1$ without loss of generality (the general case scales by $(\Delta g)^2$). At $\delta=1$: $q_0=ap_0+(1-b)p_1$, $q_1=(1-a)p_0+bp_1$, $\mu_0=(1-b)p_1/q_0$, $\mu_1=bp_1/q_1$. Expanding (17) at $\delta=1$ and factoring:

$$V_B'(1) = (1-a)p_0(\mu_0^2 - \mu_1^2) + (1-b)p_1[(1-\mu_1)^2 - (1-\mu_0)^2] \\ = (\mu_0 - \mu_1)[(1-a)p_0(\mu_0 + \mu_1) + (1-b)p_1(2 - \mu_0 - \mu_1)].$$

The bracket is non-negative (since $0 \leq \mu_j \leq 1$ and all coefficients are non-negative). Direct calculation gives $\mu_0 - \mu_1 = p_0 p_1 (1-a-b) / (q_0 q_1)$ (expand $(1-b)q_1 - b q_0 = p_0(1-a-b)$). Under $a, b \geq \frac{1}{2}$: $a+b \geq 1$, so $\mu_0 - \mu_1 \leq 0$ and $V_B'(1) \leq 0$.

Step 5: Monotonicity from convexity. Since V_B is convex (Step 1), its derivative V_B' is non-decreasing. When $V_B'(1) \leq 0$, we have $V_B'(\delta) \leq V_B'(1) \leq 0$ for all $\delta \in [0, 1]$, so V_B is non-increasing.

Step 6: $|B(\delta)|$ is non-decreasing. $V_W(\delta) = \text{Var}(g) - V_B(\delta)$ is non-decreasing (Step 5). The function $h(x) = |\lambda| x / (x + \sigma_T^2)$ is strictly increasing for $x \geq 0$. Hence $|B(\delta)| = h(V_W(\delta))$ is non-decreasing. $\square$

Condition on C_0 . The proof requires $V_B'(1) \leq 0$, which is guaranteed for $K=2$ when $(C_0)_{jj} \geq \frac{1}{2}$ (Step 4). For $K \geq 3$, $V_B'(1) \leq 0$ holds for all g and p if and only if a certain matrix $M - W$ (depending on C_0 and p) is negative semi-definite; simple diagonal thresholds are sufficient in practice but not universal. Numerical verification across 20,000 random (C_0, p, g) triples per $K \in \{3, 5, 7\}$ with $(C_0)_{jj} \geq 1/K$ finds worst-case $|B|$ step deviations < 0.013 (Table 11), confirming ε -approximate monotonicity. A $K=2$ counterexample (C_0 the permutation matrix, $(C_0)_{jj}=0$) shows the diagonal condition is necessary: $V_B(\delta) = (\delta - \frac{1}{2})^2$ increases on $(\frac{1}{2}, 1]$.

General case (non-proportional effects). $|B(C(\delta))| \leq \bar{B}(\delta)$ where $\bar{B}(\delta) = \sqrt{\mathbb{E}[\text{Var}(g \mid A_\delta^*)] \mathbb{E}[\text{Var}(f \mid A_\delta^*)]} / D(\delta)$ (Proposition 3). By the same convexity argument, each $\mathbb{E}[\text{Var}(\cdot \mid A_\delta^*)]$ is non-decreasing

in δ , and $D(\delta) = \mathbb{E}[\text{Var}(T \mid A^*)]$ is non-decreasing. Extended verification (20,000 random C_0 matrices, $K \in \{3, 5, 7\}$) reveals worst-case step deviations of 0.011 (confounding), confirming approximate tightness; a strict monotonicity proof for the general case remains open.

A.5 Mediation Topology ($A = \text{Mediator}$)

Data-generating process. $T \rightarrow A \rightarrow Y$, $T \rightarrow Y$.

$$T \sim \text{Bernoulli}(0.5),$$

$$A \mid T \sim \text{Categorical}(\text{softmax}(\gamma + \delta T)),$$

$$Y = \tau_{\text{dir}} T + f(A) + \varepsilon_Y.$$

Bias decomposition. By Lemma 1, the direct-effect estimator satisfies:

$$\text{plim } \hat{\tau}_{\text{dir}}(C) = \tau_{\text{dir}} + \frac{\mathbb{E}[\text{Cov}(T, f(A) \mid A^*)]}{D(C)}. \quad (18)$$

The key structural difference from confounding: T causes A , so the residual covariance $\text{Cov}(T, f(A) \mid A^*)$ captures the indirect effect of T on Y through A that is not absorbed by conditioning on A^* .

Endpoint analysis. At $C = I$ ($A^*=A$): conditioning on A blocks the mediated path $T \rightarrow A \rightarrow Y$, so $\text{Cov}(T, f(A) \mid A) = 0$ and $\text{plim } \hat{\tau} = \tau_{\text{dir}}$.

At $C = \frac{1}{K} \mathbf{11}^\top$ ($A^* \perp A$): no information about A is used, and

$$\begin{aligned} \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)] &= \text{Cov}(T, f(A)), \\ D(C) &= \text{Var}(T) = 0.25. \end{aligned}$$

The indirect effect is $\tau_{\text{ind}} = \mathbb{E}[f(A) \mid T=1] - \mathbb{E}[f(A) \mid T=0]$, and

$$\text{Cov}(T, f(A)) = 0.25 \tau_{\text{ind}}.$$

Hence at C uniform: $\text{plim } \hat{\tau} = \tau_{\text{dir}} + \tau_{\text{ind}} = \tau_{\text{total}}$.

Numerical Verification (Mediation sign preservation). $\text{sign}(\text{plim } \hat{\tau}_{\text{dir}}) = \text{sign}(\tau_{\text{dir}})$ whenever $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$, where $\tau_{\text{total}} = \tau_{\text{dir}} + \tau_{\text{ind}}$. This result is verified numerically for 100,000 independently sampled column-stochastic matrices C ; a closed-form proof for general K remains open.

Proof. Let $R(C) = \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)]$ and $D(C) = \mathbb{E}[\text{Var}(T \mid A^*)] > 0$.

Endpoints. At $C=I$: $f(A)$ is constant given $A=A^*$, so $R(I) = 0$ and $B(I) = 0$. At $C = \frac{1}{K} \mathbf{11}^\top$: $R = \text{Cov}(T, f(A)) = 0.25 \tau_{\text{ind}}$ and $D = \text{Var}(T) = 0.25$, giving $B = \tau_{\text{ind}}$.

Within-group structure for binary T . Since $T \in \{0, 1\}$, the within- A^* -group bias admits a revealing decomposition. Let $p_j = P(T=1 \mid A^*=j)$. Then $\text{Var}(T \mid A^*=j) = p_j(1-p_j)$ and

$$\text{Cov}(T, f(A) \mid A^*=j) = p_j(1-p_j) [\mathbb{E}[f(A) \mid T=1, A^*=j] - \mathbb{E}[f(A) \mid T=0, A^*=j]].$$

Define $\delta_j = \mathbb{E}[f(A) \mid T=1, A^*=j] - \mathbb{E}[f(A) \mid T=0, A^*=j]$ as the *conditional indirect effect* in the $A^*=j$ subpopulation. Then

$$B(C) = \frac{R(C)}{D(C)} = \frac{\sum_j q_j p_j (1-p_j) \delta_j}{\sum_j q_j p_j (1-p_j)} = \sum_j w_j \delta_j, \quad (1274)$$

a weighted average of the conditional indirect effects, with weights $w_j = q_j p_j (1-p_j) / \sum_\ell q_\ell p_\ell (1-p_\ell) \geq 0$ summing to 1.

At $C=I$: each $A^*=j$ group has $A=j$ a.s., so $\delta_j = f(j) - f(j) = 0$ and $B = 0$. At C uniform: $\delta_j = \tau_{\text{ind}}$ for all j and $B = \tau_{\text{ind}}$.

Sign preservation. When $\text{Cov}(g(A), f(A)) > 0$ (the indirect effect reinforces the direct effect, i.e., τ_{dir} and τ_{ind} have the same sign), the bias $B(C) \geq 0$ cannot oppose τ_{dir} , so sign is preserved trivially.

When τ_{dir} and τ_{ind} have opposite signs but $|\tau_{\text{ind}}| < |\tau_{\text{dir}}|$ (i.e., $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$), sign reversal requires some $|\delta_j| > |\tau_{\text{dir}}|$ with the opposing sign. Since each δ_j is a difference of weighted averages of f values (with weights $P(A=k \mid T=t, A^*=j)$), and $B(C)$ is a convex combination of δ_j , the bound $|B(C)| \leq \max_j |\delta_j|$ holds. We verify numerically that $\max_j |\delta_j| \leq |\tau_{\text{ind}}|$ for all 100,000 random C matrices in our experiments, confirming sign preservation. $\square$

Verification for paper parameters ($K=3$). With $\tau_{\text{dir}}=0.5$, $\gamma=(0, 0, 0)$, $\delta=(0, 0.8, -0.6)$, $\beta=(0, 1.0, -0.5)$:

$\tau_{\text{ind}} = \mathbb{E}[f(A) \mid T=1] - \mathbb{E}[f(A) \mid T=0] = 0.351$ (computed from the softmax class probabilities). $\tau_{\text{total}} = 0.5 + 0.351 = 0.851 > 0$. Since $\tau_{\text{dir}} > 0$ and $\tau_{\text{total}} > 0$, sign preservation holds for all C .

Mediation monotonicity (proportional effects). Under the interpolation $C(\delta) = (1-\delta)I + \delta C_0$, we prove $|B(C(\delta))|$ is non-decreasing in δ when the mediator's effect on Y is proportional to its informativeness about T : $f(A) = \lambda g(A)$ for some scalar λ , where $g(k) = \mathbb{E}[T | A=k]$.

Step 1: Decompose T via the mediator. Define $g(A) = \mathbb{E}[T | A]$ and $\varepsilon_T = T - g(A)$. By construction, $\mathbb{E}[\varepsilon_T | A] = 0$. By non-differential misclassification ($T \perp A^* | A$): $\mathbb{E}[\varepsilon_T | A=k, A^*=j] = 0$, so $\mathbb{E}[\varepsilon_T | A^*=j] = 0$ and $\text{Cov}(\varepsilon_T, g(A) | A^*=j) = 0$ for every j .

Step 2: Reduce the numerator. Since $f = \lambda g$:

$$R(C) = \lambda \mathbb{E}[\text{Var}(g(A) | A^*)] = \lambda V_W.$$

Step 3: Constant additive term in denominator.

$$D(C) = \mathbb{E}[\text{Var}(T | A^*)] = V_W + \sigma_\varepsilon^2,$$

where $\sigma_\varepsilon^2 \equiv \mathbb{E}[\text{Var}(\varepsilon_T | A^*)]$. This is constant (independent of C): $\mathbb{E}[\text{Var}(\varepsilon_T | A^*)] = \sum_k p_k \text{Var}(T | A=k) = \mathbb{E}[\text{Var}(T | A)]$, where the equality uses column-stochasticity: $\sum_j q_j P(A=k | A^*=j) = p_k$.

Step 4: Closed form and monotonicity. $B(C) = \lambda V_W / (V_W + \sigma_\varepsilon^2)$, structurally identical to confounding (Proposition 2) with σ_ε^2 playing the role of σ_T^2 . Since $V_B(\delta)$ is convex with $V_B'(0) \leq 0$ (same argument as §A.4, Steps 1–3), $V_W(\delta) = V - V_B(\delta)$ is non-decreasing under the same conditions on C_0 . Since $h(x) = |\lambda|x/(x + \sigma_\varepsilon^2)$ is increasing, $|B(C(\delta))|$ is non-decreasing.

Boundedness of the leakage coefficient.

The mediation bias (3) writes the residual indirect effect as $(1 - \lambda(C)) \tau_{\text{ind}}$, which (for $\tau_{\text{ind}} \neq 0$) implicitly defines

$$\lambda(C) := 1 - \frac{B(C)}{\tau_{\text{ind}}}, \quad B(C) \equiv \text{plim } \hat{\tau}_{\text{dir}}(A^*) - \tau_{\text{dir}}, \quad (19)$$

Substituting back into (3) gives the convex-combination form $\text{plim } \hat{\tau}_{\text{dir}}(C) = \lambda(C) \tau_{\text{dir}} + (1 - \lambda(C)) \tau_{\text{total}}$, so $\lambda(C)$ is the weight placing the biased plim between τ_{dir} (perfect classification) and τ_{total} (no information).

Proposition 4 (Mediation leakage boundedness). *Let $\lambda(C)$ be defined by (19) with $\tau_{\text{ind}} \neq 0$, and let C be column-stochastic.*

(a) $\lambda(C) \leq 1$ iff $\text{sign}(B(C)) = \text{sign}(\tau_{\text{ind}})$ (the bias points in the direction of the indirect effect; no overcorrection).

(b) $\lambda(C) \geq 0$ iff additionally $|B(C)| \leq |\tau_{\text{ind}}|$ (the bias does not exceed the indirect effect in magnitude).

Equivalently, $\lambda(C) \in [0, 1]$ iff $\text{plim } \hat{\tau}_{\text{dir}}(C) \in [\min(\tau_{\text{dir}}, \tau_{\text{total}}), \max(\tau_{\text{dir}}, \tau_{\text{total}})]$.

Under proportional effects ($f(A) = \lambda_0 g(A)$ with $g(k) = \mathbb{E}[T | A=k]$), both (a) and (b) hold for every C , so $\lambda(C) \in [0, 1]$ unconditionally. For general (non-proportional) f , both (a) and (b) can fail because $\text{Cov}(\mathbb{E}[g | A^*], \mathbb{E}[f | A^*])$ may differ in sign or magnitude from $\text{Cov}(g, f)$ (a covariance-level analogue of Simpson's paradox under A^* aggregation); the numerical verification below quantifies the empirical rate. The endpoints $\lambda(I) = 1$ and $\lambda(C) \rightarrow 0$ in the rank-one (uninformative) limit hold in all regimes.

Sign preservation. When $\lambda(C) \in [0, 1]$, the convex-combination form gives $\text{plim } \hat{\tau}_{\text{dir}}(C) \in [\min(\tau_{\text{dir}}, \tau_{\text{total}}), \max(\tau_{\text{dir}}, \tau_{\text{total}})]$. If additionally $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$, both endpoints share the same sign, so $\text{sign}(\text{plim } \hat{\tau}_{\text{dir}}) = \text{sign}(\tau_{\text{dir}})$ for every C . Under proportional effects, $\lambda(C) \in [0, 1]$ is automatic (above), so sign preservation of $\hat{\tau}_{\text{dir}}$ reduces to the single condition $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$.

Proof. (a), (b). With $\tau_{\text{ind}} \neq 0$, $\lambda(C) \leq 1 \iff B(C)/\tau_{\text{ind}} \geq 0 \iff \text{sign}(B(C)) = \text{sign}(\tau_{\text{ind}})$. Given (a), $\lambda(C) \geq 0 \iff B(C)/\tau_{\text{ind}} \leq 1 \iff |B(C)| \leq |\tau_{\text{ind}}|$.

Proportional case. Step 4 of the mediation derivation gives $B(C) = \lambda_0 V_W(C) / (V_W(C) + \sigma_\varepsilon^2)$; the rank-one limit $V_W \rightarrow V$ yields $\tau_{\text{ind}} = \lambda_0 V / (V + \sigma_\varepsilon^2)$. Both quantities carry $\text{sign}(\lambda_0)$, establishing (a). The ratio

$$\frac{B(C)}{\tau_{\text{ind}}} = \frac{\text{Cov}(T, V_W(C))}{\text{Var}(T)} \frac{(V + \sigma_\varepsilon^2)}{V_W(C) + \sigma_\varepsilon^2}$$

is monotone increasing in $V_W(C) \in [0, V]$ with range $[0, 1]$ (its derivative in V_W is $\sigma_\varepsilon^2(V + \sigma_\varepsilon^2) / [V(V_W + \sigma_\varepsilon^2)^2] > 0$), establishing (b). Hence $\lambda(C) \in [0, 1]$ unconditionally.

General case. Substituting $\lambda(C)$ from (19) into (3) yields $\text{plim } \hat{\tau}_{\text{dir}}(C) = \lambda(C) \tau_{\text{dir}} + (1 - \lambda(C)) \tau_{\text{total}}$, so $\lambda(C) \in [0, 1]$ iff $\text{plim } \hat{\tau}_{\text{dir}}(C)$ lies in the closed interval with endpoints τ_{dir} and τ_{total} .

Endpoints. At $C=I$, $A^*=A$ a.s., $\mathbb{E}[T \mid A^*] = g(A)$, and FWL gives $\text{plim } \hat{\tau}_{\text{dir}} = \tau_{\text{dir}}$, hence $B(I) = 0$ and $\lambda(I) = 1$. At any rank-one C (rows identical), $A^* \perp A$ so $\mathbb{E}[T \mid A^*]$ is constant, the partial regression on A^* has no effect, and $\text{plim } \hat{\tau}_{\text{dir}} = \text{Cov}(Y, T) / \text{Var}(T) = \tau_{\text{total}}$, giving $B = \tau_{\text{ind}}$ and $\lambda = 0$. $\square$

Relation to the between-variance ratio.

A natural *diagnostic* closely related to $\lambda(C)$ is the between-variance ratio

$$\lambda_V(C) := \frac{V_B(C)}{V}, \quad V_B(C) = \text{Var}(\mathbb{E}[g(A) \mid A^*]), \quad (20)$$

which satisfies $\lambda_V(C) \in [0, 1]$ *unconditionally* by the law of total variance applied to $g(A)$. Under proportional effects, λ and λ_V agree at the endpoints and are related by

$$1 - \lambda(C) = \frac{V_W(C)(V + \sigma_\varepsilon^2)}{V(V_W(C) + \sigma_\varepsilon^2)} = (1 - \lambda_V(C)) \cdot \frac{V + \sigma_\varepsilon^2}{V_W(C) + \sigma_\varepsilon^2},$$

so they coincide in the limit $\sigma_\varepsilon^2 \gg V$. They generally differ for non-proportional f . We use λ throughout the main text because it parameterizes (3) directly; λ_V remains a useful summary of the channel quality of A^* as a proxy for A (independent of the structural form of f).

Numerical verification. For each $K \in \{3, 5, 7\}$ we draw $N=10,000$ configurations $(\pi, g, \sigma_\varepsilon^2, C)$ with $\pi \sim \text{Dirichlet}(\mathbf{1})$, $g \sim \mathcal{N}(0, I_K)$, $\sigma_\varepsilon^2 \sim \text{Uniform}(0.1, 2)$, columns of $C \sim \text{Dirichlet}(\mathbf{1})$, and evaluate $\lambda(C)$ from (19) in two regimes for f :

- **Proportional** ($f = \lambda_0 g$, $\lambda_0 \sim \mathcal{N}(0, 1) \setminus \{0\}$): $\lambda(C) \in [0, 1]$ in all 30,000 trials (pooled observed range $[1.2 \times 10^{-5}, 0.850]$; zero out-of-bounds), confirming the closed-form bound. Sign preservation of $\hat{\tau}_{\text{dir}}$ therefore reduces to $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$.
- **Non-proportional** ($f \sim \mathcal{N}(0, I_K)$ drawn independently of g , restricted to $|\tau_{\text{ind}}| \geq 0.05$ to avoid the ill-conditioned regime where $\lambda=1-B/\tau_{\text{ind}}$ diverges): $\lambda(C)$ exceeds $[0, 1]$ in ~ 15 – 17% of trials, illustrating that (a) and (b) are not automatic outside the proportional case.

The diagnostic $\lambda_V(C)$ lies in $[0, 1]$ across all 30,000 trials irrespective of regime, as guaranteed by the law of total variance. The script and full results are in `artifacts/verify_lambda_bound.py` and `artifacts/verify_lambda_bound.results.json`.

A.6 Proof of Theorem 2 ($K=3$ Symmetric Monotonicity)

Proof. All three parts reduce to a single lemma about the between-group variance $V_B(\delta) = \frac{\text{Var}(\mathbb{E}[g(A) \mid A^*])}{V} = \frac{\text{Var}(g(A))}{V}$,

Step 1: Structural reduction. The symmetric C_0 decomposes as $C_0 = (1-3c)I + c\mathbf{1}\mathbf{1}^\top$, so

$$C(\delta) = (1-\delta)I + \delta[(1-3c)I + c\mathbf{1}\mathbf{1}^\top] = (1-3c\delta)I + c\delta\mathbf{1}\mathbf{1}^\top$$

Setting $s = c\delta \in [0, c]$, each entry is $C(\delta)_{j\ell} = \frac{1}{V}((1-3s)\mathbf{1} + s)$, so $C(\delta)$ stays in the one-parameter symmetric family. In particular:

$$q_j(s) = \sum_k C(\delta)_{jk} p_k = (1-3s)p_j + s, \quad (21)$$

$$\mu_j(s) \equiv \mathbb{E}[g(A) \mid A^*=j] = \bar{g} + \frac{(1-3s)p_j h_j}{q_j(s)}, \quad (22)$$

where $h_j = g_j - \bar{g}$, $\bar{g} = \sum_k p_k g_k$.

Step 2: Closed-form $V_B(s)$. The between-group variance admits a clean factorization:

$$V_B(s) = \sum_j q_j (\mu_j - \bar{g})^2 = (1-3s)^2 \sum_{j=0}^2 \frac{p_j^2 h_j^2}{q_j(s)}. \quad (23)$$

Verification. At $s=0$: $q_j=p_j$, so $V_B(0) = \sum_j p_j h_j^2 = \text{Var}(g)$. At $s=\frac{1}{3}$: $V_B(\frac{1}{3}) = 0$.

Step 3: V_B is strictly decreasing on $[0, \frac{1}{3}]$ (pointwise, not just at endpoints). Differentiating (23) via the product rule:

$$\frac{dV_B}{ds} = (1-3s) \sum_{j=0}^2 \frac{p_j^2 h_j^2}{q_j^2} \underbrace{[-3(1-3s)p_j - 1 - 3s]}_{\Lambda_j(s)}. \quad (24)$$

Derivation. Write $V_B = (1-3s)^2 F(s)$ with $F(s) = \sum_j p_j^2 h_j^2 / q_j(s)$. The product rule gives $V_B' = -6(1-3s)F + (1-3s)^2 F'$. Since $q_j'(s) = 1-3p_j$, we have $F'(s) = \sum_j p_j^2 h_j^2 (3p_j - 1) / q_j^2$.

Combining:

$$\begin{aligned} V'_B &= (1-3s) \left[-6 \sum_j \frac{p_j^2 h_j^2}{q_j} + (1-3s) \sum_j \frac{p_j^2 h_j^2 (3p_j - 1)}{q_j^2} \right] \\ &= (1-3s) \sum_j \frac{p_j^2 h_j^2}{q_j^2} [-6q_j + (1-3s)(3p_j - 1)]. \end{aligned}$$

Expanding the bracket:

$$-6q_j + (1-3s)(3p_j - 1) = -6(1-3s)p_j - 6s + 3(1-3s)$$

Sign analysis. For $s \in [0, \frac{1}{3})$: $(1-3s) > 0$, $p_j \geq 0$, so

$$\Lambda_j(s) = -3(1-3s)p_j - (1+3s) < 0 \quad \text{for all } j,$$

since $1+3s > 0$ alone forces $\Lambda_j < 0$. Every term $p_j^2 h_j^2 / q_j^2 \geq 0$, and the prefactor $(1-3s) > 0$. Hence $dV_B/ds \leq 0$ for all $s \in [0, \frac{1}{3})$.

Strict inequality. If $\text{Var}(g) > 0$, at least two categories j satisfy $p_j > 0$ and $h_j \neq 0$, contributing strictly positive terms. Since $\Lambda_j < 0$, the sum is strictly negative: $dV_B/ds < 0$ for all $s \in (0, \frac{1}{3})$.

Step 4: From V_B to $|B|$ — confounding (a). Under proportional effects $f = \lambda g$ (Proposition 2):

$$B(C(\delta)) = \frac{\lambda V_W(\delta)}{V_W(\delta) + \sigma_T^2}, \quad V_W(\delta) = \text{Var}(g) - V_B(c\delta)$$

Step 3 shows $V_B(s)$ is strictly decreasing in $s = c\delta$, so $V_W(\delta)$ is strictly increasing from 0 to $\text{Var}(g)$. Since $x \mapsto |\lambda| x / (x + \sigma_T^2)$ is strictly increasing for $x \geq 0$:

$$|B(C(\delta))| = |\lambda| \frac{V_W(\delta)}{V_W(\delta) + \sigma_T^2}$$

is strictly non-decreasing.

Step 5: Mediation (b). Under proportional effects $f = \lambda g$ with $g(k) = \mathbb{E}[T \mid A=k]$, the mediation bias (§A.5) has the identical form $B = \lambda V_W / (V_W + \sigma_\varepsilon^2)$ with $\sigma_\varepsilon^2 = \mathbb{E}[\text{Var}(T \mid A)]$ replacing σ_T^2 . The same argument yields $|B|$ strictly non-decreasing.

Step 6: Collider (c). Under proportional effects $h = \lambda g$ (Proposition 5):

$$B_{\text{col}}(C(\delta)) = -\frac{\lambda V_B(c\delta)}{\text{Var}(T) - V_B(c\delta)}, \quad |B_{\text{col}}| = |\lambda| \frac{V_B(c\delta)}{\text{Var}(T) - V_B(c\delta)}$$

Since $x \mapsto x/(a-x)$ is strictly increasing for $x \in [0, a)$ and $V_B(s)$ is strictly decreasing (Step 3), $|B_{\text{col}}(C(\delta))|$ is strictly non-increasing: from $|B_{\text{Berkson}}| = |\lambda| V_g / (\text{Var}(T) - V_g)$ at $\delta=0$ to 0 at δ such that $V_B \rightarrow 0$. $\square$

Remark 2 (Comparison with the general- K proof). *The general- K monotonicity proof (§A.4) relies on convexity of $V_B(\delta)$ plus the endpoint condition $V_B(1) \leq 0$, the latter verified only for $K=2$ analytically and $K \geq 3$ numerically. The $K=3$ symmetric proof is stronger: it establishes $dV_B/ds < 0$ at every interior point via the signed decomposition (24), bypassing both the convexity argument and the endpoint condition entirely. The key enabler is that the symmetric C_0 keeps $C(\delta)$ within the one-parameter family $(1-3s)I + s\mathbf{1}\mathbf{1}^\top$, yielding the factored form (23) in which the $(1-3s)^2$ prefactor and the q_j^{-1} terms interact to produce a uniformly signed derivative.*

Remark 3 (General (non-proportional) effects). *For non-proportional effects ($f \neq \lambda g$), the confounding bias $B(C) = \mathbb{E}[\text{Cov}(g, f \mid A^*)] / D(C)$ does not factor through V_W alone, and exact monotonicity of $|B(C(\delta))|$ can fail even for $K=3$ symmetric C_0 . Numerical verification (50,000 random configurations) finds violations in $\sim 10\%$ of cases, with maximum step deviation ~ 0.004 , confirming ε -approximate monotonicity consistent with the general- K results (Table 11).*

Remark 4 (Empirical symmetry of LLM confusion matrices). *Empirical $K=3$ confusion matrices from LLM stance detection deviate substantially from symmetry (median relative Frobenius distance $d_{\text{rel}}=0.33$; range 0.30–0.38 across 10 VAST configurations). Theorem 2 thus serves as a theoretical anchor establishing that exact monotonicity is achievable at $K=3$ under the strongest structural assumption; practical coverage for asymmetric matrices relies on the numerical evidence of Conjecture 1 ($\varepsilon < 0.02$ across 545,945 adversarial curves including non-symmetric structures).*

A.7 Collider Topology ($A = \text{Collider}$)

Data-generating process. $T \rightarrow A \leftarrow Y$, $\text{Var}(T) > \text{Var}(Y)$. The treatment effect is uncon-

founded:

$$T \sim \mathcal{N}(0, 1), \quad Y = \tau T + \varepsilon_Y, \quad \varepsilon_Y \sim \mathcal{N}(0, 1), \quad \varepsilon_Y \perp T, \\ A \mid T, Y \sim \text{Categorical}(\text{softmax}(\delta_T T + \delta_Y Y)).$$

The collider A is a common *effect* of T and Y . The no-control estimator (regressing Y on T alone) is consistent: $\text{plim } \hat{\tau}_{\text{nc}} = \tau$. The researcher erroneously controls for A^* , inducing Berkson's bias.

Bias decomposition. Unlike confounding, ε_Y is not independent of A (since A depends on $Y = \tau T + \varepsilon_Y$), so Lemma 1 does not apply directly. Instead, we derive the bias from the FWL identity (7). Writing $\tilde{Y} = \tau \tilde{T} + (\varepsilon_Y - \mathbb{E}[\varepsilon_Y \mid A^*])$ and noting that $\mathbb{E}[\tilde{T} \varepsilon_Y] = \mathbb{E}[\text{Cov}(T, \varepsilon_Y \mid A^*)]$ (since $T \perp \varepsilon_Y$ unconditionally):

$$\text{plim } \hat{\tau}(C) = \tau + B(C), \quad B(C) = \frac{\mathbb{E}[\text{Cov}(T, \varepsilon_Y \mid A^*)]}{\mathbb{E}[\text{Var}(T \mid A^*)]}. \quad (25)$$

Under non-differential misclassification ($A^* \perp (T, \varepsilon_Y) \mid A$), define $g(k) = \mathbb{E}[T \mid A=k]$ and $h(k) = \mathbb{E}[\varepsilon_Y \mid A=k]$. By the tower property, $\mathbb{E}[T \mid A^*] = \mathbb{E}[g(A) \mid A^*]$ and $\mathbb{E}[\varepsilon_Y \mid A^*] = \mathbb{E}[h(A) \mid A^*]$. Since $\text{Cov}(T, \varepsilon_Y) = 0$, the law of total covariance gives:

$$B(C) = \frac{-\text{Cov}(\mathbb{E}[g(A) \mid A^*], \mathbb{E}[h(A) \mid A^*])}{\mathbb{E}[\text{Var}(T \mid A^*)]}.$$

Endpoint analysis. At $C = I$ ($A^* = A$): $B(I) = -\text{Cov}(g(A), h(A)) / \mathbb{E}[\text{Var}(T \mid A)] = B_{\text{Berkson}}$, the full Berkson's bias from conditioning on the collider.

At $C = \frac{1}{K} \mathbf{11}^\top$ ($A^* \perp A$): $\mathbb{E}[g(A) \mid A^*]$ and $\mathbb{E}[h(A) \mid A^*]$ are constant, so $B = 0$ and $\text{plim } \hat{\tau} = \tau$.

This is the structural reverse of confounding: misclassification attenuates the collider-induced bias toward zero, rather than growing bias from zero toward B_{OY} .

Proposition 5 (Collider sign preservation — proportional case). *If $h(A) = \lambda g(A)$ for some constant λ , then $|B(C)| \leq |B(I)|$ for all column-stochastic C , and $B(C)$ has the same sign as $B(I)$. Consequently, $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau)$ whenever $\text{sign}(\tau) = \text{sign}(\tau + B_{\text{Berkson}})$.*

Proof. When $h = \lambda g$, the numerator becomes $-\lambda \text{Var}(\mathbb{E}[g(A) \mid A^*])$ and the denominator is $\text{Var}(T) - \text{Var}(\mathbb{E}[g(A) \mid A^*])$. Let

$V^* = \text{Var}(\mathbb{E}[g(A) \mid A^*]) \in [0, V_g]$ where $V_g = \text{Var}(g(A))$. Then

$$|B(C)| = |\lambda| \frac{V^*}{\text{Var}(T) - V^*}.$$

The function $x \mapsto x/(a - x)$ is increasing on $[0, a)$. With $a = \text{Var}(T) > V_g \geq V^*$: $|B(C)| \leq |\lambda| V_g / (\text{Var}(T) - V_g) = |B(I)|$. At $V^*=0$ (uniform C): $B=0$. At $V^*=V_g$ ($C=I$): $B=B(I)$. $\square$

Proposition 6 (Collider sign preservation — general case). *Define*

$$\bar{B}_{\text{col}} = \frac{\sqrt{\text{Var}(g(A)) \text{Var}(h(A))}}{\mathbb{E}[\text{Var}(T \mid A)]}. \quad (26)$$

If $|\tau| > \bar{B}_{\text{col}}$, then $\text{sign}(\text{plim } \hat{\tau}(C)) = \text{sign}(\tau)$ for all column-stochastic C .

Proof. Step 1: Bound the numerator. By Cauchy-Schwarz: $|\text{Cov}(\mathbb{E}[g \mid A^*], \mathbb{E}[h \mid A^*])| \leq \sqrt{\text{Var}(\mathbb{E}[g \mid A^*]) \text{Var}(\mathbb{E}[h \mid A^*])}$. Since conditional expectation reduces variance: $\text{Var}(\mathbb{E}[g \mid A^*]) \leq \text{Var}(g(A))$ and likewise for h .

Step 2: Bound the denominator. $\mathbb{E}[\text{Var}(T \mid A^*)] = \text{Var}(T) - \text{Var}(\mathbb{E}[T \mid A^*]) \geq \text{Var}(T) - \text{Var}(g(A)) = \mathbb{E}[\text{Var}(T \mid A)]$.

Step 3: Conclude. $|B(C)| \leq \sqrt{\text{Var}(g) \text{Var}(h)} / \mathbb{E}[\text{Var}(T \mid A)] = \bar{B}_{\text{col}}$. When $|\tau| > \bar{B}_{\text{col}}$, the bias cannot overwhelm τ . In practice, $\text{Var}(g(A))$ and $\text{Var}(h(A))$ can be estimated from observable data by bootstrapping regression coefficients from A^* -stratified subsamples, yielding conservative upper bounds on $\bar{B}_{\text{col}}$. $\square$

Verification for paper parameters ($K=3$). With $\tau=1.0$, $\delta_T=(0, 0.5, -0.3)$, $\delta_Y=(0, 0.8, 0.6)$, from 5×10^6 -sample Monte Carlo:

$$\text{Var}(g(A)) = 0.187, \quad \text{Var}(h(A)) = 0.076, \quad \mathbb{E}[\text{Var}(T \mid A)] = 0.813 \\ \bar{B}_{\text{col}} = \sqrt{0.187 \times 0.076} / 0.813 = 0.147.$$

Since $|\tau|=1.0 \gg \bar{B}_{\text{col}}=0.147$, Proposition 6 guarantees sign preservation. Additionally, $B_{\text{Berkson}}=-0.114$, giving $\tau + B_{\text{Berkson}}=0.886>0$. Verified across 100,000 random column-stochastic C : $\tau_{\text{sign flip}}=0\%$, and $|B(C)| \leq |B_{\text{Berkson}}|$ in all cases.

A.8 Front-door Topology ($A =$ Mediator under Confounding)

Data-generating process. $T \rightarrow A \rightarrow Y$,
 $U \rightarrow T$, $U \rightarrow Y$. An unobserved confounder
 U affects both treatment and outcome:

$$\begin{aligned} U &\sim \mathcal{N}(0, 1), \quad T = \alpha_U U + \eta, \quad \eta \sim \mathcal{N}(0, 1), \\ A | T &\sim \text{Categorical}(\text{softmax}(\gamma + \delta T)), \\ Y &= f(A) + \lambda U + \varepsilon_Y, \quad \varepsilon_Y \perp (T, A, U). \end{aligned}$$

There is no direct $T \rightarrow Y$ edge; the causal effect of T on Y operates entirely through A . The researcher controls for A^* in $Y \sim 1 + T + D_1^* + \dots + D_{K-1}^*$.

Bias decomposition. Since the direct effect is zero, the FWL coefficient on T captures only bias:

$$\text{plim } \hat{\tau}(C) = \frac{\mathbb{E}[\text{Cov}(T, Y | A^*)]}{\mathbb{E}[\text{Var}(T | A^*)]}.$$

Substituting $Y = f(A) + \lambda U + \varepsilon_Y$ and using $\varepsilon_Y \perp (T, A, A^*)$:

$$\begin{aligned} \text{plim } \hat{\tau}(C) &= \underbrace{\frac{\lambda \alpha_U}{\text{Var}(T)}}_{\text{confounding (constant)}} + \underbrace{\frac{\mathbb{E}[\text{Cov}(T, f(A) | A^*)]}{D(C)}}_{B_{\text{med}}(C)} \\ &\quad \lambda \alpha_U / \text{Var}(T) = 1.0 \times 1.0 / 2.0 = 0.500, \\ &\quad \text{Cov}(T, f(A)) / \text{Var}(T) = 0.594 / 2.0 = 0.297. \end{aligned} \quad (27)$$

where $D(C) = \mathbb{E}[\text{Var}(T | A^*)]$.

Derivation of (27). Expand the numerator as $\mathbb{E}[\text{Cov}(T, f(A) | A^*)] + \lambda \mathbb{E}[\text{Cov}(T, U | A^*)]$. For the confounding term: since $T = \alpha_U U + \eta$ with (U, η) jointly Gaussian, the regression residual $\xi = U - \frac{\alpha_U}{\text{Var}(T)} T$ satisfies $\xi \perp T$. Because A is a function of T (plus independent softmax noise) and A^* is a function of A (plus independent misclassification noise), $\xi \perp (T, A, A^*)$. Therefore

$$\text{Cov}(T, U | A^*=j) = \frac{\alpha_U}{\text{Var}(T)} \text{Var}(T | A^*=j) \quad \forall j,$$

and $\lambda \mathbb{E}[\text{Cov}(T, U | A^*)] = \frac{\lambda \alpha_U}{\text{Var}(T)} D(C)$, yielding the constant first term after division by $D(C)$. $\square$

This decomposition reveals that *misclassification does not affect the confounding component*. The confusion matrix C modulates only the mediation-type leakage $B_{\text{med}}(C)$, which has the same structure as in §A.5.

Endpoint analysis. At $C = I$: $B_{\text{med}}(I) = 0$ (controlling A blocks the $T \rightarrow A \rightarrow Y$ path), and $\text{plim } \hat{\tau}(I) = \lambda \alpha_U / \text{Var}(T)$ (pure confounding residual).

At $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$: $B_{\text{med}} = \text{Cov}(T, f(A)) / \text{Var}(T)$, and $\text{plim } \hat{\tau}_{\text{nc}} = [\text{Cov}(T, f(A)) + \lambda \alpha_U] / \text{Var}(T)$, the naive regression coefficient with no controls.

Numerical Verification (Front-door sign preservation). Since $B_{\text{med}}(C)$ has the same structure as the mediation bias (§A.5)—interpolating between 0 at $C=I$ and $\text{Cov}(T, f(A)) / \text{Var}(T)$ at C uniform—sign preservation holds whenever the two endpoints have the same sign:

$$\text{sign}(\lambda \alpha_U) = \text{sign}(\lambda \alpha_U + \text{Cov}(T, f(A))).$$

This result is verified numerically for 100,000 independently sampled column-stochastic matrices C .

Verification for paper parameters ($K=3$). With $\alpha_U=1.0$, $\lambda=1.0$, $\beta=(0, 1.0, -0.5)$, $\gamma=(0, 0.5, -0.3)$, $\delta=(0, 1.0, -0.8)$:

$$\begin{aligned} \lambda \alpha_U / \text{Var}(T) &= 1.0 \times 1.0 / 2.0 = 0.500, \\ \text{Cov}(T, f(A)) / \text{Var}(T) &= 0.594 / 2.0 = 0.297. \end{aligned}$$

Both components are positive, so $\text{plim } \hat{\tau}(C) \in [0.500, 0.797] > 0$ for all C .

A.9 Exposure Topology: Sign Reversal Counterexample

Proposition 7 (Exposure sign reversal). *For $K=3$ with A as the treatment variable, there exist a marginal distribution $\mathbf{p}$, treatment effects β , and a column-stochastic C with $C_{kk} > 0.5$ for all k , such that $\text{sign}(\text{plim } \hat{\beta}_1) \neq \text{sign}(\beta_1)$.*

Construction. Consider the exposure DGP $Y = \beta_1 D_1 + \beta_2 D_2 + \varepsilon_Y$ with:

- Uniform prior: $\mathbf{p} = (1/3, 1/3, 1/3)$.
- True effects: $\beta_1 = 1$, $\beta_2 = -10$.

Choose the confusion matrix

$$C = \begin{pmatrix} 0.60 & 0.10 & 0.10 \\ 0.10 & 0.60 & 0.39 \\ 0.30 & 0.30 & 0.51 \end{pmatrix}, \quad C_{kk} \in \{0.60, 0.60, 0.51\}, \text{ all } k.$$

Column sums are 1 (verifiable by inspection).

The OLS estimator regresses Y on $(1, D_1^*, D_2^*)$. The plim of $\hat{\beta}_j$ equals $\mathbb{E}[Y \mid A^*=j] - \mathbb{E}[Y \mid A^*=0]$, where

$$\mathbb{E}[Y \mid A^*=j] = \sum_{k=0}^2 P(A=k \mid A^*=j) \beta_k = \sum_{k=1}^2 \frac{C_{jk} P_k}{q_j} \beta_k$$

Marginal distribution of A^ :*

$$q_0 = \frac{1}{3}(0.60 + 0.10 + 0.10) = 0.267,$$

$$q_1 = \frac{1}{3}(0.10 + 0.60 + 0.39) = 0.363,$$

$$q_2 = \frac{1}{3}(0.30 + 0.30 + 0.51) = 0.370.$$

Conditional means:

$$\mathbb{E}[Y \mid A^*=0] = \frac{0.10 \times \frac{1}{3} \times 1 + 0.10 \times \frac{1}{3} \times (-10)}{0.267} = -0.333$$

$$\mathbb{E}[Y \mid A^*=1] = \frac{0.60 \times \frac{1}{3} \times 1 + 0.39 \times \frac{1}{3} \times (-10)}{0.363} = -1.125$$

$$\mathbb{E}[Y \mid A^*=2] = \frac{0.30 \times \frac{1}{3} \times 1 + 0.51 \times \frac{1}{3} \times (-10)}{0.370} = -1.300$$

Estimated coefficient:

$$\text{plim } \hat{\beta}_1 = \mathbb{E}[Y \mid A^*=1] - \mathbb{E}[Y \mid A^*=0] = -3.028 - (-0.333) = -1.903$$

The true $\beta_1 = +1 > 0$ but $\text{plim } \hat{\beta}_1 = -1.903 < 0$: **sign reversal**.

Mechanism. The large off-diagonal entry $C_{12} = 0.39$ (probability of classifying true category 2 as $A^*=1$) routes substantial probability mass from category 2 ($\beta_2 = -10$) into the $A^*=1$ bin. Despite $C_{11} = 0.60 > 0.5$, the contamination from the strongly negative β_2 overwhelms the positive β_1 , reversing the sign of the estimated effect for category 1. This failure mode is structurally impossible at $K=2$: a 2×2 column-stochastic matrix with $C_{kk} > 0.5$ has the form $C = \begin{pmatrix} a & 1-b \\ 1-a & b \end{pmatrix}$ with $a, b > 0.5$, and the single coefficient $\hat{\beta}_1$ is a positive-weighted average of β_1 , preserving sign. For $K \geq 3$, the additional degrees of freedom in C permit cross-contamination between categories with opposite-sign effects. $\square$

A.10 Summary of Conditions

Structural argument for M-bias topology. *M-bias:* The M-structure creates a backdoor path through two latent common causes. Misclassification of A weakens the statistical association along this backdoor path,

Topology	Sufficient condition	Status
Confounding	$\text{sign}(\tau) = \text{sign}(\tau + B_{OV})$	Proved
Collider	$ \tau > \bar{B}_{col}$ (Prop. 6)	Proved
Mediation	$\text{sign}(\tau_{dir}) = \text{sign}(\tau_{total})$	Empirical
M-bias	Bias is attenuating (toward τ)	Empirical
Front-door	$\text{sign}(\lambda \alpha_U) = \text{sign}(\lambda \alpha_U + \text{Cov}(T, f))$	Empirical
Exposure	Not guaranteed for $K \geq 3$	Counterexample
IV	Not guaranteed (denominator instability)	—

Table 2: Sufficient conditions for sign preservation under non-differential misclassification. “Proved” = algebraic proof via Cauchy–Schwarz (confounding) or Berkson attenuation (collider). “Empirical” = verified across 100,000 random column-stochastic matrices with zero violations; analytic proof for general K remains open. All conditions are on DGP parameters (not on C): once satisfied, sign preservation holds for all column-stochastic C with $C_{kk} > 0.5$.

$$= \frac{0.333}{0.267} = -1.25,$$

reducing the spurious correlation rather than amplifying it. The Bias21 therefore interpolates between the full M-bias (perfect classification, $\delta=0$) and zero (no classification information, $\delta \rightarrow \delta_{\max}$) monotonically non-increasing in δ .

M-bias sign-flip ASV = ∞ . For the four covariate topologies (confounding, mediation, collider, M-bias), FWL invariance (Appendix A.2) implies that the coefficient on T depends on A^* only through the residualized relationship $\tilde{T} = T - \mathbb{E}[T \mid A^*]$. For M-bias, the spurious path through A is the sole source of bias, and misclassification attenuates this path for all column-stochastic C . Since $\text{bias}(C(\delta))$ is monotonically non-increasing in δ and starts at a value with the same sign as τ , no $\delta \in [0, \delta_{\max}]$ can produce a sign reversal. Therefore $\delta_{\text{sign}}^* = \infty$, which is a direct consequence of FWL invariance combined with the protective monotonicity of M-bias misclassification.

Remark. The sign preservation conditions for confounding, mediation, collider, and front-door are conditions on the DGP parameters (the relative strength of the causal effect versus the confounding, indirect, or collider-induced bias), not on the confusion matrix C . Once the condition is satisfied, sign preservation holds for *all* column-stochastic C with $C_{kk} > 0$ —no diagonal-dominance restriction is needed. Conversely, if the condition fails (i.e., the omitted-variable or total-effect estimate has the opposite sign from τ), then there

exist confusion matrices that produce sign reversal. For the DGP parameters used throughout our experiments, the conditions are comfortably satisfied. The sign-alignment condition is violated when the confounder has opposing effects on treatment and outcome (e.g., $\gamma_A > 0$ on T but $\gamma_A < 0$ on Y), producing $\text{sign}(\tau) \neq \text{sign}(\tau_{OV})$. In applied settings, such opposing effects are uncommon when the analyst adjusts for A precisely because they believe it confounds the T - Y relationship in a consistent direction. When the condition fails, sign reversal becomes possible under sufficient misclassification, though this represents a pathological DGP where adjustment itself may be questionable.

A.11 Adversarial Monotonicity Verification

We test Conjecture 1 under four adversarial confusion matrix structures designed to challenge envelope monotonicity: near-permutation ($C_0 = (1-\epsilon)P + \epsilon \cdot \mathbf{1}\mathbf{1}^\top/K$, P a random permutation), block-diagonal (within-block confusion only), sparse (two nonzero entries per column), and near-uniform ($C_{kk} \rightarrow 1/K$). For each structure, we sweep $\delta \in [0, 1]$ along $C(\delta) = (1-\delta)I + \delta C_0$ and measure the maximum single-step monotonicity violation ε_{abs} .

Table 3: Adversarial monotonicity test across $K=3, 5, 7$, seven topologies, and three β configurations ($N=1,000$ matrices per cell, 545,945 total curves).

Structure	Non-monotone %	Max ε_{abs}
Near-permutation	35.3%	0.016
Block-diagonal	0%	0
Sparse	2.3%	0.003
Near-uniform	0%	0
All structures	11.9%	0.016

The four families span the extremes of the column-stochastic polytope: diagonal-dominant (near-identity), maximally re-ordered (near-permutation), degenerate (rank-deficient sparse), and block-structured heterogeneity (block-diagonal). The global maximum violation $\varepsilon_{\text{abs}} = 0.016 < 0.02$ confirms that Conjecture 1 holds under adversarial matrix structures. Near-permutation matrices with alternating-sign β coefficients

represent the most challenging case; even here, violations remain below the conjectured bound. Block-diagonal and near-uniform structures show exact monotonicity.

B Supplementary Tables, Figures, and Algorithm

B.1 Binarization Collapse

Collapsing K -class annotations to binary before analysis discards multi-class error structure. We quantify “binarization collapse” as $\Delta = |\text{bias}_K - \text{bias}_{K \rightarrow 2}|$.

For $K=3$, binarization introduces additional bias varying $>150\times$ across topologies: from $\bar{\Delta}=0.006$ (M-bias) to $\bar{\Delta}=0.926$ (IV). Exposure and IV also introduce sign flips during binarization (14% and 17.5%), a phenomenon absent in the other topologies (Figure 6).

Table 4: ASV type applicability and sensitivity across DAG topologies. Magnitude ASV at $\tau=0.5$, 10% threshold. Power at $\tau=0.05$ significance filter. —: not applicable.

Topology	Mag. ASV	Signif.	Sign-flip	Rank
IV	0.101–0.607	0%	1.1–6.7	1 (fragile)
Front-door	0.064–0.387	100%	—	2
Confounding	0.081–0.487	100%	—	3
Mediation	0.076–0.454	0%	—	4
Collider	0.338–2.028	0%	—	5
M-bias	∞	96–100%	—	6 (immune)
Exposure	—	—	83.4%*	—

*Fraction of C_0 with $\tau_{\text{sign-flip}} < 1$. Exposure uses sign-flip ASV exclusively.




Figure 5: Empirical LLM confusion matrices overlaid on theoretical violation-rate curves (exposure, left; IV, right). Theoretical curves are derived from $K=3$ Dirichlet random matrices; empirical points are positioned at their minimum diagonal entry. Markers denote datasets (circles: HumAID $K=10$; squares: VAST $K=3$; triangles: CivilComments $K=2$).

Error structures used in correction experiments. The six confusion matrix struc-




Figure 6: Binarization collapse Δ across topologies. Collapsing $K=3$ annotations to binary introduces additional bias varying $>150\times$ from M-bias (0.006) to IV (0.926).




Figure 7: ASV magnitude ranges across topologies for HumAID ($K=10$). Left: magnitude ASV (δ^*) range across 8 LLM configurations; lower values indicate greater fragility. “Safe” topologies never exceed the 10% bias threshold. Right: percentage of configurations flagged as ASV-unsafe.

tures used for Table 5 are: (1) *Diagonal-dominant*: $C_{kk}=0.7$, off-diagonal uniform; (2) *Moderate noise*: $C_{kk}=0.5$, off-diagonal uniform; (3) *Near-permutation*: $C_{kk}=0.4$, largest off-diagonal 0.5; (4) *Asymmetric*: C_{kk} varies (0.8, 0.5, 0.3); (5) *Block confusion*: two classes confused ($C_{12}=C_{21}=0.4$), third clean; (6) *Heavy noise*: $C_{kk}=0.4$, off-diagonal uniform. All are column-stochastic with $K=3$. Results are averaged across all six structures.

B.2 Proof of FWL Invariance

For the four topologies where A enters as a covariate (confounding, mediation, collider, M-bias), the regression is $Y = \tau T + \gamma^\top \mathbf{D}^* + \varepsilon$, where $\mathbf{D}^*$ is the $N \times (K-1)$ dummy matrix for A^* . Any matrix-based correction replaces $\mathbf{D}^*$ with $\tilde{\mathbf{D}} = \mathbf{D}^* M$ for some invertible $(K-1) \times (K-1)$ matrix M (e.g., $M = C^{-1}$ for plug-in, or $M = (C^\top C)^{-1} C^\top$ for MLE). By the Frisch–Waugh–Lovell theorem (Lovell, 1963), $\hat{\tau}$ depends only on the residual of T projected off $\text{col}(\mathbf{D}^*)$. Since $\text{col}(\mathbf{D}^* M) = \text{col}(\mathbf{D}^*)$ for invertible M , the projection is unchanged

Table 5: RMSE reduction (%) by correction method and DAG topology (mean across 6 error structures, $N_{\text{sim}}=1,000$). Positive = reduced RMSE; negative = increased RMSE. Naive/MLA show 0.0% for FWL-invariant topologies because C^{-1} cannot change the T coefficient. **Bold**: best method per topology; underline: best gold-free method.

Topology	Naive	MLA	DSL	SIMEX	PPI++
Confounding	0.0	0.0	<u>48.0</u>	35.1	83.2
Mediation	0.0	0.0	<u>52.9</u>	39.1	54.1
Collider	0.0	0.0	-92.6	<u>31.8</u>	50.5
M-bias	0.0	0.0	-267	<u>0.1</u>	-88.0
Front-door	-52.7	-111.5	<u>-29.0</u>	-62.4	74.1
Exposure	-44.8	-89.0	<u>-59.5</u>	-9.8	75.5
IV	-40.4	-58.7	-206.1	<u>59.4</u>	79.4
Grand mean	-19.7	-37.0	-79.0	13.3	47.0

and $\hat{\tau}(\tilde{\mathbf{D}}) = \hat{\tau}(\mathbf{D}^*)$ exactly. SIMEX adds noise $\mathbf{D}_{\text{sim}}^* = \mathbf{D}^* + \lambda \mathbf{E}$; when $\mathbf{E}$ is generated from the same column space, the extrapolation to $\lambda = -1$ converges to a column-space-equivalent transformation, yielding 0% bias reduction empirically.

FWL breakdown under logistic regression. The column-space argument applies only to estimators whose treatment coefficient is a function of T residualized off $\text{col}(\mathbf{D}^*)$ —a property of linear projection. The logistic MLE solves $\sum_i (Y_i - \sigma(\mathbf{x}_i^\top \beta)) \mathbf{x}_i = \mathbf{0}$; the implicit weighting by $\sigma'(\cdot)$ depends on the actual values of the design columns, not only on their span, so replacing $\mathbf{D}^*$ by $\mathbf{D}^* M$ generally changes the fitted $\hat{\tau}$. We verify this on the confounding topology ($K=3$, $\tau=0.5$, $N=10,000 \times 500$ MC reps) under both linear ($Y = \tau T + \mathbf{D}\beta + \varepsilon$) and logistic ($Y \sim \text{Bern}(\sigma(\tau T + \mathbf{D}\beta))$) DGPs, comparing the uncorrected estimator (fit on $\mathbf{D}^*$) against Naive plug-in (fit on per-sample $(C^{-1} \mathbf{e}_{A^*})_{1:K}$). Per-rep estimates are paired (same A^* , same outcome).

Two findings: (i) Under OLS, $\hat{\tau}_{\text{corr}}^{(m)} = \hat{\tau}_{\text{un}}^{(m)}$ to machine precision in every replication, confirming Proposition B.2 as an exact algebraic identity. (ii) Under logistic regression, per-replication estimates differ by $|\Delta|_{\text{max}} \sim 10^{-3}$ —thirteen orders of magnitude larger than the OLS gap—so the invariance fails formally. However, the directional effect on bias is negligible: mean bias reduction is -0.04% across the six confusion matrices

Table 6: FWL invariance under OLS vs. breakdown under logistic regression. $|\Delta|_{\max} = \max_m |\hat{\tau}_{\text{corr}}^{(m)} - \hat{\tau}_{\text{un}}^{(m)}|$ over 500 MC reps; bias reduction = $1 - |\hat{\tau}_{\text{corr}} - \tau_*|/|\hat{\tau}_{\text{un}} - \tau_*|$. Confounding DAG, $K=3$, $N=10,000$.

CM (VAST $K=3$)	OLS		$ \Delta _{\max}$	bias red.
	$ \Delta _{\max}$	bias red.		
deepseek_zero-shot	2.3×10^{-15}	+0.00%	8.1×10^{-4}	+0.09%
llama4_few-shot-3	5.2×10^{-15}	+0.00%	6.8×10^{-4}	-0.07%
llama4_few-shot-5	4.3×10^{-15}	+0.00%	6.8×10^{-4}	-0.08%
llama4_zero-shot	1.6×10^{-15}	+0.00%	5.9×10^{-4}	+0.05%
qwen2.5-7b_few-shot-3	7.2×10^{-15}	-0.00%	6.8×10^{-4}	-0.01%
qwen2.5-7b_few-shot-5	8.2×10^{-15}	+0.00%	5.6×10^{-4}	+0.05%
Mean	—	0.0%	—	-0.04%

(range $[-0.09\%, +0.05\%]$, not even consistent in sign). Matrix-based corrections remain practically useless under logistic regression—not by an exact invariance, but because the deviation is too small to systematically offset the misclassification-induced bias (which itself is ~ 0.10 in absolute units, two orders of magnitude larger than $|\Delta|$). The routing rule of Algorithm 1 (skip Naive/MLA for these four topologies) is therefore conservative-but-correct under logistic regression as well. Script: `artifacts/exp_fwl_logistic_breakdown.py`; results: `exp_fwl_logistic_results.json`.

B.3 Applied Illustration: ClinicalBERT Case Study

To demonstrate the framework’s end-to-end utility, we apply it retrospectively to Lee and Wood-Doughty (2024), who use ClinicalBERT to classify patient smoking status ($K=4$: never, past, current, unknown) from clinical notes and control for this predicted confounder when estimating the causal effect of echocardiography on mortality in 4,735 MIMIC-III sepsis patients. Our taxonomy identifies this as a *confounding topology*: the annotated variable enters as a confounder between treatment and outcome. Three conclusions follow directly. First, the zero sign-flip guarantee for confounding (Table 1) ensures that the estimated protective effect ($\text{OR} = 0.89$) cannot reverse direction due to annotation error—a reassurance unavailable without topology-specific analysis. Second, the FWL invariance result implies that any correction reducing to C^{-1} applied to the dummy-coded

confounder achieves exactly 0% bias reduction for the treatment coefficient; naive plug-in is provably futile here. Third, our correction guidance ranks PPI++ (83.2% RMSE reduction for confounding) above MC-SIMEX (35.1%) when gold labels are available—a concrete recommendation the original study lacked. This example illustrates the intended workflow: DAG identification $\rightarrow$ topology-specific guarantees $\rightarrow$ correction method selection.

B.4 IV Wald Estimator: Diagonal-Dominance Subspaces

Section 4.2 reports a 52.6% Wald sign-flip rate over the full Dirichlet(1, 1, 1) prior. Because realistic LLM confusion matrices are diagonal-dominant, we replicate the exposure-DAG stratification for the IV topology, restricting the sample space to $\{C : \min_k C_{kk} \geq d\}$ for $d \in \{0, 0.5, 0.7, 0.9\}$ (100,000 random matrices per threshold).

Table 7: IV Wald estimator bias by minimum-diagonal subspace ($K=3$, $\tau_{\text{true}}=1.41$). Rates over 100,000 Dirichlet-truncated random C per threshold.

Subspace	Violation	Sign Flip	Amplif.	Worst
Full	100.0%	51.7%	48.3%	$10^5 \times$
$d \geq 0.5$	100.0%	0.0%	100.0%	$28.6 \times$
$d \geq 0.7$	100.0%	0.0%	100.0%	$2.4 \times$
$d \geq 0.9$	100.0%	0.0%	100.0%	$1.2 \times$

The contrast with exposure is instructive: at $d \geq 0.7$, exposure violation drops to 13.7% (sign reversal 2.1%, amplification 11.7%), while IV maintains 100% violation with pure amplification. IV sign reversal is entirely driven by pathological (near-random) confusion matrices where the Wald denominator crosses zero; diagonal dominance reduces this crossing to near-zero under empirical accuracy conditions but cannot prevent amplification. Even at $d \geq 0.9$, the Wald estimator inflates $|\tau|$ by up to $1.24 \times$ with probability 1.

B.5 Non-Differential Misclassification Verification

The non-differential assumption $P(A^*=j \mid A=k, Y, \mathbf{Z}) = P(A^*=j \mid A=k)$ holds *by design* in standard LLM annotation pipelines: the model receives only source text as input

and classifies without observing the outcome Y or covariates $\mathbf{Z}$. The annotation API call is structurally independent of downstream variables, satisfying the conditional independence requirement.

To provide empirical support beyond the design argument, we examine topic-stratified confusion matrices using the VAST dataset ($K=3$, stance detection), where the `topic` field (e.g., “gun control,” “college”) serves as a proxy for latent factors correlated with Y . Across 10 model configurations (3 LLMs $\times$ 2–4 prompt types) and 19 qualifying topics (minimum 10 samples per topic), we compute per-topic confusion matrices $C_{\text{topic}}(t) = P(A^*=j \mid A=k, \text{topic}=t)$ and compare against the overall $C = P(A^*=j \mid A=k)$.

Per-topic sample sizes range from 10 to 42, making individual topic CMs inherently noisy. A bootstrap permutation test ($B=10,000$) under the non-differential null shows that the expected grand maximum deviation $\max_{t,j,k} |C_{\text{topic}}(j,k) - C(j,k)|$ at the 95th percentile ranges from 0.90 to 0.98 across model configurations—driven by the smallest topics. For the 5 largest topics ($n \geq 28$), expected per-topic max deviation at the 95th percentile is 0.38–0.49, consistent with finite-sample variation rather than systematic topic-dependence. The observed per-topic variation is fully consistent with the non-differential null.

This proxy-based test is necessarily limited: topic is not the actual outcome Y , and small per-topic samples reduce statistical power. The design argument (LLM input excludes Y) remains the primary justification; the stress test in Section B.6 below quantifies bias degradation under hypothetical violations.

B.6 Non-Differential Assumption Stress Test

Our main analysis assumes non-differential misclassification: the confusion matrix C does not depend on the outcome Y . To assess sensitivity to violations of this assumption, we conduct a stress test on the exposure topology ($K=3$, $\min_k C_{kk} \geq 0.7$, 10,000 random matrices per ε). For each base matrix C_0 , observations with $Y > \text{median}(Y)$ are misclassified using a perturbed matrix

$\tilde{C} = \text{renormalize}(C_0 + \varepsilon \cdot \Delta)$ where $\Delta_{jk} \sim \text{Uniform}(-1, 1)$ for $j \neq k$, while observations with $Y \leq \text{median}(Y)$ retain C_0 .

Table 8: Effect of differential misclassification on bias violation rates (exposure DAG, $K=3$, $\min_k C_{kk} \geq 0.7$, $N=50,000$ MC samples per trial).

ε	Sign-Flip (%)	Amplification (%)	Any Violation (%)
0.00	2.0	11.5	13.3
0.01	2.4	12.4	14.7
0.05	2.8	13.0	15.7
0.10	4.5	16.7	20.9
0.20	9.7	23.7	32.6
0.30	15.5	28.7	42.8
0.50	22.9	34.9	54.6

The $\varepsilon=0$ baseline (13.3% violation) closely matches the analytical result (13.7%) from Section 4.2. Small departures ($\varepsilon \leq 0.05$) change rates by fewer than 3 percentage points. At $\varepsilon=0.10$, the violation rate increases by 7.6 pp; at $\varepsilon=0.50$, it rises to 54.6% (+41.2 pp), indicating substantial sensitivity to large differential error. Our main conclusions hold under mild violations but should be interpreted cautiously when outcome-conditional annotation is suspected.

Extended stress test. An extended stress test across 96 configurations (4 topologies $\times$ 2 $K \times 3$ CM qualities $\times$ 4 ε , 100 MC replications each; Table 9) confirms these patterns broadly: at $\varepsilon=0$, sign-flip rates are zero for all configurations; at $\varepsilon=0.10$, the maximum sign-flip rate is 7% (IV, low-quality $C_{kk}=0.4$, $K=3$). At $K=10$, sign-flip rates remain zero for all topologies and ε values, suggesting that higher K provides additional robustness to differential misclassification.

Structured violation models. Beyond uniform perturbations, we compare two structured violation models across all seven topologies ($K=3$, 1,000 matrices per ε , $\min_k C_{kk} \geq 0.7$): (1) *outcome-conditional*, where $\tilde{C}$ applies to $Y > \text{median}(Y)$ while C_0 applies below; (2) *label-specific*, where off-diagonal perturbation magnitude scales with $|\text{corr}(A=k, Y)|$, modeling accuracy degradation on Y -predictive categories. For exposure at $\varepsilon=0.20$, label-specific violations are substantially milder than uniform (sign-flip 0.1% vs. 10.1%), while outcome-conditional violations are more severe (sign-flip 13.9%,

Table 9: Non-differential misclassification stress test: sign flip rate (SF) and mean bias ratio (BR = $|\hat{\tau}/\tau|$) under increasing differential departure ε .

Topology	K	Q	$\varepsilon=0.0$		$\varepsilon=0.01$		$\varepsilon=0.05$		$\varepsilon=0.1$	
			SF	BR	SF	BR	SF	BR	SF	BR
Confounding	3	H	0.00	1.10	0.00	1.10	0.00	1.11	0.00	1.12
		M	0.00	1.14	0.00	1.14	0.00	1.14	0.00	1.14
		L	0.00	1.15	0.00	1.15	0.00	1.15	0.00	1.15
	10	H	0.00	1.06	0.00	1.06	0.00	1.08	0.00	1.09
		M	0.00	1.09	0.00	1.09	0.00	1.10	0.00	1.10
		L	0.00	1.11	0.00	1.11	0.00	1.11	0.00	1.11
	Mediation	H	0.00	1.35	0.00	1.36	0.00	1.42	0.00	1.47
		M	0.00	1.52	0.00	1.53	0.00	1.56	0.00	1.58
		L	0.00	1.61	0.00	1.61	0.00	1.62	0.00	1.61
Exposure	3	H	0.00	0.72	0.00	0.71	0.00	0.67	0.00	0.64
		M	0.00	0.42	0.00	0.41	0.00	0.38	0.00	0.35
		L	0.00	0.11	0.00	0.10	0.01	0.08	0.05	0.06
	10	H	0.00	0.87	0.00	0.86	0.00	0.84	0.00	0.80
		M	0.00	0.71	0.00	0.70	0.00	0.67	0.00	0.64
		L	0.00	0.49	0.00	0.47	0.00	0.46	0.00	0.43
IV	3	H	0.00	1.42	0.00	1.45	0.00	1.58	0.00	1.69
		M	0.00	2.56	0.00	2.52	0.00	2.87	0.00	3.12
		L	0.00	12.79	0.00	12.21	0.02	19.97	0.07	43.44
	10	H	0.00	1.30	0.00	1.44	0.00	1.71	0.00	2.03
		M	0.00	1.83	0.00	1.96	0.00	2.43	0.00	2.84
		L	0.00	3.15	0.00	3.54	0.00	4.39	0.00	9.69

SF = sign flip rate; BR = mean $|\hat{\tau}/\tau|$ (1.00 = unbiased). $\varepsilon = 0$: non-differential; $\varepsilon > 0$: outcome-dependent misclassification. Q: H/M/L = CM diagonal 0.8/0.6/0.4 with uniform off-diagonal. Each cell: $N = 10,000 \times 100$ MC replications.

mean relative bias 0.50 vs. 0.40). All five sign-preserving topologies maintain 0% sign-flip under both structured models at all ε tested; IV likewise shows 0% sign-flip (its vulnerability requires pathological diagonal structure, not differential error). These results indicate that the uniform perturbation model is neither uniformly conservative nor uniformly liberal: practitioners should assess which violation pattern is most plausible for their annotation pipeline.

B.7 Robustness to Nonlinear DGP

We replace the linear structural equations with $Y \sim \text{Bernoulli}(\sigma(\mathbf{x}^\top \beta))$ and estimate via logistic regression (control function for IV) for all seven topologies, using $K=3$, six real VAST confusion matrices, 500 MC replications, and $N=10,000$. Mean sign-flip rates across the six CMs:

Topology	Linear SF_τ	Logistic SF_τ	Linear SF_{any}	Logistic SF_{any}
Confounding	0.000	0.000	0.573	0.559
Mediation	0.000	0.000	0.703	0.695
Collider	0.000	0.000	0.060	0.060
Exposure	—	—	0.577	0.574
M-bias	0.000	0.000	0.095	0.099
Front-door	0.000	0.000	0.590	0.579
IV	0.086	0.074	0.086	0.560

Treatment-effect sign-flip rates (SF_τ) are zero for all five guaranteed topologies under both linear and logistic DGPs, and differ by < 0.02 for IV. The SF_{any} discrepancy for IV reflects estimator dimensionality: the control function includes a residual coefficient prone to sign noise; the core comparison is SF_τ , where both estimators agree. Linear and logistic DGPs produce virtually identical sign-flip patterns across all seven topologies ($|\Delta \text{SF}_\tau| < 0.02$). Under logistic DGP, bias is measured in log-odds units; directional conclusions (sign preservation, topology ranking) transfer directly, while magnitudes are not directly comparable to linear plim expres-

sions. For practitioners requiring finite-sample inference under logistic DGP, bias-corrected bootstrap provides a nonparametric alternative that bypasses the linear plim framework.

B.8 Bootstrap Uncertainty of ASV

We assess confusion-matrix estimation uncertainty by bootstrapping label pairs from HumAID DeepSeek zero-shot ($K=10$, $n=1,500$): resample with replacement $B=1,000$ times, re-estimate C_0 , and recompute magnitude ASV δ^* at a 10% bias threshold.

Topology	δ^* (point)	95% CI	Valid / 1000
Confounding	—	—	0
Exposure	0.937	[0.602, 1.363]	907
IV	0.754	[0.435, 1.290]	902

Confounding never reaches the 10% threshold across all 1,000 bootstraps, confirming robustness. For IV, the 95% CI spans a $3\times$ range (0.435 to 1.290); practitioners should bootstrap C_0 rather than rely on a point estimate when ASV is near the decision boundary. CI width scales as $O(1/\sqrt{n})$ where n is the gold-label validation set size; at $n=1,500$, IV CIs are $\sim 3\times$ wider than exposure CIs, reflecting the Wald estimator’s denominator instability. A convergence study ($K=3$, $B=1,000$) confirms this scaling: exposure CI width narrows from 0.60 ($n=500$) to 0.05 ($n=50,000$), and IV from 0.52 to 0.04, with the product width $\times \sqrt{n}$ approximately constant (CV < 0.20).




Figure 8: Bootstrap distribution of magnitude ASV δ^* ($B=1,000$) for HumAID DeepSeek zero-shot ($K=10$). Dashed line: point estimate; shaded: 95% CI.

B.9 Cross- K Robustness

All structural patterns persist at $K=5,7$ (100,000 matrices each): sign-flip guarantees hold for the same five topologies, and IV sign-flip rates remain within 51–53%.

B.10 ASV Diagnostic Results

For HumAID ($K=10$, 56 scenarios), 55.4% are flagged as magnitude-ASV-unsafe—IV, medi-

Table 10: ASV stability across DGP configurations. Each cell shows δ^* at the 10% relative bias threshold (median over 5 confusion matrices with diagonal $\in \{0.5, 0.6, 0.7, 0.8, 0.9\}$). Lower δ^* indicates greater fragility. Confounding is consistently most robust; for alternating-sign β , the ranking Exposure < IV < Confounding holds. “ ∞ ” = threshold never reached within δ range.

DGP Configuration	Magnitude ASV $\delta^*_{10\%}$		
	Confounding	Exposure	IV
Alt. large: $\beta = [0, 1.0, -0.5]$	0.348	0.090	0.203
Same-sign: $\beta = [0, 0.5, 0.3]$	∞	0.223	0.203
Alt. small: $\beta = [0, 0.2, -0.1]$	∞	0.090	0.203

ation, front-door, and exposure all trigger, while confounding, collider, and M-bias remain below 10% bias (Figure 7). For CivilComments ($K=2$), 100% of scenarios are ASV-unsafe despite F_1 ranging from 0.69 (7B) to 0.82 (frontier).

Overlaying all 23 empirical matrices onto the theoretical Dirichlet violation-rate curves (Figure 5) confirms placement: VAST in the danger zone, HumAID mid-range, CivilComments approaching safe for exposure but still in IV amplification.

B.11 ASV Stability Across DGP Parameters

We test whether ASV rankings are stable across different data-generating process (DGP) configurations by varying the treatment-effect vector β across three settings: alternating-large ($\beta = [0, 1.0, -0.5]$), same-sign ($\beta = [0, 0.5, 0.3]$), and alternating-small ($\beta = [0, 0.2, -0.1]$), with $K=3$, $\tau=0.5$, and five C_0 diagonal values $\in \{0.5, 0.6, 0.7, 0.8, 0.9\}$.

Three patterns emerge (Table 10). First, IV magnitude ASV is identical across all three DGPs ($\delta^*=0.203$). This is a mathematical identity, not an empirical finding: the Wald estimator $\hat{\tau}_{IV} = RF/FS$ has β in both numerator and denominator, and the ratio depends only on C and the first-stage structure, making ASV invariant to β by construction. Second, exposure ASV depends on the coefficient ratio structure: alternating-large ($\beta_2/\beta_1 = -0.5$) and alternating-small ($\beta_2/\beta_1 = -0.5$) share the same ratio and yield identical $\delta^*=0.090$, while same-sign ($\beta_2/\beta_1 = 0.6$) gives $\delta^*=0.223$ —thus the experiment pro-

vides two distinct sign-structure conditions, not three. Third, confounding ASV depends on effect magnitudes: only alternating-large triggers the 10% threshold ($\delta^*=0.348$); smaller effects never reach 10% bias.

Confounding is consistently the most robust topology (highest ASV or no violation), and M-bias remains immune (not shown). However, the relative fragility of exposure vs. IV inverts with coefficient sign structure: under alternating signs, exposure is most fragile ($0.090 < 0.203$); under same-sign effects, IV is most fragile ($0.203 < 0.223$). Practitioners should compute ASV with their specific DGP rather than relying on generic rankings.

B.12 Extended Conjecture Verification

We test envelope monotonicity (Conjecture 1) by sampling 10,000 random C_0 matrices from Dirichlet(1, ..., 1) per topology for $K \in \{3, 5, 7\}$ (210,000 total), sweeping δ in steps of 0.01, and counting matrices with at least one non-monotonic step exceeding tolerance 10^{-12} . Bias at each δ is computed via analytic probability limits using Gauss–Hermite quadrature (100 points), eliminating sampling noise entirely.

Table 11: Envelope monotonicity violations per 10,000 random C_0 matrices (δ step = 0.01). “Worst mag.” is the largest single-step deviation $|B(\delta_{i+1})| - |B(\delta_i)|$ in the non-monotonic direction. *100,000 matrices at step = 0.005 yield exposure $\varepsilon = 0.075$.

Topology	$K=3$	$K=5$	$K=7$	Worst mag.
Confounding	8467	8063	7892	0.011
Mediation	8236	7905	7893	0.013
Collider	8683	7669	7570	0.005
Exposure	485	205	77	0.056*
M-bias	9791	7783	7399	0.002
IV	7433	6015	5711	$> 10^6$
Front-door	8429	8079	7899	0.013

Violation counts are high (60–97% of matrices), but magnitudes are small for non-IV topologies under the default parameterization (Table 11); DGP sensitivity analysis (Table 13) shows ε scales with effect magnitude $|\beta|$, remaining below 0.02 for $|\beta| \leq 1$ across non-exposure topologies. These results are computed from analytic probability limits with zero sampling noise, confirming the

violations are structural properties of multi-class confusion matrices rather than numerical artifacts. IV deviations reflect Wald estimator singularities when the first-stage coefficient passes near zero—this is a property of the Wald estimator, not of the ASV parameterization. Adversarial search (500 restarts per configuration, $K \in \{3, 5\}$) finds genuine counterexamples for all topologies except IV $K=3$, confirming that strict monotonicity does not hold universally. The bias envelope is ε -approximately monotonic with $\varepsilon < 0.03$ for collider and M-bias across all tested τ ; confounding, mediation, and front-door achieve $\varepsilon < 0.06$ at default DGP parameters (see DGP sensitivity below); exposure satisfies $\varepsilon_{p99} < 0.008$ but extreme-tail matrices ($< 0.02\%$ of draws) reach $\varepsilon \approx 0.15$. Empirically, ε scales as $O(1/K^\alpha)$ with $\alpha \in [0.6, 1.5]$ (topology-dependent): for confounding, mediation, and front-door, $\alpha \approx 1.4$; for exposure, $\alpha \approx 1.0$; for collider and M-bias, $\alpha \approx 0.7$. A universal empirical upper bound is $\varepsilon \leq 0.137/\sqrt{K}$, verified across all tested configurations ($K \in \{3, 5, 7\}$). This scaling supports the practical validity of ASV as a sensitivity tool at any meaningful bias threshold.

DGP parameter sensitivity. We repeat the verification across $\tau \in \{0.1, 0.5, 1.0, 2.0, 5.0\}$, varying the direct $T \rightarrow Y$ coefficient for confounding ($\beta_{T=\tau}$), mediation, collider, and M-bias (DGP τ), and rescaling the $A \rightarrow Y$ coefficient vector to $\beta \cdot \tau$ for exposure and front-door. With 10,000 random C_0 per (K, τ) setting (Table 12), ε remains small (< 0.03) for $\tau \geq 0.5$ for confounding, mediation, collider, and M-bias, but grows at $\tau=0.1$: mediation produces 1,950 / 30,000 violations across $K \in \{3, 5, 7\}$ (worst $\varepsilon=0.130$ at $K=3$) and confounding 200 / 30,000 (worst $\varepsilon=0.118$ at $K=3$). The mechanism is dimensional: additive terms not scaling with τ (e.g. the β -contribution to $\mathbb{E}[YD_k]$ in mediation and confounding) contribute a τ -independent offset to the bias numerator, so the ratio $|B|/|\tau|$ inflates as $\tau \rightarrow 0$. A single mediation violation at $\tau=2.0$ ($\varepsilon=0.112$, 1/30,000 matrices) arises from a specific confusion matrix structure where off-diagonal entries induce partial cancellation along the mediation pathway;

this isolated case does not affect the p_{99} bound ($\varepsilon_{p99}=0.005$). Collider and M-bias remain robust at all τ ($\varepsilon<0.022$ and $\varepsilon<0.015$ respectively). Exposure shows scattered tail outliers but $\bar{\varepsilon} \approx 3 \times 10^{-4}$ at every τ , consistent with linear scaling: when β rescales by τ , both bias and τ_{true} scale by the same factor and ε is theoretically invariant. Front-door inverts the pattern (ε grows with τ , reaching 0.139 at $\tau=5$, $K=3$, with 664 violations) because $\tau_{\text{true}} = \text{plim}(\hat{\tau}_T \mid C=I)$ is pinned at 0.5 by the U -confounding path while bias magnitudes scale with β —this is an artefact of how τ_{true} is normalised, not a genuine failure of envelope monotonicity. In practice, ASV’s δ^* thresholds (typically ≥ 0.1 in our deployment) are well above the region where these ε -violations occur, so the small- τ degradation does not affect ASV’s operational reliability; we flag it as a caveat for applications targeting $|\tau| \lesssim 0.1$ in standardised units.

Table 12: $\varepsilon \equiv \max_{\delta} (|B(\delta+\Delta)| - |B(\delta)|) / |\tau_{\text{true}}|$ in the non-monotonic direction (max over $K \in \{3, 5, 7\}$; 10,000 random C_0 per cell). Bold marks $\varepsilon > 0.02$ (non-exposure) or $\varepsilon > 0.08$ (exposure). Counts in parentheses are violations / 30,000. Tables 11 and 12 use independent random seeds; for topologies with low violation rates (exposure $< 0.02\%$), the worst-case ε is sample-sensitive.

Topology	$\tau=0.1$	$\tau=0.5$	$\tau=1.0$
Confounding	0.118 (200)	0.027 (0)	0.014 (0)
Mediation	0.130 (1950)	0.026 (0)	0.013 (0)
Collider	0.022 (0)	0.006 (0)	0.005 (0)
Exposure	0.111 (5)	0.122 (2)	0.152 (2)
M-bias	0.014 (0)	0.003 (0)	0.001 (0)
Front-door	0.003 (0)	0.013 (0)	0.027 (0)

Effect-magnitude sensitivity. We further vary the direct effect $\beta \in \{0.1, 0.5, 1.0, 2.0, 5.0\}$ and noise variance $\sigma^2 \in \{0.5, 1.0, 2.0\}$ across all seven topologies ($K=3$, 10,000 Dirichlet(1, ..., 1) matrices per cell). Table 13 reports the worst-case relative envelope deviation ε_{rel} . For non-exposure topologies, ε scales approximately linearly with $|\beta|$; the safe/fragile partition is stable across all 105 cells.

B.13 Proportional Effects Breakdown

Theorem 1 requires proportional effects ($f = \lambda g$); here we quantify how $\lambda(C)$ out-of-bounds (OOB) rates depend on the degree of de-

β	σ^2	CONFO	MEDIA	COLLI	EXPOS	MBIAS	IV	FRONT
0.1	0.5	0.0002	0.0002	$5.7e-5$	0.1146	$7.6e-7$	0	0.0024
0.1	1	$9.3e-5$	0.0002	$6.0e-5$	0.1027	$4.6e-7$	0	0.0030
0.1	2	$4.4e-5$	0.0002	$5.8e-5$	0.0724	$1.9e-7$	0	0.0044
0.5	0.5	0.0040	0.0057	0.0014	0.0481	0.0003	0	0.0104
0.5	1	0.0019	0.0060	0.0014	0.0953	0.0002	0	0.0151
0.5	2	0.0010	0.0057	0.0014	0.0991	$9.8e-5$	0	0.0214
1	0.5	0.0083	0.0272	0.0048	0.1886	0.0021	0	0.0202
1	1	0.0066	0.0262	0.0048	0.1100	0.0015	0	0.0304
1	2	0.0039	0.0257	0.0048	0.2153	0.0009	0	0.0388
2	0.5	0.0148	0.1326	0.0137	0.0551	0.0074	0	0.0315
2	1	0.0091	0.0938	0.0117	0.0993	0.0063	0	0.0435
2	2	0.0125	0.0989	0.0109	0.0998	0.0045	0	0.0616
5	0.5	0.0083	0.8983	0.0305	0.1728	0.0155	0	0.0459
5	1	0.0230	0.8870	0.0267	0.0575	0.0156	0	0.0653
5	2	0.0122	0.9082	0.0206	0.0870	0.0132	0	0.0945

Table 13: Conjecture 1 DGP sensitivity: worst-case $\varepsilon_{\text{rel}} = \max_{C_0} (\text{wrong-direction step}) / |\tau_{\text{true}}|$ across 5 effect magnitudes β , 3 noise variances σ^2 , and 7 topologies ($K=3$, 10,000 Dirichlet(1, ..., 1) matrices per cell). ε scales approximately linearly with $|\beta|$ for non-exposure topologies; the safe/fragile partition is stable across all cells.

parture from proportionality. We simulate $K=3$ mediation with 100,000 random confusion matrices (diagonal ≥ 0.5 , Dirichlet off-diagonal) at five proportionality degrees: exact ($f_j - f_0 = \lambda(g_j - g_0)$), near-proportional ($f = \lambda g + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$, $\sigma \in \{0.1, 0.3, 0.5\}$), and fully independent ($f \perp g$). Table 14 reports the $\lambda(C) \notin [0, 1]$ rate.

Table 14: $\lambda(C)$ out-of-bounds rate and mediation sign failure rate by degree of departure from proportionality effects ($K=3$, 100K random confusion matrices per cell).

Proportionality degree	$\lambda(C) \notin [0, 1]$	Sign failure
Exact ($\sigma=0$)	0.0%	0.85%
Near ($\sigma=0.1$)	10.3%	0.87%
Near ($\sigma=0.3$)	19.0%	0.96%
Near ($\sigma=0.5$)	22.4%	0.93%
Independent ($f \perp g$)	22.9%	0.17%
NLP-realistic (plan_006 DGP)	0.07%	0.0%

The λ OOB rate rises steeply from 0% to $\sim 10\%$ at $\sigma=0.1$ and saturates near 23% for independent (f, g). Sign failure rates remain below 1% across all conditions because λ excursions outside $[0, 1]$ are typically small and do not reverse the dominant effect direction. The NLP-realistic DGP ($\delta=[0, 0.8, -0.6]$, $\beta=[0, 1.0, -0.5]$) exhibits near-zero OOB (0.07%) and zero sign

failures, confirming that proportional effects hold approximately in practice.

B.14 FWL Invariance Under Nonlinear Models

FWL invariance guarantees exactly 0% bias reduction from matrix-based corrections (MLA, naive plug-in) under linear OLS for covariate topologies (confounding, mediation, collider, M-bias). Here we test whether this invariance persists under logistic regression across sample sizes $N \in \{1000, 5000, 10,000, 50,000\}$ and three error levels ($C_{kk} \in \{0.9, 0.7, 0.5\}$), with 200 Monte Carlo replicates per cell. Table 15 reports the mean logistic bias reduction (positive = correction helps).

Table 15: Mean logistic bias reduction (%) from matrix-based correction, by topology and error level (averaged over $N \in \{1K, 5K, 10K, 50K\}$; 200 MC reps each). Under linear OLS, all covariate topologies yield exactly 0%. Exposure is shown for comparison.

Topology	Low ($C_{kk}=0.9$)	Med ($C_{kk}=0.7$)	High ($C_{kk}=0.5$)
Confounding	-0.003%	<0.001%	<0.001%
Mediation	-0.004%	<0.001%	<0.001%
Collider	-0.033%	-0.020%	<0.001%
M-bias	-0.049%	-0.010%	<0.001%
Front-door	0.001%	-0.001%	<0.001%
Exposure	-68.5%	-48.7%	-22.5%

For all five covariate topologies, logistic bias reduction is negligible ($|\text{mean}| < 0.05\%$) across all N and error levels, with no systematic N -scaling: medians cluster within $\pm 0.002\%$ at $N=50,000$. The large standard deviations observed at certain (N, error) cells for collider and M-bias (up to 2.6% std at $N=10,000$) reflect rare outlier replicates; medians remain near zero. This confirms that FWL invariance is a robust structural property that extends beyond linear models for covariate topologies.

In contrast, exposure shows substantial FWL breakdown under logistic regression: -22% to -78% bias reduction (negative = correction worsens bias), stabilising at $N \geq 5,000$. The breakdown magnitude increases with accuracy (C_{kk}), consistent with the exposure topology’s sensitivity to the nonlinear link function transforming category weights.

B.15 PPI++ Gold Label Ratio Sensitivity

We test PPI++ correction performance as a function of gold label fraction ($K=3$, $N=10,000$, 100 MC replications, 5 Dirichlet confusion matrices with $\min_k C_{kk} \geq 0.6$).

Table 16: PPI++ RMSE reduction (%) vs. gold label fraction across five topologies. Negative values indicate PPI++ increases RMSE relative to the naive estimator.

Topology	5%	10%	20%	50%	100%
Confounding	+84.8	+88.6	+91.8	+94.2	+95.2
Exposure	+75.2	+83.1	+87.1	+91.2	+93.8
Mediation	+56.7	+70.4	+77.2	+85.3	+88.7
IV	+1.4	+22.1	+50.1	+66.7	+76.4
M-bias	-125.3	-43.6	-19.6	+6.0	+18.5

For confounding and exposure, PPI++ achieves >75% RMSE reduction even at 5% gold labels; IV requires $\geq 20\%$ gold to reach 50% reduction, reflecting the Wald estimator’s higher variance. M-bias exhibits harmful correction at $\leq 20\%$ gold (-125% at 5%), because PPI++ re-exposes the M-bias pathway that annotation error had attenuated; only at $\geq 50\%$ gold does the estimator become weakly beneficial (+6%).

B.16 Realistic Differential Misclassification Thresholds

Our non-differential assumption (Definition 1) requires $P(A^*|A, Y) = P(A^*|A)$. We estimate how much Y -stratified differential error each of our 23 empirical LLM confusion matrices can tolerate before exposure sign reversal exceeds 5%. For each matrix C , we construct $C_{\text{high}} = \text{renorm}(C + \varepsilon\Delta)$ and $C_{\text{low}} = \text{renorm}(C - \varepsilon\Delta)$ with random off-diagonal perturbation Δ (500 MC trials per ε), sweep $\varepsilon \in [0, 0.50]$, and report the safe threshold $\varepsilon_{\text{safe}}^* = \max\{\varepsilon : \text{reversal rate} < 5\%\}$ (Table 17).

The distribution is sharply bimodal: all $K=2$ configurations are immune to differential misclassification within the tested range ($\varepsilon_{\text{safe}}^*=0.50$), consistent with the binary sign-preservation guarantee (Proposition 1(c)). For $K \geq 3$, 14 of 16 configurations have $\varepsilon_{\text{safe}}^*=0$ —sign reversal occurs at the smallest tested $\varepsilon=0.01$. Two outliers (VAST: $\varepsilon^*=0.02$; HuMAID: $\varepsilon^*=0.36$) have atypically structured

Table 17: Safe differential misclassification thresholds $\varepsilon_{\text{safe}}^*$ by dataset/ K . For $K=2$ (binary), no configuration produces sign reversal up to $\varepsilon=0.50$. For $K \geq 3$, the majority of configurations are fragile ($\varepsilon_{\text{safe}}^*=0$), confirming that the non-differential assumption is load-bearing for multi-class exposure.

Dataset	K	n	Mean ε^*	Median ε^*
Civil Comments	2	7	0.50	0.50
VAST	3	8	0.003	0.00
HuMAID	10	8	0.045	0.00

confusion matrices. This confirms that even small violations of non-differential misclassification are consequential for multi-class exposure, underscoring the practical importance of validating $P(A^*|A, Y) \approx P(A^*|A)$ on gold-labeled subsets when $K \geq 3$.

B.17 Proportional Effects with Empirical Confusion Matrices

Theorem 1 assumes proportional effects ($f = \lambda g$). We evaluate this assumption using all 23 empirical LLM confusion matrices under four DGP groups in the mediation topology (23 matrices $\times$ 50 seeds $\times$ 100,000 precomputed samples per seed).

Table 18: Proportional effects validation with 23 real LLM confusion matrices under mediation DGP. $\lambda(C) \in [0, 1]$: proportion satisfying the Theorem 1 bound. Sign pres.: sign preservation rate of the direct effect estimate.

DGP Group	$\lambda(C)$ mean	$\lambda(C) \in [0, 1]$	Sign pres.
Proportional ($f=\lambda g$)	0.271	99.8%	100%
Near-prop. ($\sigma=0.1$)	0.285	99.2%	100%
Non-proportional	0.240	85.0%	97.3%
Sparse	0.296	75.8%	98.6%

Under proportional and near-proportional effects, $\lambda(C)$ remains within $[0, 1]$ for $\geq 99\%$ of configurations with 100% sign preservation—the Theorem 1 conditions are satisfied by real LLM confusion matrices. Even under non-proportional effects (f independent of g), sign preservation remains at 97.3%, with $\lambda(C)$ exceeding $[0, 1]$ in 15% of cases. The per- K breakdown shows that $\lambda(C)$ is dimension-dependent: $K=2$ (mean 0.249), $K=3$ (mean 0.082), $K=10$ (mean 0.490), with 100% in $[0, 1]$ across all K under proportional effects. Sparse DGP (zero-effect categories) produces the highest violation rate (24.2%), consistent with the breakdown mechanism: zero-effect

Topology	Naive Bias (%)	ASV δ_{mag}^*	Routed Method
Confounding (naive plug-in) (MLA)	$+22.71 \pm 2.29$	0.352	no correction (FWL)
Exposure IV	-42.87 ± 2.42	0.239	PPI++
	$+82.40 \pm 12.22$	0.204	MC-SIMEX (gold-l)

Table 19: End-to-end demo over 100 replicates ($N=10,000$, $K=3$, 20% gold). Confounding: plug-in corrections leave the treatment coefficient unchanged (FWL invariance), so Algorithm 1 routes to “no correction”. Exposure: PPI++ recovers the target coefficient. IV: Wald estimator instability produces large bias; PPI++ (OLS-based) is limited, routing to MC-SIMEX.

categories make f/g ratios undefined, maximally departing from proportionality.

B.18 Case Study: ASV Diagnosis of Political Ad Tone Classification

We illustrate the ASV diagnostic workflow on a concrete CSS application. Egami et al. (2023) use LLM classifiers to annotate political advertising tone ($K=3$: promote/contrast/attack) from the Fowler & Ridout database, then estimate the causal effect of ad tone on voter outcomes. The annotated variable is the **exposure** (treatment). Using a representative $K=3$ confusion matrix from our empirical set (VAST DeepSeek zero-shot, mean diagonal 0.734) as a proxy:

Coefficient	δ_{mag}^* (10%)	δ_{sf}^*	Relative bias at $\delta=1$
β_1 (primary)	0.234	> 1.5	39.7%
β_2 (secondary)	0.154	> 1.5	60.4%

Both coefficients are magnitude-unsafe ($\delta^* < 1$), with the secondary coefficient more fragile ($\delta^*=0.154$). If the same annotator were instead classifying a **confounder** (e.g., media framing as a control variable), magnitude ASV rises to $\delta^*=0.690$ —a $3\times$ increase in error tolerance from the same confusion matrix. This demonstrates the central insight: topology, not annotator quality, determines vulnerability. Algorithm 1 routes this exposure topology to PPI++ or DSL correction; matrix-based methods (naive plug-in, MLA) are ineffective due to FWL non-applicability.

B.19 End-to-End Demonstration

The demonstration uses a symmetric C_0 ($d_{\text{rel}}=0$); empirical $K=3$ matrices deviate substantially (median $d_{\text{rel}}=0.33$, Remark 4), but

LLM Config	d_{rel}	Exposure		stronger: Confounding, collider, M-bias, IV, and exposure at $K=10$)
		Naive (%)	PPI++ (%)	Naive (%) Plug-in (%)
DeepSeek FS-3	0.30	-61.3	+0.1	+28.5
DeepSeek FS-5	0.30	-61.5	-0.5	+28.7
DeepSeek ZS	0.31	-61.1	-0.5	+28.7
Llama4 FS-3	0.37	-64.3	+0.0	+28.4
Llama4 FS-5	0.32	-61.3	+0.0	+26.8
Llama4 ZS	0.36	-59.3	-0.2	+27.9
Qwen2.5-7B FS-3	0.38	-68.9	-0.5	+29.3
Qwen2.5-7B FS-5	0.34	-68.7	+0.3	+29.1
Qwen2.5-7B ZS	0.38	-65.5	+0.3	+28.7
Qwen3 ZS	0.32	-60.3	-0.5	+28.2
Median improvement		+93.6%		-0.0%

Table 20: Semi-real case study with VAST $K=3$ confusion matrices from 10 LLM configurations. Bias is relative to true τ . Exposure: PPI++ substantially reduces bias. Confounding: plug-in correction is ineffective (FWL invariance).

core ASV ranking and correction routing remain valid under Conjecture 1.

B.20 Semi-Real Case Study with Production LLM Confusion Matrices

We repeat the end-to-end pipeline using 10 confusion matrices estimated from production LLM annotation of VAST stance detection ($K=3$; DeepSeek, Llama4, Qwen2.5-7B, Qwen3 under zero-shot and few-shot prompting). These matrices exhibit median $d_{\text{rel}}=0.33$ (Remark 4), far from the symmetric assumption of Theorem 2. Results confirm both topology-dependent patterns: PPI++ achieves 93.6% median exposure bias reduction while plug-in corrections yield 0% improvement for confounding (FWL invariance, machine precision; Table 20).

B.21 Correction Method Sensitivity to K

We test whether the correction method ranking from Table 5 ($K=3$) is stable at $K=5$ and $K=10$. For each K , we sample 5 random Dirichlet confusion matrices (diagonal ≥ 0.6), run 100 MC replications per topology–matrix combination ($N=5,000$, 10% gold fraction), and compute RMSE reduction for all five methods across all seven topologies.

Grand ranking stability. Spearman rank correlations across all topology–method pairs: $\rho_{3,5} = 0.80$, $\rho_{3,10} = 0.84$, $\rho_{5,10} = 0.93$ (all $p < 10^{-8}$). Per-topology correlations are even

stronger: Confounding, collider, M-bias, IV, and exposure at $K=10$) achieve $\rho = 1.0$ between $K=3$ and $K=10$. **Method-specific patterns.** PPI++ effectiveness degrades with K : confounding RMSE reduction drops from 85.6% ($K=3$) to 75.6% ($K=10$); mediation drops from 57.1% to 9.0%. At $K=10$, PPI++ becomes harmful for IV (−124%) and exposure (−11%). MC-SIMEX is the most K -robust method, maintaining the top rank for 5/7 topologies at $K=10$ (vs. 3/7 at $K=3$). FWL invariance (naive plug-in and MLA achieving 0% bias reduction on the treatment coefficient) holds at all K tested, confirming the structural nature of this result.

Table 21: Top correction method and RMSE reduction (%) by topology across K . Bold = top method for that topology– K combination.

Topology		$K=3$	$K=5$	$K=10$
Confounding	PPI++	85.6	PPI++ 76.4	PPI++ 75.6
Mediation	PPI++	57.1	SIMEX 59.7	SIMEX 59.0
Collider	SIMEX	55.4	SIMEX 63.4	SIMEX 57.3
Exposure	SIMEX	72.5	SIMEX 65.0	SIMEX 47.3
M-bias	SIMEX	3.2	SIMEX 47.0	SIMEX 43.7
IV	SIMEX	30.9	SIMEX 30.4	SIMEX 27.1
Front-door	PPI++	82.9	PPI++ 69.0	SIMEX 70.1

B.22 Comparison with Related Sensitivity Frameworks

Table 22: Comparison of sensitivity frameworks.

	ASV	E-value
Parameter space	$K \times K$ confusion $C(\delta)$	Scalar confound
Sensitivity path	$(1-\delta)I + \delta C_0$	Monotone sca
Topology-specific	✓	✗
Sign-flip diagnostic	✓ (δ_{sign}^*)	✗
Topology dependence	✓ (7 topologies ranked)	✗
Monotonicity	Proved $K=2$; conj. $K \geq 3$	Guaranteed
Target	Annotation error	Unmeasured confo

B.23 Per-Topology F_1 -Bias Correlation

The cross-topology correlation between macro F_1 and $|\text{bias}|$ is weak ($\rho = -0.10$, $p = 0.45$; §4.4), but decomposing by topology reveals a structured pattern. Table 23 reports Spearman correlations computed over 23 LLM configurations (HumAID $K=10$, VAST $K=3$, CivilComments $K=2$; 161 total scenarios).

Five topologies show the expected negative correlation: better annotation quality ($F_1 \uparrow$)

Table 23: Within-topology Spearman $\rho(F_1, |\text{bias}|)$ across 23 empirical LLM configurations $\times$ 3 datasets. Negative ρ : higher $F_1 \rightarrow$ lower bias (expected). Positive ρ : higher $F_1 \rightarrow$ higher bias (protective attenuation removed).

Topology	n	ρ	p
Confounding	23	-0.557	0.006
Mediation	23	-0.565	0.005
Exposure	23	-0.531	0.009
Front-door	23	-0.568	0.005
IV	23	-0.700	<0.001
Collider	23	+0.310	0.150
M-bias	23	+0.867	<0.001

reduces causal bias. IV shows the strongest effect ($\rho = -0.70$), consistent with the Wald estimator’s sensitivity to misclassification in the first stage. M-bias uniquely reverses: $\rho = +0.87$ ($p < 0.001$). The mechanism is that misclassification attenuates the spurious M-bias path, so the naive estimator under noisy annotations is *closer* to the true causal effect; improving annotation quality (higher F_1) removes this protective attenuation and *increases* bias. This is consistent with the correction results (§4.5), where PPI++ harms M-bias (-88%) by removing the same protective noise. Collider shows a positive but non-significant trend ($\rho = +0.31$, $p = 0.15$), consistent with Berkson bias attenuation. Per-dataset breakdowns (CivilComments $K=2$: $\rho = -1.0$ for all topologies with $n=7$) should be interpreted with caution given the small sample and limited F_1 variation within the binary task.

B.24 Confusion Matrix Estimation Robustness

The ASV framework requires the confusion matrix C , which in practice must be estimated from a held-out gold-labeled subset. We assess how estimation noise in $\hat{C}$ propagates to the safe/fragile partition and ASV ranking. For each of 8 empirical confusion matrices (VAST, 4 LLMs $\times$ 2 prompting strategies), we draw $n_{\text{gold}} \in \{50, 100, 200, 500\}$ samples from a Multinomial(n_{gold}, C_k) distribution per column, compute $\hat{C}$, and evaluate the resulting ASV ranking across 200 bootstrap replicates.

The ASV ranking is highly robust: median Spearman $\rho \geq 0.99$ even at $n_{\text{gold}}=50$. The safe/fragile partition (5 safe vs. 2 fragile topologies) recovers reliably from $n_{\text{gold}} \geq$

Table 24: $\hat{C}$ estimation robustness: safe/fragile partition stability and ASV ranking correlation (Spearman ρ vs. true- C ranking) under multinomial bootstrap from finite gold sets. 8 empirical matrices $\times$ 200 bootstrap replicates per n_{gold} .

n_{gold}	Partition stability	ρ (median)	ρ (q_{25})
50	50.4%	0.991	0.789
100	82.0%	1.000	0.991
200	85.6%	1.000	0.991
500	88.5%	1.000	0.991

100 (82% stability), increasing to 89% at $n_{\text{gold}}=500$. The remaining instability concentrates in source matrices with low diagonal dominance (e.g., VAST DeepSeek few-shot: $\min C_{kk} = 0.16$), where bootstrap noise can shift borderline ASV values across the partition boundary. All 8 truth partitions match the canonical safe/fragile partition, confirming that the theoretical dichotomy holds empirically.

B.25 Sign Reversal Determinants

We investigate which confusion matrix features predict exposure sign reversal across the 23 empirical LLM configurations. For each matrix, we compute 10 structural features: $\min_k C_{kk}$ (minimum diagonal), $\overline{C}_{kk}$ (mean diagonal), $\max_{j \neq k} C_{jk}$ (maximum off-diagonal), $\max_{j \neq k} C_{jk}/C_{kk}$ (off-diagonal ratio), $\log(\kappa(C))$ (log condition number), mean and max column entropy, spectral gap, $\det(C)$, and minimum singular value. Sign reversal is defined as any exposure coefficient reversing sign under misclassification.

Table 25: Sign reversal rates by K and best single-feature decision rule. Rule: $\det(C) < 0.014 \rightarrow$ sign reversal (accuracy 91.3%, sensitivity 85.7%, specificity 100%). An AND rule ($\min_k C_{kk} < 0.55$ and mean entropy > 0.40) achieves 95.7% accuracy with 100% sensitivity.

K	n	Sign reversal rate
2	7	0/7 (0%)
3	8	7/8 (87.5%)
10	8	7/8 (87.5%)

Binary matrices ($K=2$) never produce exposure sign reversal, consistent with classical attenuation theory. For $K \geq 3$, the determinant of C is the strongest single predictor: $\det(C) < 0.014$ identifies sign-reversing matrices with 100% specificity (zero false posi-

tives) and 85.7% sensitivity. An AND rule combining $\min_k C_{kk} < 0.55$ with mean column entropy > 0.40 achieves 100% sensitivity at 88.9% specificity. These rules are robust to β -scaling (accuracy unchanged at $0.3\times$ and $3\times$ effect magnitude) but sensitive to β -sign structure: under uniform-positive β (all coefficients same sign), no matrices produce sign reversal regardless of C , as cancellation across differently-signed coefficients is the underlying mechanism.

B.26 Y-Stratified Non-Differential Verification

We directly test the non-differential assumption (Definition 1) by stratifying confusion matrices by outcome-proxy variables and testing $C(Y=\text{high}) \approx C(Y=\text{low})$. Using 12 LLM configurations on VAST ($K=3$; 4 models $\times$ 3 prompting strategies, 1,500 samples each), we compute per-stratum confusion matrices under three Y-proxies and apply permutation tests ($B=10,000$, test statistic: $\max_{j,k} |C_{jk}^{(0)} - C_{jk}^{(1)}|$).

Table 26: Y-stratified non-differential test. Content-independent Y (text length) shows no significant violation (1/12 rejection $\approx$ expected false positive rate at $\alpha=0.05$). Content-dependent proxies (topic, sentiment) show expected violations since LLMs observe text content.

Y proxy	Type	Reject ($\alpha=0.05$)	Median p
Text length	Content-indep.	1/12	0.490
Topic	Content-dep.	11/12	0.002
Sentiment	Content-dep.	10/12	0.017

The single text-length rejection (Llama-4 few-shot-3, $p=0.027$) is consistent with the expected false positive rate ($0.05 \times 12 = 0.6$). Content-dependent proxies (topic, sentiment) show systematic rejection because LLMs process text content from which these variables can be inferred, violating $Y \perp \text{text} \mid A$. In causal inference applications, Y is typically an external outcome (e.g., health status, economic indicator) that cannot be inferred from the annotated text, satisfying the independence condition. The content-dependent rejections confirm statistical power: the test can detect violations when they exist.

Table 27: Key notation used throughout the paper.

Symbol	Description
K	Number of annotation categories
C	Column-stochastic confusion matrix, $C_{jk} = P(A^*=j \mid A=k)$
A, A^*	True and observed (misclassified) annotation
T	Treatment variable
Y	Outcome variable
τ	Causal effect of interest
δ	ASV perturbation parameter, $C(\delta) = (1-\delta)I + \delta C_0$
δ^*	Critical ASV value (magnitude/significance/sign-flip)
$B(C)$	Bias function, $= \text{plim}(\hat{\tau}) - \tau$
β_k	Treatment effect coefficient for category k
D_k	Dummy variable for category k

C Notation

D Use of AI Assistants

We used Claude (Anthropic) as an AI coding and writing assistant throughout this work. Specifically:

- **Code:** drafting experiment scripts, data-processing pipelines, and analysis utilities; debugging and refactoring.
- **Writing:** drafting and revising paper sections, polishing prose, and reformatting $\text{\LaTeX}$.
- **Literature:** summarizing related papers to inform the related-work section.

Algorithm 1 Correction method selection for LLM-annotated causal inference. This topology-based classification is robust to ε -approximate deviations; practitioners computing quantitative ASV thresholds should note the ε -approximate bounds (Conjecture 1, Section 3.3). Assumes linear structural equations; FWL-based routing may not apply under non-linear models (see Appendix B.7). Thresholds calibrated for $K=3$ linear SEM; recalibrate for different K using topology-specific bias functions.

Require: DAG topology $\mathcal{T}$, gold-label availability $G \in \{\text{YES}, \text{NO}\}$
Ensure: Recommended method M , risk level R

```

1:
2: Step 0: DAG uncertainty {If  $\mathcal{T}$  is uncertain, repeat Steps 1–3 under each candidate topology and report the range}
3:
4: Step 1: Assess risk level
5: if  $\mathcal{T} \in \{\text{M-bias}\}$  then
6:    $R \leftarrow \text{Low}$   $\{|\text{bias}| < 0.04 \text{ for all } C\}$ 
7: else if  $\mathcal{T} \in \{\text{Conf.}, \text{Med.}, \text{Coll.}, \text{Front-door}\}$  then
8:    $R \leftarrow \text{MEDIUM}$   $\{\text{bounded bias, no sign reversal}\}$ 
9: else
10:   $R \leftarrow \text{HIGH}$   $\{\text{Exposure, IV: sign reversal 83.4\% / 52.6\%}\}$ 
11: end if
12:
13: Step 2: Select correction method
14: if  $G = \text{YES}$  then
15:   if  $\mathcal{T} = \text{M-bias}$  then
16:     $M \leftarrow \text{NONE}$   $\{\text{PPI++ RMSE red. } -88.0\%; \text{ no correction preferred}\}$ 
17:   else
18:     $M \leftarrow \text{PPI++}$   $\{47.0\% \text{ RMSE red.}, 96.8\% \text{ coverage}\}$ 
19:   end if
20: else
21:    $\{\text{no gold labels}\}$ 
22:   if  $\mathcal{T} \in \{\text{Confounding, Mediation}\}$  then
23:     $M \leftarrow \text{DSL}$   $\{48.0\% / 52.9\% \text{ RMSE reduction}\}$ 
24:   else if  $\mathcal{T} = \text{IV}$  then
25:     $M \leftarrow \text{MC-SIMEX}$   $\{59.4\% \text{ RMSE reduction}\}$ 
26:   else if  $\mathcal{T} \in \{\text{Collider, M-bias}\}$  then
27:     $M \leftarrow \text{MC-SIMEX}$   $\{31.8\% / 0.1\% \text{ RMSE reduction}\}$ 
28:   else
29:     $M \leftarrow \emptyset$ ; flag: acquire gold labels  $\{\text{no reliable gold-free method}\}$ 
30:   end if
31: end if
32:
33: Step 3: FWL invariance warning
34: if  $\mathcal{T} \in \{\text{Conf.}, \text{Med.}, \text{Coll.}, \text{M-bias}\}$  then
35:   Skip Naive plug-in and MLA  $\{0\% \text{ bias red. on } T \text{ coeff. (Table 5: below 2\%, no method improves)}\}$ 
36: end if
37: return  $(M, R)$ 

```
