

Multi-Class Annotation Errors in Causal Inference: Topology-Dependent Bias, Sensitivity, and Correction

Anonymous ACL submission

Abstract

When large language model (LLM)-generated annotations feed into causal inference, annotation quality alone is insufficient to assess causal validity; the causal directed acyclic graph (DAG) topology determines whether errors attenuate, amplify, or reverse the treatment effect estimate. We derive bias expressions for seven canonical topologies under K -class non-differential misclassification in linear structural equation models and show that $K \geq 3$ introduces a structural possibility of sign reversal absent in the binary case. Under proportional effects, five topologies preserve the treatment effect sign under regularity conditions; exposure and instrumental variable (IV) can reverse it. Under empirical LLM confusion matrices, sign reversal is confined to exposure topology. We introduce Annotation Sensitivity Values (ASVs), with proved ($K=2$, $K=3$ symmetric) and extensively validated (general $K \geq 3$) monotonicity that yields a vulnerability ranking from IV (maximally fragile) to M-bias (immune), and show that matrix-based corrections are provably ineffective under linear regression for four topologies due to Frisch–Waugh–Lovell invariance. A case study with 10 real LLM confusion matrices confirms: topology-aware correction reduces exposure bias from $\sim 61\%$ to $< 1\%$, while topology-blind corrections are provably futile under linear regression for four topologies¹.

1 Introduction

Large language models (LLMs) are rapidly replacing human annotators in computational social science (Gilardi et al., 2023; Ziems et al., 2024; Egami et al., 2023; Baumann et al., 2025). Most tasks are inherently multi-class (stance detection: 3–4 categories, sentiment: 5, disaster response: 10+), and annotation errors propagate into downstream causal

estimates in ways that depend on both error structure and causal design. Baumann et al. (2025) show that 31% of hypotheses yield incorrect conclusions with state-of-the-art LLMs, with 60% error correlation across models (Kim et al., 2025). Standard metrics do not capture this risk: within F_1 0.61–0.69, causal bias varies by an order of magnitude across topologies.

The transition at $K \geq 3$ is not gradual but structural. For $K=2$, the stochastic constraint forces any confusion matrix with positive diagonal to preserve coefficient signs, an algebraic property of the 2×2 matrix dimension (Brenner, 1993). This guarantee vanishes at $K \geq 3$: the higher-dimensional confusion matrix admits negative effective weights on categories with opposite-sign coefficients; sign reversal can occur even under high accuracy (Veierød and Laake, 2001). Under a uniform Dirichlet prior, 83.4% of $K=3$ exposure confusion matrices produce sign reversal of at least one secondary coefficient; under realistic accuracy (diagonal ≥ 0.7), 2.1% still do. The primary treatment coefficient is never reversed under the tested β configurations where a single category dominates ($\beta_1 \gg \beta_2$), an empirical regularity we discuss in Section 3.

F_1 does not exhibit significant rank correlation with causal bias ($\rho = -0.10$, $p = 0.45$ across 56 HumAID disaster-response scenarios ($K=10$); bias varies $> 10\times$ (from $< 6\%$ to 42.1%) within F_1 0.614–0.690). Prior work noted polytomous sign reversal (Dosemeci et al., 1990; Veierød and Laake, 2001) and confounder amplification (Brenner, 1993), but covered at most one topology (Lee and Wood-Doughty, 2024). We present the first systematic analysis jointly addressing multi-class ($K \geq 3$) misclassification across seven directed acyclic graph (DAG) topologies, with analytical predictions matching Monte Carlo within 1.6% mean absolute percentage error (MAPE) across 161 scenarios.

We address these gaps with three contributions:

¹Code and data are available at https://anonymous.4open.science/r/causal_annot_error_attrib-C146/.

1. **Multi-class bias taxonomy.** We derive bias expressions for seven canonical DAG topologies² (Figure 1) under K -class non-differential misclassification and establish a sign-preservation/reversal dichotomy: five topologies preserve the treatment effect sign under topology-specific regularity conditions (Proposition 1), while exposure and IV can reverse it.
2. **Annotation Sensitivity Values.** We define three ASV diagnostics (magnitude, significance, and sign-flip) with proved envelope monotonicity (Theorems 1, 2) and ε -approximate monotonicity for general $K \geq 3$ (Conjecture 1). This reveals a vulnerability spectrum from IV (maximally fragile) to M-bias (immune).
3. **Topology-aware correction routing.** Frisch–Waugh–Lovell (FWL) invariance renders matrix-based corrections provably ineffective for four topologies (confounding, mediation, collider, M-bias). We evaluate five correction methods and provide a topology-aware decision procedure (Algorithm 1), validated with real LLM confusion matrices.

2 Related Work

Measurement error in causal inference. Non-differential misclassification of a binary variable biases estimates toward the null (Bross, 1954; Brenner, 1993; Ogburn and VanderWeele, 2012), formalized in the DAG framework by Hernán and Cole (2009) and Kuroki and Pearl (2014). This guarantee breaks for polytomous variables (Dosemeci et al., 1990; Jurek et al., 2005; Veierød and Laake, 2001), yet extensions to differential error, mismeasured covariates, and nonclassical structures (Brennan et al., 2021; Shu and Yi, 2019; Valeri and VanderWeele, 2013; Schennach, 2020) treat topologies in isolation. No prior work compares how the *same* confusion matrix produces qualitatively different bias depending on the mismeasured variable’s causal role.

LLM annotation and downstream causal validity. LLMs are rapidly replacing human annotators (Gilardi et al., 2023; Ziems et al., 2024; Tan

et al., 2024), yet consequences for causal inference remain underexplored. Baumann et al. (2025) and Kim et al. (2025) show that 31% of hypotheses yield incorrect conclusions and error correlations reach 60% across models, but neither connects error patterns to DAG structure. TeBlunthuis et al. (2024) provides the most complete binary analysis (four topologies, $K=2$) and Lee and Wood-Doughty (2024) extends to categorical classifiers, but only for confounding. The multi-class ($K \geq 3$) regime, where sign reversal becomes structurally possible, remains unaddressed.

Sensitivity analysis and correction methods. Sensitivity frameworks for unmeasured confounding (VanderWeele and Ding, 2017; Cinelli and Hazlett, 2020; Oster, 2019) do not address annotation error, while matrix-adjustment (Blackwell et al., 2017; Lash et al., 2009) and prediction-powered methods (Angelopoulos et al., 2023; Egami et al., 2023; Zrnić and Candès, 2024; Gligorić et al., 2025) target misclassification but remain topology-agnostic. We show this agnosticism is costly: FWL invariance renders matrix-based corrections provably ineffective for four topologies, and topology-blind application can amplify bias (−139% for IV).

3 Method

3.1 Setup and Notation

Let A denote a true K -class categorical variable (e.g., sentiment with $K=3$ levels) and A^* its observed, potentially misclassified version produced by an LLM annotator. The misclassification process is characterized by a $K \times K$ confusion matrix C with entries $C_{jk} = P(A^*=j \mid A=k)$, where each column sums to one (column-stochastic: each column gives the distribution of predicted labels for a given true label). Following Hernán and Cole (2009), we represent this measurement process using *starred-node* DAGs: the true variable A is latent, and the observed A^* is connected to A via a dashed measurement edge.

Notation conventions. Our analysis characterizes probability limits ($N \rightarrow \infty$) under the FWL theorem (Frisch and Waugh, 1933), which explains why matrix-based corrections fail for covariate topologies (Section 4.5). Notation is collected in Table 28 (Appendix C).

Definition 1 (Non-differential misclassification). *Misclassification of A is non-differential with re-*

²These span the four structural roles of the annotated variable: treatment (exposure), outcome-side covariate (confounding, mediation, M-bias, front-door), selection variable (collider), and instrumental variable (IV).




Figure 1: Seven canonical DAG topologies with starred-node annotation (Hernán and Cole, 2009). A (latent) $\rightarrow A^*$ (observed) via a dashed measurement edge; the causal role of A varies: treatment (exposure), covariate (confounding, mediation, M-bias, front-door), selection (collider), instrument (IV).

spect to outcome Y and covariates $\mathbf{Z}$ if $P(A^*=j \mid A=k, Y, \mathbf{Z}) = P(A^*=j \mid A=k) = C_{jk}$ for all j, k .

In batch annotation pipelines, the LLM processes each text independently without access to Y , making non-differentiality an architectural invariant (Ziems et al., 2024); Y -stratified permutation tests confirm this (1/12 rejection at $\alpha=0.05$; Appendix B.5, B.26). Non-differentiality may be violated when Y correlates with text content, in which case our Y -stratified test provides empirical verification. Under stress testing, sign reversal occurs in 14/16 $K \geq 3$ exposure configurations at $\eta=0.01$ structured departure (Appendix B.6).

Throughout, τ denotes the causal estimand of interest, whose definition varies by topology: the confounder’s coefficient, the treatment effect, the indirect effect, or the Wald ratio (Appendix A.3ff.).

We study seven canonical DAG topologies (Figure 1) (Pearl, 2009; Hernán and Robins, 2020), deriving probability limits under linear structural equation models (SEMs) (Carroll et al., 2006); logistic regression robustness is verified in Appendix B.7 and B.14.

3.2 Multi-Class Bias Expressions

For each topology, we derive the plim³ of the naive estimator that substitutes A^* for A . Let τ denote the target causal parameter, whose precise definition is topology-specific (e.g., the coefficient on X for confounding, the category-specific β_j for exposure, or the direct effect τ_{direct} for mediation), and $\hat{\tau}(A^*)$ the estimator using misclassified annotations.

Confounding ($A = \text{confounder}$). When A confounds the $X \rightarrow Y$ relationship, the ordinary least squares (OLS) estimator controlling for A^* has probability limit

$$\text{plim } \hat{\tau}(A^*) = (\mathbb{E}[Z'Z])^{-1} \mathbb{E}[Z'X] \beta, \quad (1)$$

where $Z = [X, \text{dummies}(A^*)]$, β contains the true coefficients, and the target parameter τ is the coefficient on X . Misclassification produces residual confounding that monotonically approaches the omitted variable bias, but the sign of $\hat{\tau}$ is preserved: $\tau_{\text{sign_flip}} = 0$ structurally.⁴

³The probability limit (plim) is the value to which an estimator converges as sample size $N \rightarrow \infty$; it characterizes systematic bias separately from sampling variability.

⁴Consistent with Brenner (1993): misclassification degrades confounding control, producing amplification ($|\hat{\tau}| >$




Figure 2: Bias distributions under 100,000 random $K=3$ confusion matrices (uniform Dirichlet; realistic-accuracy results in Section 4.2). Red = sign reversal. Exposure: 83.4% sign reversal; IV: 52.6%.

Exposure ($A = \text{treatment variable}$). When A is the exposure of interest, the biased coefficient for category j becomes a weighted combination of true coefficients:

$$\text{plim } \hat{\beta}_j(A^*) = \sum_{k=1}^K w_{jk} \beta_k, \quad (2)$$

where w_{jk} depends on the confusion matrix C and the population marginal $\pi_k = P(A=k)$. Unlike binary misclassification (Brenner, 1993), $K \geq 3$ permits w_{jk} to assign negative weight to categories with opposite-sign coefficients, producing *sign reversal* (83.4% under uniform Dirichlet prior); this is structurally impossible at $K=2$.

Primary coefficient preservation. Across 23 empirical LLM configurations, the primary coefficient (the category with the largest true effect $|\beta_k|$) never reverses sign, while 61% reverse a secondary (non-dominant) coefficient, confirmed under 100,000 Dirichlet matrices with diagonal ≥ 0.7 across six beta-ratio conditions (1.0 to 50.0) with 0% primary sign-flip (Appendix A.9).

Mediation ($A = \text{mediator}$). When A mediates the $X \rightarrow Y$ relationship (VanderWeele, 2015), the

$|\tau|$, universal in this topology) but never sign reversal ($\text{sign}(\hat{\tau}) \neq \text{sign}(\tau)$).

direct effect estimator controlling for A^* has probability limit

$$\text{plim } \hat{\tau}_{\text{direct}}(A^*) = \tau_{\text{direct}} + (1 - \lambda(C)) \tau_{\text{indirect}}, \quad (3)$$

where $\lambda(C) \in [0, 1]$ under proportional effects ($f = \lambda g$; formally stated in Theorem 1; Appendix A.5, Proposition 4), with $\lambda(I) = 1$ under perfect classification. Misclassification ($\lambda < 1$) causes the residual indirect effect to be absorbed into the direct path estimate, producing upward bias in $|\hat{\tau}_{\text{direct}}|$ but preserving the sign ($\tau_{\text{sign_flip}} = 0$).

Collider, M-bias, and front-door. Three topologies exhibit protective or immune behavior. **Collider:** conditioning on A^* instead of A attenuates Berkson bias by weakening the spurious path, so annotation error is protective ($\tau_{\text{sign_flip}} = 0$). **M-bias:** in the M-shaped DAG (Greenland, 2003), misclassifying A attenuates the non-causal path, making this topology immune ($\text{ASV} = \infty$). **Front-door:** the front-door estimator (Bellemare et al., 2024) requires two stages ($X \rightarrow A$ and $A \rightarrow Y$), each attenuated by C ; the compound effect produces moderate bias with $\tau_{\text{sign_flip}} = 0$ (Appendix A.8).

Instrumental variable ($A = \text{instrument}$). The Wald estimator (Angrist and Imbens, 1995) $\hat{\tau}_{\text{IV}} = \hat{\text{cov}}(A^*, Y) / \hat{\text{cov}}(A^*, X)$ has C in the denominator. Misclassification can drive first-stage coefficients




Figure 3: ASV vulnerability ranking across topologies. Three ASV types (magnitude, significance, sign-flip) reveal a topology-dependent spectrum from IV (maximally fragile) to M-bias (immune, ASV = ∞).

toward zero or reverse their sign; under random Dirichlet matrices, 52.6% produce sign reversal; IV is the most fragile topology (Appendix B.4).

Structural sign-flip guarantee.

Proposition 1 (Sign preservation: analytic results). *Under non-differential misclassification with $C_{kk} > 0$ for all k : (a) Confounding preserves sign when $\text{sign}(\tau) = \text{sign}(\tau + B_{OV})$, where $B_{OV} = \text{Cov}(g(A), f(A)) / \text{Var}(T)$ is the omitted-variable bias (Cauchy–Schwarz; Appendix A.3). (b) Collider preserves sign when $|\tau| > \sqrt{\text{Var}(g) \text{Var}(h)} / \mathbb{E}[\text{Var}(T|A)]$ (Berkson attenuation; Appendix A.7); empirically verified across 100,000 matrices with zero violations. (c) Exposure can produce sign reversal for $K \geq 3$, even with $C_{kk} > 0.5$ (counterexample in Appendix A.9).*

Sign preservation for non-exposure topologies holds under topology-specific conditions (Table 3, Appendix A.10), verified across 100,000 random matrices with zero violations; analytic proofs for general K remain open.

3.3 Annotation Sensitivity Values

We adapt the sensitivity analysis paradigm (E-values (VanderWeele and Ding, 2017), partial- R^2 bounds (Cinelli and Hazlett, 2020), coefficient stability (Oster, 2019)) to $K \times K$ annotation error and define Annotation Sensitivity Values (ASVs): we prove ε -approximate monotonicity along $C(\delta)$

(Theorem 1), which reduces K^2 dimensions to a scalar; define three ASV diagnostics for topology-dependent failure modes; and combine ASVs with FWL invariance to route corrections.

Parameterization. We set $C(\delta) = (1-\delta)I + \delta C_0$, where C_0 is a reference confusion matrix and $\delta \in [0, 1]$ controls severity ($\delta=0$: perfect; $\delta=1$: matches C_0). This linear path is one of many in the column-stochastic polytope from I to C_0 ; ASV thresholds are path-dependent, but our bootstrap procedure (Section 4.3) captures estimation uncertainty along the empirical path. We recommend bootstrapping C_0 and reporting ASV confidence intervals (Appendix B.8).

Three ASV types. We define three complementary measures: (1) **Magnitude ASV** $\delta_{\text{mag}}^* = \min\{\delta : |\text{plim } \hat{\tau}(C(\delta)) - \tau| > \epsilon\}$, analogous to E-values; (2) **Significance ASV** $\delta_{\text{sig}}^* = \min\{\delta : |\text{plim } \hat{\tau}(C(\delta))| < z_{\alpha/2} \cdot SE(N)\}$, relevant for attenuating topologies; (3) **Sign-flip ASV** $\delta_{\text{sf}}^* = \min\{\delta : \text{sign}(\hat{\tau}) \neq \text{sign}(\tau)\}$, for exposure and IV only. The significance ASV depends on N through $SE(N)$, exhibiting a sharp phase transition at a critical N^* below which all rejections are blocked.

Theorem 1 (Envelope monotonicity, $K=2$). *For $K=2$ confounding and mediation topologies with proportional effects ($f = \lambda g$ for some scalar λ) under $C(\delta) = (1-\delta)I + \delta C_0$ with $(C_0)_{jj} \geq \frac{1}{2}$,*

$|B(C(\delta))|$ is monotonically non-decreasing in $\delta \in [0, 1]$.

The proof (Appendix A.4, A.5) exploits perspective function convexity. All 23 empirical configurations preserve sign; 97.3% do so under non-proportional effects (Appendix B.13).

Theorem 2 ($K=3$ symmetric monotonicity). *For $K=3$ with symmetric confusion matrix $C_{jk} = c$ for $j \neq k$ and $C_{kk} = 1-2c$, $c \in (0, 1/3)$, and proportional effects ($f=\lambda g$), envelope monotonicity holds exactly: $|B(C(\delta))|$ is strictly monotone in δ for confounding, mediation, and collider topologies whenever $\text{Var}(g) > 0$.*

Theorem 2 anchors exact monotonicity at $K=3$; for general asymmetric matrices ($d_{\text{rel}}=0.33$; Remark 5), Conjecture 1 provides $\varepsilon < 0.017$ coverage. The proof exploits the reduction $s = c\delta$ (Appendix A.6).

From proofs to numerical evidence. Exact monotonicity (Theorems 1–2) does not extend to general $K \geq 3$ column-stochastic structures; we state a conjecture supported by numerical evidence:

Conjecture 1 (ε -approximate envelope monotonicity). *For general $K \geq 3$ non-IV topologies under $C(\delta) = (1-\delta)I + \delta C_0$, $|B(C(\delta))|$ is ε -approximately monotonic with $\varepsilon < 0.02$ for $|\beta| \leq 1$, scaling with effect magnitude (Table 12, 14; Appendix B.12). The safe/fragile partition (five sign-preserving topologies versus exposure and IV) is robust across all tested parameterizations.*

Numerical support spans 100,000 Dirichlet matrices and 545,945 adversarial curves (Appendix A.11).

Remark 1 (ASV and E-values as complementary tools). *ASV and E-values (VanderWeele and Ding, 2017) address orthogonal sources (annotation error vs. unmeasured confounding; Appendix B.22) and are most valuable jointly.*

Under alternating-sign effects ($\beta=[0, 1.0, -0.5]$), magnitude ASV ranges from 0.09 (exposure) to ∞ (M-bias), a $4\times$ difference determined by topology (Table 11). Sign-flip ASV < 1 flags sign-reversal-prone pairs; magnitude ASV $< \delta_{\text{observed}}$ indicates bias exceeds tolerance (Algorithm 1); the safe/fragile partition is stable across parameterizations ($\varepsilon < 0.02$; Appendix B.11).

Table 1: Bias under $K=3$ non-differential misclassification (100,000 Dirichlet(1, 1, 1) matrices). Violation: $|\text{bias}| > 0$; Sign Flip: any coefficient reverses sign; Amplif.: $|\hat{\tau}| > |\tau|$. $\tau_{\text{sf}}=0?$: sign-flip guarantee holds unconditionally (Yes), conditionally (Cond.; Proposition 1), or not (No). Collider 100% amplification reflects Berkson bias; misclassification is protective.

Topology	Violation	Sign Flip	Amplif.	$\tau_{\text{sf}}=0?$
Confounding	100%	0%	100%	Cond.
Exposure	85%	83.4%	14.1%	No
Mediation	100%	0%	100%	Cond.
Front-door	100%	0%	100%	Cond.
Collider	100%	0%	100%	Cond.
M-bias	53.2%	0%	10.4%	Yes
IV	100%	52.6%	100%	No




Figure 4: Empirical LLM confusion matrices (markers) overlaid on theoretical violation-rate curves from $K=3$ Dirichlet analysis. Left: exposure; right: IV. Real-world accuracy places all models in the low-risk region for IV sign reversal but not for exposure.

4 Experiments

We validate through Monte Carlo simulation ($K=3$), cross- K robustness checks ($K=5, 7$), and empirical validation with real LLM confusion matrices from three natural language processing (NLP) tasks ($K=2, 3, 10$).

4.1 Setup

For each topology, we sample 100,000 confusion matrices from a Dirichlet(1, 1, 1) prior, compute analytical probability limits, and validate against Monte Carlo simulation ($N=10,000$). ASV experiments use $C(\delta) = (1-\delta)I + \delta C_0$.

4.2 Bias Taxonomy

Theory-MC agreement. Analytical probability limits match Monte Carlo estimates with mean absolute percentage error (MAPE) of 1.6%, averaged across 161 scenarios (23 LLM configurations $\times$ 7 topologies; per-topology range: 0.3%–3.7%), val-

Table 2: ASV summary across topologies ($K=3$, $\tau=0.5$). —: not applicable.

Topology	Mag ASV	Power	Sign-flip
IV	0.10–0.61	0%	1.1–6.7
Front-door	0.06–0.39	100%	—
Confounding	0.08–0.49	100%	—
Mediation	0.08–0.45	0%	—
Collider	0.34–2.03	0%	—
M-bias	∞	98%	—
Exposure	—	—	83.4%

idating the closed-form derivations and justifying probability limits as practical bias diagnostics.

Exposure DAG: sign reversal under multi-class misclassification. Under realistic accuracy (diagonal ≥ 0.7), the violation rate is 13.7% with 2.1% sign reversal; at diagonal ≥ 0.9 , violation persists at 11.9% (Appendix B.4). Under the full Dirichlet prior, 85% of matrices violate the binary “toward-null” guarantee and 83.4% cause sign reversal; $\det(C) < 0.014$ predicts sign reversal with 91% accuracy (Appendix B.25). The $K=2$ control confirms classical attenuation (Brenner, 1993). This establishes the multi-class failure as structurally distinct.

Structural guarantee for non-exposure topologies. Five topologies exhibit zero sign-flip rate (Table 1): confounding, mediation, and front-door show 100% amplification (Brenner, 1993; Dosemeci et al., 1990) but never reverse the sign; collider is protective ($|B(C)|$ non-increasing; Conjecture 1); M-bias shows weakest sensitivity (53.2% violation). All patterns persist at $K=5, 7$ (Appendix B.9).

IV: Wald estimator instability. The IV topology is the sole non-exposure exception: 52.6% of random matrices cause sign reversal, stable across $K=3, 5, 7$. Diagonal dominance ($C_{kk} > 0.5$) eliminates sign reversal but preserves amplification (Appendix B.4); however, all 23 empirical LLM configurations show zero IV sign reversal.

Figure 4 overlays all 23 empirical LLM configurations onto the theoretical violation-rate curves; real-world accuracy ($\min_k C_{kk} \geq 0.6$) places all models in the low-risk region for IV but leaves exposure vulnerable.

4.3 ASV Framework Validation

ASV results. The ASV envelope is ε -approximately monotonic ($\varepsilon < 0.02$; Con-




Figure 5: Topology determines downstream causal bias, not F_1 . HumAID $K=10$ (56 points): within F_1 range 0.614–0.690, bias varies 10 $\times$ (IV 42.1% vs. collider $< 5.6\%$).

jecture 1); IV is maximally fragile, M-bias is immune (ASV = ∞ ; Figure 3). Sign-flip ASV: 83.4% of exposure matrices yield $\delta_{sf}^* < 1$; IV ranges 1.1–6.7 (Table 2). Significance ASV exhibits a phase transition at $N^* \approx 6,000$ –10,000; bootstrap CIs (Appendix B.8, B.24) confirm stable topology ranking, with the safe/fragile partition recovering from $n_{gold} \geq 100$ (82% stability; Appendix B.10, B.18).

4.4 Empirical Validation

We validate using real LLM confusion matrices from HumAID (Alam et al., 2021) ($K=10$), VAST ($K=3$), and CivilComments ($K=2$): four LLMs, 23 configurations, 161 bias scenarios. HumAID ($K=10$) matches predictions: IV (12.4–42.1%) and exposure (9.0–17.4%) dominate; remaining topologies $< 15\%$. F_1 does not exhibit significant rank correlation with bias ($\rho = -0.10$, $p = 0.45$; Figure 5), but per-topology patterns diverge ($\rho \leq -0.53$ for five topologies vs. 0.87 for M-bias; Appendix B.23). Of exposure coefficients, 14/23 (61%) exhibit sign reversal, all secondary; zero across 138 non-exposure scenarios.

4.5 Correction Guidance

We evaluate five correction methods (naive plug-in, maximum likelihood adjustment (MLA), design-based supervised learning (DSL) (Egami et al., 2023), Monte Carlo simulation extrapolation (MC-SIMEX) (Küchenhoff et al., 2006), PPI++ (Angelopoulos et al., 2023)) across all topologies (210 configurations; Table 6, Appendix B).




Figure 6: Binarization collapse Δ : additional bias from collapsing $K=3$ to binary. Topology determines severity ($>145\times$ range); exposure and IV exhibit sign flips.

FWL invariance. Under linear regression, FWL invariance renders matrix-based methods (naive plug-in, MLA) provably ineffective for covariate topologies (0% bias reduction; Appendix B.2): the treatment coefficient depends only on the projection orthogonal to the covariate space, so relabeling $A^* \rightarrow A$ leaves the coefficient unchanged. Under logistic regression this invariance breaks, but the method ranking persists ($|\Delta| < 0.02$; Appendix B.14).

Method ranking and decision procedure. PPI++ achieves the highest root mean squared error (RMSE) reduction (47.0% at $K=3$) but requires gold labels (Appendix B.15) and harms M-bias (−88%); MC-SIMEX is gold-free and K -robust (27–70%), though it harms front-door (−62.4%); rankings are stable across K ($\rho \geq 0.80$; Appendix B.21). Algorithm 1 routes through risk assessment, method selection, and FWL screening. End-to-end demos (Appendix B.19, B.20, B.3) validate: PPI++ reduces exposure bias from 42.9% to 0.2%, while topology-agnostic IV correction amplifies bias (−139%). Collapsing $K=3$ to binary is not a safe fallback: the additional bias Δ varies $>145\times$ by topology (Figure 6; Appendix B.1).

5 Conclusion

We characterized how multi-class annotation errors interact with causal DAG topology: at $K \geq 3$, exposure and IV can reverse treatment effect signs while five topologies preserve them. ASVs provide a topology-dependent vulnerability ranking; matrix-based corrections are provably ineffective for four topologies (FWL invariance). Practitioners should identify the topology first: for exposure, PPI++ reduces bias from $\sim 61\%$ to $<1\%$ (Table 21);

Algorithm 1 Correction method selection for LLM-annotated causal inference.

Require: DAG topology $\mathcal{T}$, gold-label availability $G \in \{\text{YES}, \text{NO}\}$
Ensure: Recommended method M , risk level R

- 1: **Step 1: Risk assessment** {If $\mathcal{T}$ uncertain, repeat per candidate}
- 2: $R \leftarrow$ LOW if M-bias; MED if Conf/Med/Coll/FD; HIGH if Exp/IV
- 3: **Step 2: Method selection**
- 4: **if** $G = \text{YES}$ **then**
- 5: $M \leftarrow$ PPI++ (47% RMSE red.) except M-bias $\rightarrow$ NONE (−88%)
- 6: **else**
- 7: $M \leftarrow$ DSL if Conf/Med; MC-SIMEX if IV/Coll/M-bias; $\emptyset$ otherwise
- 8: **end if**
- 9: **Step 3: FWL screening**
- 10: **if** $\mathcal{T} \in \{\text{Conf.}, \text{Med.}, \text{Coll.}, \text{M-bias}\}$ **then**
- 11: Skip Naive plug-in and MLA {0% bias red. on T }
- 12: **end if**
- 13: **return** (M, R)

for IV, topology-agnostic correction can amplify bias (−139% in our demo; Table 20). High F_1 does not guarantee low causal bias; topology assessment must precede downstream inference.

The practical workflow we recommend is: (1) draw the DAG and locate the annotated variable, (2) compute ASVs to quantify vulnerability, (3) apply Algorithm 1 for topology-aware method selection, and (4) report ASV confidence intervals alongside effect estimates. This pipeline requires only a gold-label subsample and the confusion matrix—resources already available in most LLM annotation workflows.

Our framework opens several directions. Extending the analysis to *differential* misclassification—where error rates depend on Y or covariates—would cover settings where annotator difficulty correlates with outcome severity. Nonlinear identification strategies (matching, doubly robust estimation) break the FWL invariance that underlies our correction routing; characterizing which topologies remain tractable under nonlinear models is an important next step. Proving Conjecture 1 for general K would place the ASV framework on exact rather than ϵ -approximate foundations.

Limitations

Our closed-form results assume linear SEMs with proportional effects ($f = \lambda g$); sign preservation empirically holds at 97.3% under non-proportional effects (Appendix B.17), but recommendations do not directly extend to nonlinear models (matching, doubly-robust). The framework assumes a

known DAG with non-differential misclassification (Appendix B.6, B.16); DAG misspecification under measurement error may compound bias. ASV computation requires gold labels for C^{-1} ; error propagation under small gold-label budgets (<500) and general $K \geq 3$ envelope monotonicity (Conjecture 1) remain open.

References

Firoj Alam, Umair Qazi, Muhammad Imran, and Ferda Ofli. 2021. HumAID: Human-annotated disaster incidents data from Twitter with deep learning benchmarks. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 15, pages 933–942.

Anastasios N. Angelopoulos, Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnić. 2023. [Prediction-powered inference](#). *Science*, 382(6671):669–674.

Joshua D. Angrist and Guido W. Imbens. 1995. [Two-stage least squares estimation of average causal effects in models with variable treatment intensity](#). *Journal of the American Statistical Association*, 90(430):431–442.

Joachim Baumann, Paul Röttger, Aleksandra Urman, Albert Wendsjö, Flor Miriam Plaza-del Arco, Johannes B. Gruber, and Dirk Hovy. 2025. Large language model hacking: Quantifying the hidden risks of using LLMs for text annotation. *arXiv preprint arXiv:2509.08825*. Preprint.

Marc F. Bellemare, Jeffrey R. Bloem, and Noah Wexler. 2024. [The paper of how: Estimating treatment effects using the front-door criterion](#). *Oxford Bulletin of Economics and Statistics*, 86(4):951–993.

Matthew Blackwell, James Honaker, and Gary King. 2017. [A unified approach to measurement error and missing data: Overview and applications](#). *Sociological Methods & Research*, 46(3):303–341.

Alana T. Brennan, Kelly D. Getz, Daniel R. Brooks, and Matthew P. Fox. 2021. [An underappreciated misclassification mechanism: Implications of nondifferential dependent misclassification of covariate and exposure](#). *Annals of Epidemiology*, 58:104–123.

Hermann Brenner. 1993. [Bias due to non-differential misclassification of polytomous confounders](#). *Journal of Clinical Epidemiology*, 46(1):57–63.

Irwin Bross. 1954. [Misclassification in 2 x 2 tables](#). *Biometrics*, 10(4):478–486.

Raymond J. Carroll, David Ruppert, Leonard A. Stefanski, and Ciprian M. Crainiceanu. 2006. *Measurement Error in Nonlinear Models: A Modern Perspective*, 2nd edition. Chapman & Hall/CRC.

Carlos Cinelli and Chad Hazlett. 2020. [Making sense of sensitivity: Extending omitted variable bias](#). *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 82(1):39–67.

Mustafa Dosemeci, Sholom Wacholder, and Jay H. Lubin. 1990. [Does nondifferential misclassification of exposure always bias a true effect toward the null value?](#) *American Journal of Epidemiology*, 132(4):746–748.

Naoki Egami, Musashi Hinck, Brandon M. Stewart, and Hanying Wei. 2023. Using imperfect surrogates for downstream inference: Design-based supervised learning for social science applications of large language models. In *Advances in Neural Information Processing Systems* 36.

Ragnar Frisch and Frederick V. Waugh. 1933. Partial time regressions as compared with individual trends. *Econometrica*, 1(4):387–401.

Fabrizio Gilardi, Meysam Alizadeh, and Maël Kubli. 2023. [ChatGPT outperforms crowd workers for text-annotation tasks](#). *Proceedings of the National Academy of Sciences*, 120(30):e2305016120.

Kristina Gligorić, Tijana Zrnić, Cinoo Lee, Emmanuel J. Candès, and Dan Jurafsky. 2025. Can unconfident LLM annotations be used for confident conclusions? In *Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics*.

Sander Greenland. 2003. [Quantifying biases in causal models: Classical confounding vs collider-stratification bias](#). *Epidemiology*, 14:300–306.

Miguel A. Hernán and Stephen R. Cole. 2009. [Invited commentary: Causal diagrams and measurement bias](#). *American Journal of Epidemiology*, 170(8):959–962.

Miguel A. Hernán and James M. Robins. 2020. *Causal Inference: What If*. Chapman & Hall/CRC.

Anne M. Jurek, Sander Greenland, George Maldonado, and Timothy R. Church. 2005. [Proper interpretation of non-differential misclassification effects: Expectations vs observations](#). *International Journal of Epidemiology*, 34(3):680–687.

Elliot Kim, Avi Garg, Kenny Peng, and Nikhil Garg. 2025. Correlated errors in large language models. In *Proceedings of the 42nd International Conference on Machine Learning*.

Helmut Küchenhoff, Samuel M. Mwalili, and Emmanuel Lesaffre. 2006. A general method for dealing with misclassification in regression: The misclassification SIMEX. *Biometrics*, 62(1):85–96.

Manabu Kuroki and Judea Pearl. 2014. [Measurement bias and effect restoration in causal inference](#). *Biometrika*, 101(2):423–437.

Timothy L. Lash, Matthew P. Fox, and Aliza K. Fink. 2009. *Applying Quantitative Bias Analysis to Epidemiologic Data*. Statistics for Biology and Health. Springer.

Samuel Lee and Zach Wood-Doughty. 2024. Controlling for unobserved confounding with large language model classification of patient smoking status. In *AIM-FM Workshop, NeurIPS 2024*. Workshop paper.

Michael C. Lovell. 1963. Seasonal adjustment of economic time series and multiple regression analysis. *Journal of the American Statistical Association*, 58(304):993–1010.

Elizabeth L. Ogburn and Tyler J. VanderWeele. 2012. On the nondifferential misclassification of a binary confounder. *Epidemiology*, 23(3):433–439.

Emily Oster. 2019. Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics*, 37(2):187–204.

Judea Pearl. 2009. *Causality: Models, Reasoning and Inference*, 2nd edition. Cambridge University Press.

Susanne M. Schennach. 2020. Mismeasured and unobserved variables. In Steven N. Durlauf, Lars Peter Hansen, James J. Heckman, and Rosa L. Matzkin, editors, *Handbook of Econometrics*, volume 7A, pages 487–565. Elsevier.

Di Shu and Grace Y. Yi. 2019. Weighted causal inference methods with mismeasured covariates and misclassified outcomes. *Statistics in Medicine*, 38(10):1835–1854.

Zhen Tan, Dawei Li, Song Wang, Alimohammad Beigi, Bohan Jiang, Amrita Bhattacharjee, Mansoor Karami, Jundong Li, Lu Cheng, and Huan Liu. 2024. Large language models for data annotation and synthesis: A survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 930–957.

Nathan TeBlunthuis, Valerie Hase, and Chung-hong Chan. 2024. Misclassification in automated content analysis causes bias in regression. Can we fix it? Yes we can! *Communication Methods and Measures*, 18(3):278–299.

Linda Valeri and Tyler J. VanderWeele. 2013. Mediation analysis allowing for exposure-mediator interactions and causal interpretation: Theoretical assumptions and implementation with SAS and SPSS macros. *Psychological Methods*, 18(2):137–150.

Tyler J. VanderWeele. 2015. *Explanation in Causal Inference: Methods for Mediation and Interaction*. Oxford University Press, New York.

Tyler J. VanderWeele and Peng Ding. 2017. Sensitivity analysis in observational research: Introducing the E-value. *Annals of Internal Medicine*, 167:268–274.

Marit B. Veierød and Petter Laake. 2001. Exposure misclassification: Bias in category specific Poisson regression coefficients. *Statistics in Medicine*, 20(5):771–784.

Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. 2024. Can large language models transform computational social science? *Computational Linguistics*, 50(1):237–291.

Tijana Zrnić and Emmanuel J. Candès. 2024. Cross-prediction-powered inference. *Proceedings of the National Academy of Sciences*, 121(15):e2322083121.

A Proofs of Proposition 1 and Theorem 1

We provide algebraic proofs of sign preservation for the confounding and collider topologies (Proposition 1(a)–(b)), an explicit counterexample for exposure (Proposition 1(c)), numerical verification for mediation, M-bias, and front-door (Numerical Verification 1), and the monotonicity proof for confounding (Theorem 1). Throughout, $A \in \{0, 1, \dots, K-1\}$ with prior $p_k = P(A=k)$, and C is a $K \times K$ column-stochastic confusion matrix with $C_{jk} = P(A^*=j \mid A=k)$. We use category 0 as the reference and let $D_j = 1(A=j)$, $D_j^* = 1(A^*=j)$ for $j = 1, \dots, K-1$.

A.1 Notation and Key Identities

Under non-differential misclassification, the misclassified moments factor linearly through C . Define the population vectors:

$$q_j = P(A^*=j) = \sum_k C_{jk} p_k, \quad (4)$$

$$\text{ct}_j = \mathbb{E}[T \cdot D_j^*] = \sum_k C_{jk} \mathbb{E}[T \cdot D_k], \quad (5)$$

$$\text{cy}_j = \mathbb{E}[Y \cdot D_j^*] = \sum_k C_{jk} \mathbb{E}[Y \cdot D_k]. \quad (6)$$

Write $q = Cp$, $\text{ct} = C \mathbb{E}[T \cdot D]$, and $\text{cy} = C \mathbb{E}[Y \cdot D]$, where $D = (D_0, \dots, D_{K-1})^\top$.

A.2 Frisch–Waugh–Lovell Decomposition

For topologies where T is directly observed (confounding, mediation, collider, M-bias, front-door), the estimator regresses Y on $Z^* = (1, T, D_1^*, \dots, D_{K-1}^*)$. By the Frisch–Waugh–Lovell (FWL) theorem, the coefficient on T satisfies:

$$\text{plim } \hat{\tau}(C) = \frac{\mathbb{E}[\tilde{T} \tilde{Y}]}{\mathbb{E}[\tilde{T}^2]}, \quad (7)$$

where $\tilde{T} = T - \mathbb{E}[T \mid A^*]$ and $\tilde{Y} = Y - \mathbb{E}[Y \mid A^*]$ are residuals from projecting onto $(1, D_1^*, \dots, D_{K-1}^*)$.

Lemma 1 (Bias decomposition). Suppose $Y = \tau T + f(A) + \varepsilon_Y$ with $\varepsilon_Y \perp (T, A, A^*)$. Then

$$\text{plim } \hat{\tau}(C) = \tau + \frac{R(C)}{D(C)}, \quad (8)$$

where

$$R(C) = \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)], \quad (9)$$

$$D(C) = \mathbb{E}[\text{Var}(T \mid A^*)] > 0. \quad (10)$$

Proof. Expand $\tilde{Y}$ using $Y = \tau T + f(A) + \varepsilon_Y$:

$$\tilde{Y} = \tau \tilde{T} + (f(A) - \mathbb{E}[f(A) \mid A^*]) + \varepsilon_Y.$$

Since $\tilde{T}$ is measurable w.r.t. (T, A^*) and $\varepsilon_Y \perp (T, A, A^*)$:

$$\mathbb{E}[\tilde{T} \tilde{Y}] = \tau \mathbb{E}[\tilde{T}^2] + \mathbb{E}[\tilde{T} (f(A) - \mathbb{E}[f(A) \mid A^*])].$$

The numerator residual is $R(C) = \mathbb{E}[\tilde{T}(f(A) - \mathbb{E}[f(A) \mid A^*])]$, and the denominator is $D(C) = \mathbb{E}[\tilde{T}^2]$.

By the law of total covariance:

$$\begin{aligned} R(C) &= \sum_{j=0}^{K-1} q_j \text{Cov}(T, f(A) \mid A^*=j) \\ &= \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)]. \end{aligned}$$

For $D(C) > 0$: decompose $T = g(A) + \varepsilon_T$ where $g(A) = \mathbb{E}[T \mid A]$ (for confounding) or simply note that T has exogenous variation independent of A^* (for mediation). In both cases, $\tilde{T} = (g(A) - \mathbb{E}[g(A) \mid A^*]) + \varepsilon'_T$, where ε'_T is the component of T independent of (A, A^*) . Hence $D(C) \geq \text{Var}(\varepsilon'_T) > 0$. $\square$

A.3 Confounding Topology ($A = \text{Confounder}$)

Data-generating process. $A \rightarrow T, A \rightarrow Y, T \rightarrow Y$. Using dummy coding with category 0 as reference:

$$\begin{aligned} T &= \sum_{j=1}^{K-1} \alpha_j D_j + \varepsilon_T, \quad \varepsilon_T \sim \mathcal{N}(0, \sigma_T^2), \\ Y &= \tau T + \sum_{j=1}^{K-1} \beta_j^{(A)} D_j + \varepsilon_Y, \\ \varepsilon_Y &\sim \mathcal{N}(0, 1), \end{aligned}$$

with $\varepsilon_T \perp A$ and $\varepsilon_Y \perp (T, A)$. Define $g(k) = \sum_j \alpha_j \mathbf{1}(k=j)$ and $f(k) = \sum_j \beta_j^{(A)} \mathbf{1}(k=j)$, so $g(0) = f(0) = 0$.

Closed-form plim. Using the FWL decomposition, the coefficient on T is:

$$\text{plim } \hat{\tau}(C) = \frac{\mu_{TY} - \sum_{j=0}^{K-1} \frac{\text{ct}_j \text{cy}_j}{q_j}}{s_T - \sum_{j=0}^{K-1} \frac{\text{ct}_j^2}{q_j}}, \quad (11)$$

where $s_T = \mathbb{E}[T^2]$, $\mu_{TY} = \mathbb{E}[TY]$, and $\text{ct}_j, \text{cy}_j, q_j$ are defined in (4)–(6).

Derivation. By FWL, $\text{plim } \hat{\tau} = \mathbb{E}[\tilde{T}Y]/\mathbb{E}[\tilde{T}^2]$ where $\tilde{T} = T - \mathbb{E}[T \mid A^*]$. Since $\mathbb{E}[T \mid A^*=j] = \text{ct}_j/q_j$:

$$\begin{aligned} \mathbb{E}[\tilde{T}^2] &= \mathbb{E}[T^2] - \sum_j q_j (\text{ct}_j/q_j)^2 \\ &= s_T - \sum_j \text{ct}_j^2/q_j, \end{aligned}$$

$$\begin{aligned} \mathbb{E}[\tilde{T}Y] &= \mathbb{E}[TY] - \sum_j (\text{ct}_j/q_j) \mathbb{E}[Y \cdot D_j^*] \\ &= \mu_{TY} - \sum_j \text{ct}_j \text{cy}_j/q_j. \end{aligned}$$

$\square$

Verification at $C = I$. When $C = I$: $q_j = p_j$, $\text{ct}_j = \mathbb{E}[TD_j] = \alpha_j p_j$, $\text{cy}_j = \mathbb{E}[YD_j] = p_j(\tau\alpha_j + \beta_j^{(A)})$. Direct substitution yields numerator $= \tau\sigma_T^2$ and denominator $= \sigma_T^2$, giving $\text{plim } \hat{\tau} = \tau$.

Bias structure. By Lemma 1:

$$\begin{aligned} \text{plim } \hat{\tau}(C) &= \tau + B(C), \\ B(C) &= \frac{\mathbb{E}[\text{Cov}(g(A), f(A) \mid A^*)]}{D(C)}. \end{aligned} \quad (12)$$

By the law of total covariance:

$$\begin{aligned} \mathbb{E}[\text{Cov}(g, f \mid A^*)] &= \text{Cov}(g(A), f(A)) \\ &\quad - \text{Cov}(\mathbb{E}[g(A) \mid A^*], \mathbb{E}[f(A) \mid A^*]). \end{aligned} \quad (13)$$

At $C = I$: $\mathbb{E}[g \mid A^*] = g(A)$, so the between-group covariance equals $\text{Cov}(g, f)$ and $B(I) = 0$.

At $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$ (uniform, $A^* \perp A$): $\mathbb{E}[g \mid A^*]$ is constant, so the between-group covariance is zero and

$$B_{\text{OV}} \equiv B\left(\frac{1}{K} \mathbf{1}\mathbf{1}^\top\right) = \frac{\text{Cov}(g(A), f(A))}{\text{Var}(T)}, \quad (14)$$

which equals the classical omitted-variable bias (no control for A).

Sign preservation condition. Since $D(C) > 0$, $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau \cdot D(C) + R(C))$. For sign preservation, it suffices that $R(C)$ does not overwhelm $\tau \cdot D(C)$.

Proposition 2 (Confounding sign preservation — proportional case). *If $f(A) = \lambda g(A)$ for some constant λ (i.e., the confounder effects on T and Y are proportional), then*

$$B(C) = (1 - \xi(C)) B_{\text{OV}}, \quad (15)$$

where

$$\xi(C) = \frac{V_B \sigma_T^2}{V(V - V_B + \sigma_T^2)} \in [0, 1],$$

where $V_B = \text{Var}(\mathbb{E}[g(A) \mid A^*])$ and $V = \text{Var}(g(A))$. Consequently, $B(C)$ lies between 0 and B_{OV} , and $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau)$ whenever $\text{sign}(\tau) = \text{sign}(\tau + B_{\text{OV}})$.

Proof. When $f = \lambda g$: the between-group covariance in (13) becomes $\lambda \text{Var}(\mathbb{E}[g \mid A^*])$, and $\text{Cov}(g, f) = \lambda \text{Var}(g)$. So $R(C) = \lambda(\text{Var}(g) - \text{Var}(\mathbb{E}[g \mid A^*])) = \lambda V_W$, where $V_W = \text{Var}(g) - V_B$ with $V_B = \text{Var}(\mathbb{E}[g \mid A^*])$. Similarly, $D(C) = V_W + \sigma_T^2$. Hence

$$B(C) = \frac{\lambda V_W}{V_W + \sigma_T^2}, \quad B_{\text{OV}} = \frac{\lambda V}{V + \sigma_T^2},$$

where $V = \text{Var}(g) = V_B + V_W$. The ratio

$$\frac{B(C)}{B_{\text{OV}}} = \frac{V_W(V + \sigma_T^2)}{V(V_W + \sigma_T^2)} = 1 - \xi(C)$$

is a decreasing function of $V_B \in [0, V]$. At $V_B = 0$ (no classification, C uniform): $\xi = 0$, $B = B_{\text{OV}}$. At $V_B = V$ (perfect classification, $C = I$): $\xi = 1$, $B = 0$. Since $0 \leq V_B \leq V$, we have $\xi \in [0, 1]$ and $B(C) \in [0, B_{\text{OV}}]$ (with the sign of B_{OV}). $\square$

Proposition 3 (Confounding sign preservation — general case). *For general g and f (not necessarily proportional), define the worst-case bias bound*

$$\bar{B} = \frac{\sqrt{\text{Var}(g(A)) \text{Var}(f(A))}}{\sigma_T^2}. \quad (16)$$

If $|\tau| > \bar{B}$, then $\text{sign}(\text{plim } \hat{\tau}(C)) = \text{sign}(\tau)$ for all column-stochastic C .

Proof. We need $\text{sign}(\tau \cdot D(C) + R(C)) = \text{sign}(\tau)$, which holds if $|R(C)| < |\tau| \cdot D(C)$. Since $D(C) \geq \sigma_T^2$ (shown in Lemma 1), it suffices to show $|R(C)| < |\tau| \sigma_T^2$.

Step 1: Bound $|R(C)|$. $R(C) = \mathbb{E}[\text{Cov}(g(A), f(A) \mid A^*)]$. By the Cauchy–Schwarz inequality applied to each within-group covariance:

$$\begin{aligned} |\text{Cov}(g, f \mid A^*=j)| & \\ & \leq \sqrt{\text{Var}(g \mid A^*=j) \text{Var}(f \mid A^*=j)}. \end{aligned}$$

Taking expectations and applying Cauchy–Schwarz again:

$$\begin{aligned} |R(C)| & \leq \mathbb{E}[\sqrt{\text{Var}(g \mid A^*) \text{Var}(f \mid A^*)}] \\ & \leq \sqrt{\mathbb{E}[\text{Var}(g \mid A^*)] \mathbb{E}[\text{Var}(f \mid A^*)]} \\ & \leq \sqrt{\text{Var}(g(A)) \text{Var}(f(A))}, \end{aligned}$$

where the last step uses $\mathbb{E}[\text{Var}(X \mid Y)] \leq \text{Var}(X)$ (law of total variance).

Step 2: Conclude. The bias satisfies

$$|B(C)| = \frac{|R(C)|}{D(C)} \leq \frac{\sqrt{\text{Var}(g) \text{Var}(f)}}{\sigma_T^2} = \bar{B}.$$

When $|\tau| > \bar{B}$, the bias cannot cancel τ , so $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau)$ for all C . $\square$

Relationship to the omitted-variable condition.

A weaker (necessary) condition for sign preservation is $\text{sign}(\tau) = \text{sign}(\tau_{\text{OV}})$ where $\tau_{\text{OV}} = \tau + \text{Cov}(g, f) / \text{Var}(T)$. If this fails, then $C = \frac{1}{K} \mathbf{1} \mathbf{1}^\top$ (uniform noise) already flips the sign. The rigorous sufficient condition $|\tau| > \bar{B}$ implies the weaker condition, since $|B_{\text{OV}}| = |\text{Cov}(g, f)| / \text{Var}(T) \leq \bar{B}$ by Cauchy–Schwarz and $\text{Var}(T) \geq \sigma_T^2$.

Verification for paper parameters ($K=3$).

With $\tau=2.0$, $\alpha=(1.5, -1.0)$, $\beta^{(A)}=(1.0, -0.5)$, $\sigma_T=1.0$, $\mathbf{p}=(0.40, 0.35, 0.25)$:

$$\begin{aligned} \text{Var}(g) & = \sum_k p_k \alpha_k^2 - (\sum_k p_k \alpha_k)^2 \\ & = 0.962, \end{aligned}$$

$$\begin{aligned} \text{Var}(f) & = \sum_k p_k (\beta_k^{(A)})^2 - (\sum_k p_k \beta_k^{(A)})^2 \\ & = 0.362, \end{aligned}$$

$$\bar{B} = \sqrt{0.962 \times 0.362} / 1.0^2 = 0.590.$$

Since $|\tau|=2.0 > \bar{B}=0.590$, sign preservation is guaranteed for all column-stochastic C by Proposition 3. Moreover, $B_{\text{OV}} = \text{Cov}(g, f) / \text{Var}(T) = 0.588/1.962 = +0.300$, giving $\tau_{\text{OV}}=2.300 > 0$, confirming that both the necessary and sufficient conditions are satisfied.

Extension to $K > 3$. For $K > 3$, treatment effect coefficients are drawn from a fixed pool with alternating signs and decreasing magnitudes:

$$\beta = (0, 1.0, -0.5, 0.7, -0.3, 0.5, -0.2, 0.4, -0.15, 0.3, \dots),$$

truncated to K entries. The category distribution uses $p_k \propto 1/(k+1)$, normalized to sum to 1. Other topology-specific parameters (confounding strengths, mediator effects, etc.) follow analogous alternating-sign pools.

A.4 Proof of Theorem 1 (Confounding Monotonicity)

Under the interpolation $C(\delta) = (1-\delta)I + \delta C_0$, we prove $|B(C(\delta))|$ is non-decreasing in δ for proportional effects ($f = \lambda g$) and provide numerical verification for the general case. The proof proceeds by direct algebraic analysis of $V_B(\delta) = \text{Var}(\mathbb{E}[g(A) | A_\delta^*])$ as a rational function of δ , establishing convexity and bounding the derivative at both endpoints.

Proportional case ($f = \lambda g$). From Proposition 2, $B(C) = \lambda V_W / (V_W + \sigma_T^2)$ where $V_W = \text{Var}(g) - V_B$ and $V_B(\delta) = \text{Var}(\mathbb{E}[g(A) | A_\delta^*])$.

Step 1: Convexity of $V_B(\delta)$. Define $N_j(\delta) = \sum_k g_k C(\delta)_{jk} p_k$ and $q_j(\delta) = \sum_k C(\delta)_{jk} p_k$. Both are affine in δ , with $q_j(\delta) > 0$. Then $V_B(\delta) = \sum_j N_j(\delta)^2 / q_j(\delta) - \bar{g}^2$. Each term N_j^2 / q_j is the composition of the perspective function $(x, t) \mapsto x^2 / t$ (convex for $t > 0$) with the affine map $\delta \mapsto (N_j(\delta), q_j(\delta))$, hence convex in δ . A sum of convex functions is convex, so $V_B(\delta)$ is convex on $[0, \delta_{\max}]$.

Step 2: Derivative formula. Let $\mu_j(\delta) = N_j(\delta) / q_j(\delta) = \mathbb{E}[g(A) | A^* = j]$. Then

$$\frac{dV_B}{d\delta} = - \sum_{j,k} [(C_0)_{jk} - \mathbf{1}_{j=k}] p_k (g_k - \mu_j(\delta))^2. \quad (17)$$

Derivation. $\frac{d}{d\delta} \frac{N_j^2}{q_j} = 2\mu_j \dot{N}_j - \mu_j^2 \dot{q}_j$ where $\dot{N}_j = m_j^0 - g_j p_j$ and $\dot{q}_j = q_j^0 - p_j$ (with $m_j^0 = \sum_k g_k (C_0)_{jk} p_k$, $q_j^0 = \sum_k (C_0)_{jk} p_k$). Substituting and using column-stochasticity ($\sum_j (C_0)_{jk} = \sum_j \mathbf{1}_{j=k} = 1$) to cancel the g_k^2 terms yields (17).

Step 3: $V'_B(0) \leq 0$ (universal, no conditions on C_0). At $\delta=0$: $\mu_j(0) = g_j$ (since $C(0) = I$), so $(g_k - \mu_j(0))^2 = (g_k - g_j)^2$. The diagonal terms

($j=k$) vanish. Hence

$$V'_B(0) = - \sum_{j \neq k} (C_0)_{jk} p_k (g_j - g_k)^2 \leq 0,$$

since every term is non-negative.

Step 4: $V'_B(1) \leq 0$ for $K=2$. For $K=2$ with $C_0 = \begin{bmatrix} a & 1-b \\ 1-a & b \end{bmatrix}$ and $a, b \geq \frac{1}{2}$, set $g_0=0$, $g_1=1$ without loss of generality (the general case scales by $(\Delta g)^2$). At $\delta=1$: $q_0=ap_0+(1-b)p_1$, $q_1=(1-a)p_0+bp_1$, $\mu_0=(1-b)p_1/q_0$, $\mu_1=bp_1/q_1$. Expanding (17) at $\delta=1$ and factoring:

$$\begin{aligned} V'_B(1) &= (1-a) p_0 (\mu_0^2 - \mu_1^2) \\ &\quad + (1-b) p_1 [(1-\mu_1)^2 - (1-\mu_0)^2] \\ &= (\mu_0 - \mu_1) [(1-a) p_0 (\mu_0 + \mu_1) \\ &\quad + (1-b) p_1 (2 - \mu_0 - \mu_1)]. \end{aligned}$$

The bracket is non-negative (since $0 \leq \mu_j \leq 1$ and all coefficients are non-negative). Direct calculation gives $\mu_0 - \mu_1 = p_0 p_1 (1-a-b) / (q_0 q_1)$ (expand $(1-b)q_1 - b q_0 = p_0(1-a-b)$). Under $a, b \geq \frac{1}{2}$: $a+b \geq 1$, so $\mu_0 - \mu_1 \leq 0$ and $V'_B(1) \leq 0$.

Step 5: Monotonicity from convexity. Since V_B is convex (Step 1), its derivative V'_B is non-decreasing. When $V'_B(1) \leq 0$, we have $V'_B(\delta) \leq V'_B(1) \leq 0$ for all $\delta \in [0, 1]$, so V_B is non-increasing.

Step 6: $|B(\delta)|$ is non-decreasing. $V_W(\delta) = \text{Var}(g) - V_B(\delta)$ is non-decreasing (Step 5). The function $h(x) = |\lambda| x / (x + \sigma_T^2)$ is strictly increasing for $x \geq 0$. Hence $|B(\delta)| = h(V_W(\delta))$ is non-decreasing. $\square$

Condition on C_0 . The proof requires $V'_B(1) \leq 0$, which is guaranteed for $K=2$ when $(C_0)_{jj} \geq \frac{1}{2}$ (Step 4). For $K \geq 3$, $V'_B(1) \leq 0$ holds for all g and p if and only if a certain matrix $M - W$ (depending on C_0 and p) is negative semi-definite; simple diagonal thresholds are sufficient in practice but not universal. Numerical verification across 20,000 random (C_0, p, g) triples per $K \in \{3, 5, 7\}$ with $(C_0)_{jj} \geq 1/K$ finds worst-case $|B|$ step deviations < 0.013 (Table 12), confirming ε -approximate monotonicity. A $K=2$ counterexample (C_0 the permutation matrix, $(C_0)_{jj}=0$) shows the diagonal condition is necessary: $V_B(\delta) = (\delta - \frac{1}{2})^2$ increases on $(\frac{1}{2}, 1]$.

General case (non-proportional effects). $|B(C(\delta))| \leq \bar{B}(\delta)$ where $\bar{B}(\delta) = \sqrt{\mathbb{E}[\text{Var}(g | A_\delta^*)] \mathbb{E}[\text{Var}(f | A_\delta^*)] / D(\delta)}$ (Proposition 3). By the same convexity argument,

each $\mathbb{E}[\text{Var}(\cdot \mid A_\delta^*)]$ is non-decreasing in δ , and $D(\delta) = \mathbb{E}[\text{Var}(T \mid A_\delta^*)]$ is non-decreasing. Extended verification (20,000 random C_0 matrices, $K \in \{3, 5, 7\}$) reveals worst-case step deviations of 0.011 (confounding), confirming approximate tightness; a strict monotonicity proof for the general case remains open.

Remark 2 (Conditional exact monotonicity). *Exact monotonicity ($V_B'(\delta) \leq 0$ for all δ) holds under the component-wise diagonal dominance condition:*

$$\alpha_j \geq \frac{q_j}{p_j + q_j} \quad \text{for all } j, \quad (18)$$

where $\alpha_j = (C_0)_{jj}$ and $q_j = \sum_k (C_0)_{jk} p_k$. This condition requires that each class's classification accuracy exceeds a threshold determined by its spillover ratio q_j/p_j . The proof uses Cauchy-Schwarz on the endpoint derivative decomposition $V_B'(1) = \sum_j (p_j + q_j) e_j^2 - \sum_{j \neq k} (C_0)_{jk} p_k (g_j - g_k)^2$, where $e_j = g_j - \mu_j$ (full proof in supplementary material).

For doubly stochastic C_0 with uniform p , condition (18) reduces to $\alpha_j \geq 1/2$, which for $K=2$ recovers the exact condition in Theorem 1. Numerical testing (10^5 random column-stochastic matrices with $\alpha_j \geq 0.5$) shows the condition holds in 99.3% of cases; the 0.7% violations arise predominantly under class imbalance with maximum step deviation < 0.001 . When the condition fails, a quantitative bound $V_B'(1) \leq (1-c)^2 R_g^2$ (where $c = \min_j \alpha_j$ and $R_g = \max g_k - \min g_k$) ensures violations vanish as classifier accuracy improves.

A.5 Mediation Topology ($A = \text{Mediator}$)

Data-generating process. $T \rightarrow A \rightarrow Y, T \rightarrow Y$.

$$\begin{aligned} T &\sim \text{Bernoulli}(0.5), \\ A \mid T &\sim \text{Categorical}(\text{softmax}(\gamma + \delta T)), \\ Y &= \tau_{\text{dir}} T + f(A) + \varepsilon_Y. \end{aligned}$$

Bias decomposition. By Lemma 1, the direct-effect estimator satisfies:

$$\text{plim } \hat{\tau}_{\text{dir}}(C) = \tau_{\text{dir}} + \frac{\mathbb{E}[\text{Cov}(T, f(A) \mid A^*)]}{D(C)}. \quad (19)$$

The key structural difference from confounding: T causes A , so the residual covariance $\text{Cov}(T, f(A) \mid A^*)$ captures the indirect effect of T on Y through A that is not absorbed by conditioning on A^* .

Endpoint analysis. At $C = I$ ($A^*=A$): conditioning on A blocks the mediated path $T \rightarrow A \rightarrow Y$, so $\text{Cov}(T, f(A) \mid A) = 0$ and $\text{plim } \hat{\tau} = \tau_{\text{dir}}$.

At $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$ ($A^* \perp A$): no information about A is used, and

$$\begin{aligned} \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)] &= \text{Cov}(T, f(A)), \\ D(C) &= \text{Var}(T) = 0.25. \end{aligned}$$

The indirect effect is $\tau_{\text{ind}} = \mathbb{E}[f(A) \mid T=1] - \mathbb{E}[f(A) \mid T=0]$, and

$$\text{Cov}(T, f(A)) = 0.25 \tau_{\text{ind}}.$$

Hence at C uniform: $\text{plim } \hat{\tau} = \tau_{\text{dir}} + \tau_{\text{ind}} = \tau_{\text{total}}$.

Numerical Verification (Mediation sign preservation). $\text{sign}(\text{plim } \hat{\tau}_{\text{dir}}) = \text{sign}(\tau_{\text{dir}})$ whenever $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$, where $\tau_{\text{total}} = \tau_{\text{dir}} + \tau_{\text{ind}}$. This result is verified numerically for 100,000 independently sampled column-stochastic matrices C ; a closed-form proof for general K remains open.

Proof. Let $R(C) = \mathbb{E}[\text{Cov}(T, f(A) \mid A^*)]$ and $D(C) = \mathbb{E}[\text{Var}(T \mid A^*)] > 0$.

Endpoints. At $C=I$: $f(A)$ is constant given $A=A^*$, so $R(I) = 0$ and $B(I) = 0$. At $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$: $R = \text{Cov}(T, f(A)) = 0.25 \tau_{\text{ind}}$ and $D = \text{Var}(T) = 0.25$, giving $B = \tau_{\text{ind}}$.

Within-group structure for binary T . Since $T \in \{0, 1\}$, the within- A^* -group bias admits a revealing decomposition. Let $p_j = P(T=1 \mid A^*=j)$. Then $\text{Var}(T \mid A^*=j) = p_j(1-p_j)$ and

$$\begin{aligned} \text{Cov}(T, f(A) \mid A^*=j) &= p_j(1-p_j) [\mathbb{E}[f(A) \mid T=1, A^*=j] \\ &\quad - \mathbb{E}[f(A) \mid T=0, A^*=j]]. \end{aligned}$$

Define $\delta_j = \mathbb{E}[f(A) \mid T=1, A^*=j] - \mathbb{E}[f(A) \mid T=0, A^*=j]$ as the *conditional indirect effect* in the $A^*=j$ subpopulation. Then

$$\begin{aligned} B(C) &= \frac{R(C)}{D(C)} = \frac{\sum_j q_j p_j (1-p_j) \delta_j}{\sum_j q_j p_j (1-p_j)} \\ &= \sum_j w_j \delta_j, \end{aligned}$$

a weighted average of the conditional indirect effects, with weights $w_j = q_j p_j (1-p_j) / \sum_\ell q_\ell p_\ell (1-p_\ell) \geq 0$ summing to 1.

At $C=I$: each $A^*=j$ group has $A=j$ a.s., so $\delta_j = f(j) - f(j) = 0$ and $B = 0$. At C uniform: $\delta_j = \tau_{\text{ind}}$ for all j and $B = \tau_{\text{ind}}$.

Sign preservation. When $\text{Cov}(g(A), f(A)) > 0$ (the indirect effect reinforces the direct effect, i.e., τ_{dir} and τ_{ind} have the same sign), the bias $B(C) \geq 0$ cannot oppose τ_{dir} , so sign is preserved trivially.

When τ_{dir} and τ_{ind} have opposite signs but $|\tau_{\text{ind}}| < |\tau_{\text{dir}}|$ (i.e., $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$), sign reversal requires some $|\delta_j| > |\tau_{\text{dir}}|$ with the opposing sign. Since each δ_j is a difference of weighted averages of f values (with weights $P(A=k \mid T=t, A^*=j)$), and $B(C)$ is a convex combination of δ_j , the bound $|B(C)| \leq \max_j |\delta_j|$ holds. We verify numerically that $\max_j |\delta_j| \leq |\tau_{\text{ind}}|$ for all 100,000 random C matrices in our experiments, confirming sign preservation. $\square$

Verification for paper parameters ($K=3$). With $\tau_{\text{dir}}=0.5$, $\gamma=(0, 0, 0)$, $\delta=(0, 0.8, -0.6)$, $\beta=(0, 1.0, -0.5)$:

$\tau_{\text{ind}} = \mathbb{E}[f(A) \mid T=1] - \mathbb{E}[f(A) \mid T=0] = 0.351$ (computed from the softmax class probabilities). $\tau_{\text{total}} = 0.5 + 0.351 = 0.851 > 0$. Since $\tau_{\text{dir}} > 0$ and $\tau_{\text{total}} > 0$, sign preservation holds for all C .

Mediation monotonicity (proportional effects). Under the interpolation $C(\delta) = (1-\delta)I + \delta C_0$, we prove $|B(C(\delta))|$ is non-decreasing in δ when the mediator's effect on Y is proportional to its informativeness about T : $f(A) = \lambda g(A)$ for some scalar λ , where $g(k) = \mathbb{E}[T \mid A=k]$.

Step 1: Decompose T via the mediator. Define $g(A) = \mathbb{E}[T \mid A]$ and $\varepsilon_T = T - g(A)$. By construction, $\mathbb{E}[\varepsilon_T \mid A] = 0$. By non-differential misclassification ($T \perp A^* \mid A$): $\mathbb{E}[\varepsilon_T \mid A=k, A^*=j] = 0$, so $\mathbb{E}[\varepsilon_T \mid A^*=j] = 0$ and $\text{Cov}(\varepsilon_T, g(A) \mid A^*=j) = 0$ for every j .

Step 2: Reduce the numerator. Since $f = \lambda g$:

$$R(C) = \lambda \mathbb{E}[\text{Var}(g(A) \mid A^*)] = \lambda V_W.$$

Step 3: Constant additive term in denominator.

$$D(C) = \mathbb{E}[\text{Var}(T \mid A^*)] = V_W + \sigma_\varepsilon^2,$$

where $\sigma_\varepsilon^2 \equiv \mathbb{E}[\text{Var}(\varepsilon_T \mid A^*)]$. This is constant (independent of C): $\mathbb{E}[\text{Var}(\varepsilon_T \mid A^*)] = \sum_k p_k \text{Var}(T \mid A=k) = \mathbb{E}[\text{Var}(T \mid A)]$, where the equality uses column-stochasticity: $\sum_j q_j P(A=k \mid A^*=j) = p_k$.

Step 4: Closed form and monotonicity. $B(C) = \lambda V_W / (V_W + \sigma_\varepsilon^2)$, structurally identical to confounding (Proposition 2) with σ_ε^2 playing the role of σ_T^2 . Since $V_B(\delta)$ is convex with $V_B'(0) \leq 0$ (same argument as §A.4, Steps 1–3), $V_W(\delta) = V - V_B(\delta)$

is non-decreasing under the same conditions on C_0 . Since $h(x) = |\lambda|x/(x + \sigma_\varepsilon^2)$ is increasing, $|B(C(\delta))|$ is non-decreasing.

Boundedness of the leakage coefficient. The mediation bias (3) writes the residual indirect effect as $(1 - \lambda(C)) \tau_{\text{ind}}$, which (for $\tau_{\text{ind}} \neq 0$) implicitly defines

$$\lambda(C) := 1 - \frac{B(C)}{\tau_{\text{ind}}}, \quad (20)$$

$$B(C) \equiv \text{plim } \hat{\tau}_{\text{dir}}(A^*) - \tau_{\text{dir}},$$

$$\tau_{\text{ind}} = \frac{\text{Cov}(T, f(A))}{\text{Var}(T)}.$$

Substituting back into (3) gives the convex-combination form $\text{plim } \hat{\tau}_{\text{dir}}(C) = \lambda(C) \tau_{\text{dir}} + (1 - \lambda(C)) \tau_{\text{total}}$, so $\lambda(C)$ is the weight placing the biased plim between τ_{dir} (perfect classification) and τ_{total} (no information).

Proposition 4 (Mediation leakage boundedness). *Let $\lambda(C)$ be defined by (20) with $\tau_{\text{ind}} \neq 0$, and let C be column-stochastic.*

(a) $\lambda(C) \leq 1$ iff $\text{sign}(B(C)) = \text{sign}(\tau_{\text{ind}})$ (the bias points in the direction of the indirect effect; no overcorrection).

(b) $\lambda(C) \geq 0$ iff additionally $|B(C)| \leq |\tau_{\text{ind}}|$ (the bias does not exceed the indirect effect in magnitude).

Equivalently, $\lambda(C) \in [0, 1]$ iff $\text{plim } \hat{\tau}_{\text{dir}}(C) \in [\min(\tau_{\text{dir}}, \tau_{\text{total}}), \max(\tau_{\text{dir}}, \tau_{\text{total}})]$.

Under proportional effects ($f(A) = \lambda_0 g(A)$ with $g(k) = \mathbb{E}[T \mid A=k]$), both (a) and (b) hold for every C , so $\lambda(C) \in [0, 1]$ unconditionally. For general (non-proportional) f , both (a) and (b) can fail because $\text{Cov}(\mathbb{E}[g \mid A^*], \mathbb{E}[f \mid A^*])$ may differ in sign or magnitude from $\text{Cov}(g, f)$ (a covariance-level analogue of Simpson's paradox under A^* aggregation); the numerical verification below quantifies the empirical rate. The endpoints $\lambda(I) = 1$ and $\lambda(C) \rightarrow 0$ in the rank-one (uninformative) limit hold in all regimes.

Sign preservation. When $\lambda(C) \in [0, 1]$, the convex-combination form gives $\text{plim } \hat{\tau}_{\text{dir}}(C) \in [\min(\tau_{\text{dir}}, \tau_{\text{total}}), \max(\tau_{\text{dir}}, \tau_{\text{total}})]$. If additionally $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$, both endpoints share the same sign, so $\text{sign}(\text{plim } \hat{\tau}_{\text{dir}}) = \text{sign}(\tau_{\text{dir}})$ for every C . Under proportional effects, $\lambda(C) \in [0, 1]$ is automatic (above), so sign preservation of $\hat{\tau}_{\text{dir}}$ reduces to the single condition $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$.

Proof. (a),(b). With $\tau_{\text{ind}} \neq 0$, $\lambda(C) \leq 1 \iff B(C)/\tau_{\text{ind}} \geq 0 \iff \text{sign}(B(C)) = \text{sign}(\tau_{\text{ind}})$. Given (a), $\lambda(C) \geq 0 \iff B(C)/\tau_{\text{ind}} \leq 1 \iff |B(C)| \leq |\tau_{\text{ind}}|$.

Proportional case. Step 4 of the mediation derivation gives $B(C) = \lambda_0 V_W(C)/(V_W(C) + \sigma_\varepsilon^2)$; the rank-one limit $V_W \rightarrow V$ yields $\tau_{\text{ind}} = \lambda_0 V/(V + \sigma_\varepsilon^2)$. Both quantities carry $\text{sign}(\lambda_0)$, establishing (a). The ratio

$$\frac{B(C)}{\tau_{\text{ind}}} = \frac{V_W(C) (V + \sigma_\varepsilon^2)}{V (V_W(C) + \sigma_\varepsilon^2)}$$

is monotone increasing in $V_W(C) \in [0, V]$ with range $[0, 1]$ (its derivative in V_W is $\sigma_\varepsilon^2(V + \sigma_\varepsilon^2)/[V(V_W + \sigma_\varepsilon^2)^2] > 0$), establishing (b). Hence $\lambda(C) \in [0, 1]$ unconditionally.

General case. Substituting $\lambda(C)$ from (20) into (3) yields $\text{plim } \hat{\tau}_{\text{dir}}(C) = \lambda(C) \tau_{\text{dir}} + (1 - \lambda(C)) \tau_{\text{total}}$, so $\lambda(C) \in [0, 1]$ iff $\text{plim } \hat{\tau}_{\text{dir}}(C)$ lies in the closed interval with endpoints τ_{dir} and τ_{total} .

Endpoints. At $C=I$, $A^*=A$ a.s., $\mathbb{E}[T | A^*] = g(A)$, and FWL gives $\text{plim } \hat{\tau}_{\text{dir}} = \tau_{\text{dir}}$, hence $B(I) = 0$ and $\lambda(I) = 1$. At any rank-one C (rows identical), $A^* \perp A$ so $\mathbb{E}[T | A^*]$ is constant, the partial regression on A^* has no effect, and $\text{plim } \hat{\tau}_{\text{dir}} = \text{Cov}(Y, T)/\text{Var}(T) = \tau_{\text{total}}$, giving $B = \tau_{\text{ind}}$ and $\lambda = 0$. $\square$

Relation to the between-variance ratio. A natural *diagnostic* closely related to $\lambda(C)$ is the between-variance ratio

$$\lambda_V(C) := \frac{V_B(C)}{V}, \quad (21)$$

where $V_B(C) = \text{Var}(\mathbb{E}[g(A) | A^*])$ and $V = \text{Var}(g(A))$, which satisfies $\lambda_V(C) \in [0, 1]$ unconditionally by the law of total variance applied to $g(A)$. Under proportional effects, λ and λ_V agree at the endpoints and are related by

$$\begin{aligned} 1 - \lambda(C) &= \frac{V_W(C) (V + \sigma_\varepsilon^2)}{V (V_W(C) + \sigma_\varepsilon^2)} \\ &= (1 - \lambda_V(C)) \cdot \frac{V + \sigma_\varepsilon^2}{V_W(C) + \sigma_\varepsilon^2}, \end{aligned}$$

so they coincide in the limit $\sigma_\varepsilon^2 \gg V$. They generally differ for non-proportional f . We use λ throughout the main text because it parameterizes (3) directly; λ_V remains a useful summary of the channel quality of A^* as a proxy for A (independent of the structural form of f).

Numerical verification. For each $K \in \{3, 5, 7\}$ we draw $N=10,000$ configurations $(\pi, g, \sigma_\varepsilon^2, C)$ with $\pi \sim \text{Dirichlet}(\mathbf{1})$, $g \sim \mathcal{N}(0, I_K)$, $\sigma_\varepsilon^2 \sim \text{Uniform}(0.1, 2)$, columns of $C \sim \text{Dirichlet}(\mathbf{1})$, and evaluate $\lambda(C)$ from (20) in two regimes for f :

- **Proportional** ($f = \lambda_0 g$, $\lambda_0 \sim \mathcal{N}(0, 1) \setminus \{0\}$): $\lambda(C) \in [0, 1]$ in all 30,000 trials (pooled observed range $[1.2 \times 10^{-5}, 0.850]$; zero out-of-bounds), confirming the closed-form bound. Sign preservation of $\hat{\tau}_{\text{dir}}$ therefore reduces to $\text{sign}(\tau_{\text{dir}}) = \text{sign}(\tau_{\text{total}})$.
- **Non-proportional** ($f \sim \mathcal{N}(0, I_K)$ drawn independently of g , restricted to $|\tau_{\text{ind}}| \geq 0.05$ to avoid the ill-conditioned regime where $\lambda=1-B/\tau_{\text{ind}}$ diverges): $\lambda(C)$ exceeds $[0, 1]$ in $\sim 15\text{--}17\%$ of trials, illustrating that (a) and (b) are not automatic outside the proportional case.

The diagnostic $\lambda_V(C)$ lies in $[0, 1]$ across all 30,000 trials irrespective of regime, as guaranteed by the law of total variance. Reproduction code and numerical results are included in the supplementary material.

A.6 Proof of Theorem 2 ($K=3$ Symmetric Monotonicity)

Proof. All three parts reduce to a single lemma about the between-group variance $V_B(\delta) = \text{Var}(\mathbb{E}[g(A) | A_\delta^*])$.

Step 1: Structural reduction. The symmetric C_0 decomposes as $C_0 = (1-3c)I + c\mathbf{1}\mathbf{1}^\top$, so

$$\begin{aligned} C(\delta) &= (1-\delta)I + \delta[(1-3c)I + c\mathbf{1}\mathbf{1}^\top] \\ &= (1-3c\delta)I + c\delta\mathbf{1}\mathbf{1}^\top. \end{aligned}$$

Setting $s = c\delta \in [0, c]$, each entry is $C(\delta)_{j\ell} = (1-3s)\mathbf{1}_{j=\ell} + s$, so $C(\delta)$ stays in the one-parameter symmetric family. In particular:

$$\begin{aligned} q_j(s) &= \sum_k C(\delta)_{jk} p_k = (1-3s)p_j + s, \quad (22) \\ \mu_j(s) &\equiv \mathbb{E}[g(A) | A^*=j] = \bar{g} + \frac{(1-3s)p_j h_j}{q_j(s)}, \quad (23) \end{aligned}$$

where $h_j = g_j - \bar{g}$, $\bar{g} = \sum_k p_k g_k$.

Step 2: Closed-form $V_B(s)$. The between-group variance admits a clean factorization:

$$V_B(s) = \sum_j q_j (\mu_j - \bar{g})^2 = (1-3s)^2 \sum_{j=0}^2 \frac{p_j^2 h_j^2}{q_j(s)}. \quad (24)$$

Verification. At $s=0$: $q_j=p_j$, so $V_B(0) = \sum_j p_j h_j^2 = \text{Var}(g)$. At $s=\frac{1}{3}$: $V_B(\frac{1}{3}) = 0$.

Step 3: V_B is strictly decreasing on $[0, \frac{1}{3})$ (pointwise, not just at endpoints). Differentiating (24) via the product rule:

$$\frac{dV_B}{ds} = (1-3s) \sum_{j=0}^2 \frac{p_j^2 h_j^2}{q_j^2} \underbrace{[-3(1-3s)p_j - 1 - 3s]}_{\Lambda_j(s)}. \quad (25)$$

Derivation. Write $V_B = (1-3s)^2 F(s)$ with $F(s) = \sum_j p_j^2 h_j^2 / q_j(s)$. The product rule gives $V_B' = -6(1-3s)F + (1-3s)^2 F'$. Since $q_j'(s) = 1-3p_j$, we have $F'(s) = \sum_j p_j^2 h_j^2 (3p_j-1)/q_j^2$. Combining:

$$\begin{aligned} V_B' &= (1-3s) \left[-6 \sum_j \frac{p_j^2 h_j^2}{q_j} \right. \\ &\quad \left. + (1-3s) \sum_j \frac{p_j^2 h_j^2 (3p_j-1)}{q_j^2} \right] \\ &= (1-3s) \sum_j \frac{p_j^2 h_j^2}{q_j^2} [-6q_j + (1-3s)(3p_j-1)]. \end{aligned}$$

Expanding the bracket:

$$\begin{aligned} &-6q_j + (1-3s)(3p_j-1) \\ &= -6(1-3s)p_j - 6s + 3(1-3s)p_j - (1-3s) \\ &= -3(1-3s)p_j - 1 - 3s = \Lambda_j(s). \end{aligned}$$

Sign analysis. For $s \in [0, \frac{1}{3})$: $(1-3s) > 0$, $p_j \geq 0$, so

$\Lambda_j(s) = -3(1-3s)p_j - (1+3s) < 0$ for all j , since $1+3s > 0$ alone forces $\Lambda_j < 0$. Every term $p_j^2 h_j^2 / q_j^2 \geq 0$, and the prefactor $(1-3s) > 0$. Hence $dV_B/ds \leq 0$ for all $s \in [0, \frac{1}{3})$.

Strict inequality. If $\text{Var}(g) > 0$, at least two categories j satisfy $p_j > 0$ and $h_j \neq 0$, contributing strictly positive terms. Since $\Lambda_j < 0$, the sum is strictly negative: $dV_B/ds < 0$ for all $s \in (0, \frac{1}{3})$.

Step 4: From V_B to $|B|$, confounding (a). Under proportional effects $f = \lambda g$ (Proposition 2):

$$\begin{aligned} B(C(\delta)) &= \frac{\lambda V_W(\delta)}{V_W(\delta) + \sigma_T^2}, \\ V_W(\delta) &= \text{Var}(g) - V_B(c\delta). \end{aligned}$$

Step 3 shows $V_B(s)$ is strictly decreasing in $s = c\delta$, so $V_W(\delta)$ is strictly increasing from 0 to $\text{Var}(g)$. Since $x \mapsto |\lambda| x / (x + \sigma_T^2)$ is strictly increasing for $x \geq 0$:

$$|B(C(\delta))| = |\lambda| \frac{V_W(\delta)}{V_W(\delta) + \sigma_T^2}$$

is strictly non-decreasing.

Step 5: Mediation (b). Under proportional effects $f = \lambda g$ with $g(k) = \mathbb{E}[T | A=k]$, the mediation bias (§A.5) has the identical form $B = \lambda V_W / (V_W + \sigma_\varepsilon^2)$ with $\sigma_\varepsilon^2 = \mathbb{E}[\text{Var}(T | A)]$ replacing σ_T^2 . The same argument yields $|B|$ strictly non-decreasing.

Step 6: Collider (c). Under proportional effects $h = \lambda g$ (Proposition 5):

$$\begin{aligned} B_{\text{col}}(C(\delta)) &= -\frac{\lambda V_B(c\delta)}{\text{Var}(T) - V_B(c\delta)}, \\ |B_{\text{col}}| &= |\lambda| \frac{V_B(c\delta)}{\text{Var}(T) - V_B(c\delta)}. \end{aligned}$$

Since $x \mapsto x/(a-x)$ is strictly increasing for $x \in [0, a)$ and $V_B(s)$ is strictly decreasing (Step 3), $|B_{\text{col}}(C(\delta))|$ is strictly non-increasing: from $|B_{\text{Berkson}}| = |\lambda| V_g / (\text{Var}(T) - V_g)$ at $\delta=0$ to 0 at δ such that $V_B \rightarrow 0$. $\square$

Remark 3 (Comparison with the general- K proof). *The general- K monotonicity proof (§A.4) relies on convexity of $V_B(\delta)$ plus the endpoint condition $V_B'(1) \leq 0$, the latter verified only for $K=2$ analytically and $K \geq 3$ numerically. The $K=3$ symmetric proof is stronger: it establishes $dV_B/ds < 0$ at every interior point via the signed decomposition (25), bypassing both the convexity argument and the endpoint condition entirely. The key enabler is that the symmetric C_0 keeps $C(\delta)$ within the one-parameter family $(1-3s)I + s\mathbf{11}^\top$, yielding the factored form (24) in which the $(1-3s)^2$ prefactor and the q_j^{-1} terms interact to produce a uniformly signed derivative.*

Remark 4 (General (non-proportional) effects). *For non-proportional effects ($f \neq \lambda g$), the confounding bias $B(C) = \mathbb{E}[\text{Cov}(g, f | A^*)]/D(C)$ does not factor through V_W alone, and exact monotonicity of $|B(C(\delta))|$ can fail even for $K=3$ symmetric C_0 . Numerical verification (50,000 random configurations) finds violations in $\sim 10\%$ of cases, with maximum step deviation ~ 0.004 , confirming ε -approximate monotonicity consistent with the general- K results (Table 12).*

Remark 5 (Empirical symmetry of LLM confusion matrices). *Empirical $K=3$ confusion matrices from LLM stance detection deviate substantially from symmetry (median relative Frobenius distance $d_{\text{rel}}=0.33$; range 0.30–0.38 across 10 VAST configurations). Theorem 2 thus serves as a theoretical anchor establishing that exact monotonicity is achievable at $K=3$ under the strongest structural*

assumption; practical coverage for asymmetric matrices relies on the numerical evidence of Conjecture 1 ($\varepsilon < 0.02$ across 545,945 adversarial curves including non-symmetric structures).

A.7 Collider Topology ($A = \text{Collider}$)

Data-generating process. $T \rightarrow A \leftarrow Y, T \rightarrow Y$. The treatment effect is unconfounded:

$$\begin{aligned} T &\sim \mathcal{N}(0, 1), \quad Y = \tau T + \varepsilon_Y, \\ \varepsilon_Y &\sim \mathcal{N}(0, 1), \quad \varepsilon_Y \perp T, \\ A | T, Y &\sim \text{Categorical}(\text{softmax}(\delta_T T + \delta_Y Y)). \end{aligned}$$

The collider A is a common effect of T and Y . The no-control estimator (regressing Y on T alone) is consistent: $\text{plim } \hat{\tau}_{\text{nc}} = \tau$. The researcher erroneously controls for A^* , inducing Berkson's bias.

Bias decomposition. Unlike confounding, ε_Y is not independent of A (since A depends on $Y = \tau T + \varepsilon_Y$), so Lemma 1 does not apply directly. Instead, we derive the bias from the FWL identity (7). Writing $\tilde{Y} = \tau \tilde{T} + (\varepsilon_Y - \mathbb{E}[\varepsilon_Y | A^*])$ and noting that $\mathbb{E}[\tilde{T} \varepsilon_Y] = \mathbb{E}[\text{Cov}(T, \varepsilon_Y | A^*)]$ (since $T \perp \varepsilon_Y$ unconditionally):

$$\begin{aligned} \text{plim } \hat{\tau}(C) &= \tau + B(C), \\ B(C) &= \frac{\mathbb{E}[\text{Cov}(T, \varepsilon_Y | A^*)]}{\mathbb{E}[\text{Var}(T | A^*)]}. \end{aligned} \quad (26)$$

Under non-differential misclassification ($A^* \perp (T, \varepsilon_Y) | A$), define $g(k) = \mathbb{E}[T | A=k]$ and $h(k) = \mathbb{E}[\varepsilon_Y | A=k]$. By the tower property, $\mathbb{E}[T | A^*] = \mathbb{E}[g(A) | A^*]$ and $\mathbb{E}[\varepsilon_Y | A^*] = \mathbb{E}[h(A) | A^*]$. Since $\text{Cov}(T, \varepsilon_Y) = 0$, the law of total covariance gives:

$$B(C) = \frac{-\text{Cov}(\mathbb{E}[g(A) | A^*], \mathbb{E}[h(A) | A^*])}{\mathbb{E}[\text{Var}(T | A^*)]}.$$

Endpoint analysis. At $C = I$ ($A^* = A$): $B(I) = -\text{Cov}(g(A), h(A)) / \mathbb{E}[\text{Var}(T | A)] = B_{\text{Berkson}}$, the full Berkson's bias from conditioning on the collider.

At $C = \frac{1}{K} \mathbf{11}^\top$ ($A^* \perp A$): $\mathbb{E}[g(A) | A^*]$ and $\mathbb{E}[h(A) | A^*]$ are constant, so $B = 0$ and $\text{plim } \hat{\tau} = \tau$.

This is the structural reverse of confounding: misclassification attenuates the collider-induced bias toward zero, rather than growing bias from zero toward B_{OY} .

Proposition 5 (Collider sign preservation — proportional case). *If $h(A) = \lambda g(A)$ for some constant λ , then $|B(C)| \leq |B(I)|$ for all column-stochastic C , and $B(C)$ has the same sign as $B(I)$. Consequently, $\text{sign}(\text{plim } \hat{\tau}) = \text{sign}(\tau)$ whenever $\text{sign}(\tau) = \text{sign}(\tau + B_{\text{Berkson}})$.*

Proof. When $h = \lambda g$, the numerator becomes $-\lambda \text{Var}(\mathbb{E}[g(A) | A^*])$ and the denominator is $\text{Var}(T) - \text{Var}(\mathbb{E}[g(A) | A^*])$. Let $V^* = \text{Var}(\mathbb{E}[g(A) | A^*]) \in [0, V_g]$ where $V_g = \text{Var}(g(A))$. Then

$$|B(C)| = |\lambda| \frac{V^*}{\text{Var}(T) - V^*}.$$

The function $x \mapsto x/(a - x)$ is increasing on $[0, a)$. With $a = \text{Var}(T) > V_g \geq V^*$: $|B(C)| \leq |\lambda| V_g / (\text{Var}(T) - V_g) = |B(I)|$. At $V^* = 0$ (uniform C): $B = 0$. At $V^* = V_g$ ($C = I$): $B = B(I)$. $\square$

Proposition 6 (Collider sign preservation — general case). *Define*

$$\bar{B}_{\text{col}} = \frac{\sqrt{\text{Var}(g(A)) \text{Var}(h(A))}}{\mathbb{E}[\text{Var}(T | A)]}. \quad (27)$$

If $|\tau| > \bar{B}_{\text{col}}$, then $\text{sign}(\text{plim } \hat{\tau}(C)) = \text{sign}(\tau)$ for all column-stochastic C .

Proof. Step 1: Bound the numerator. By Cauchy-Schwarz: $|\text{Cov}(\mathbb{E}[g | A^*], \mathbb{E}[h | A^*])| \leq \sqrt{\text{Var}(\mathbb{E}[g | A^*]) \text{Var}(\mathbb{E}[h | A^*])}$. Since conditional expectation reduces variance: $\text{Var}(\mathbb{E}[g | A^*]) \leq \text{Var}(g(A))$ and likewise for h .

Step 2: Bound the denominator. $\mathbb{E}[\text{Var}(T | A^*)] = \text{Var}(T) - \text{Var}(\mathbb{E}[T | A^*]) \geq \text{Var}(T) - \text{Var}(g(A)) = \mathbb{E}[\text{Var}(T | A)]$.

Step 3: Conclude. $|B(C)| \leq \sqrt{\text{Var}(g) \text{Var}(h)} / \mathbb{E}[\text{Var}(T | A)] = \bar{B}_{\text{col}}$. When $|\tau| > \bar{B}_{\text{col}}$, the bias cannot overwhelm τ . In practice, $\text{Var}(g(A))$ and $\text{Var}(h(A))$ can be estimated from observable data by bootstrapping regression coefficients from A^* -stratified subsamples, yielding conservative upper bounds on $\bar{B}_{\text{col}}$. $\square$

Verification for paper parameters ($K=3$). With $\tau=1.0$, $\delta_T=(0, 0.5, -0.3)$, $\delta_Y=(0, 0.8, 0.6)$, from 5×10^6 -sample Monte Carlo:

$$\begin{aligned} \text{Var}(g(A)) &= 0.187, \quad \text{Var}(h(A)) = 0.076, \\ \mathbb{E}[\text{Var}(T | A)] &= 0.813, \\ \bar{B}_{\text{col}} &= \sqrt{0.187 \times 0.076} / 0.813 \\ &= 0.147. \end{aligned}$$

Since $|\tau|=1.0 \gg \bar{B}_{\text{col}}=0.147$, Proposition 6 guarantees sign preservation. The Berkson bias $B_{\text{Berkson}}=-0.114$, giving $\tau + B_{\text{Berkson}}=0.886>0$. Verified across 100,000 random column-stochastic C : $\tau_{\text{sign flip}}=0\%$, and $|B(C)| \leq |B_{\text{Berkson}}|$ in all cases.

A.8 Front-door Topology ($A = \text{Mediator}$ under Confounding)

Data-generating process. $T \rightarrow A \rightarrow Y, U \rightarrow T, U \rightarrow Y$. An unobserved confounder U affects both treatment and outcome:

$$\begin{aligned} U &\sim \mathcal{N}(0, 1), \quad T = \alpha_U U + \eta, \\ \eta &\sim \mathcal{N}(0, 1), \quad \eta \perp U, \\ A | T &\sim \text{Categorical}(\text{softmax}(\gamma + \delta T)), \\ Y &= f(A) + \lambda U + \varepsilon_Y, \quad \varepsilon_Y \perp (T, A, U). \end{aligned}$$

There is no direct $T \rightarrow Y$ edge; the causal effect of T on Y operates entirely through A . The researcher controls for A^* in $Y \sim 1 + T + D_1^* + \dots + D_{K-1}^*$.

Bias decomposition. Since the direct effect is zero, the FWL coefficient on T captures only bias:

$$\text{plim } \hat{\tau}(C) = \frac{\mathbb{E}[\text{Cov}(T, Y | A^*)]}{\mathbb{E}[\text{Var}(T | A^*)]}.$$

Substituting $Y = f(A) + \lambda U + \varepsilon_Y$ and using $\varepsilon_Y \perp (T, A, A^*)$:

$$\text{plim } \hat{\tau}(C) = \underbrace{\frac{\lambda \alpha_U}{\text{Var}(T)}}_{\text{conf.}} + \underbrace{\frac{\mathbb{E}[\text{Cov}(T, f(A) | A^*)]}{D(C)}}_{B_{\text{med}}(C)}, \quad (28)$$

where $D(C) = \mathbb{E}[\text{Var}(T | A^*)]$.

Derivation of (28). Expand the numerator as $\mathbb{E}[\text{Cov}(T, f(A) | A^*)] + \lambda \mathbb{E}[\text{Cov}(T, U | A^*)]$. For the confounding term: since $T = \alpha_U U + \eta$ with (U, η) jointly Gaussian, the regression residual $\xi = U - \frac{\alpha_U}{\text{Var}(T)} T$ satisfies $\xi \perp T$. Because A is a function of T (plus independent softmax noise) and A^* is a function of A (plus independent misclassification noise), $\xi \perp (T, A, A^*)$. Therefore

$$\text{Cov}(T, U | A^*=j) = \frac{\alpha_U}{\text{Var}(T)} \text{Var}(T | A^*=j) \quad \forall j,$$

and $\lambda \mathbb{E}[\text{Cov}(T, U | A^*)] = \frac{\lambda \alpha_U}{\text{Var}(T)} D(C)$, yielding the constant first term after division by $D(C)$. $\square$

This decomposition reveals that misclassification does not affect the confounding component. The confusion matrix C modulates only the mediation-type leakage $B_{\text{med}}(C)$, which has the same structure as in §A.5.

Endpoint analysis. At $C = I$: $B_{\text{med}}(I) = 0$ (controlling A blocks the $T \rightarrow A \rightarrow Y$ path), and $\text{plim } \hat{\tau}(I) = \lambda \alpha_U / \text{Var}(T)$ (pure confounding residual).

At $C = \frac{1}{K} \mathbf{1}\mathbf{1}^\top$: $B_{\text{med}} = \text{Cov}(T, f(A)) / \text{Var}(T)$, and $\text{plim } \hat{\tau}_{\text{nc}} = [\text{Cov}(T, f(A)) + \lambda \alpha_U] / \text{Var}(T)$, the naive regression coefficient with no controls.

Numerical Verification (Front-door sign preservation). Since $B_{\text{med}}(C)$ has the same structure as the mediation bias (§A.5), interpolating between 0 at $C=I$ and $\text{Cov}(T, f(A)) / \text{Var}(T)$ at C uniform. Sign preservation holds whenever the two endpoints have the same sign:

$$\text{sign}(\lambda \alpha_U) = \text{sign}(\lambda \alpha_U + \text{Cov}(T, f(A))).$$

This result is verified numerically for 100,000 independently sampled column-stochastic matrices C .

Verification for paper parameters ($K=3$). With $\alpha_U=1.0$, $\lambda=1.0$, $\beta=(0, 1.0, -0.5)$, $\gamma=(0, 0.5, -0.3)$, $\delta=(0, 1.0, -0.8)$:

$$\begin{aligned} \lambda \alpha_U / \text{Var}(T) &= 1.0 \times 1.0 / 2.0 = 0.500, \\ \text{Cov}(T, f(A)) / \text{Var}(T) &= 0.594 / 2.0 = 0.297. \end{aligned}$$

Both components are positive, so $\text{plim } \hat{\tau}(C) \in [0.500, 0.797] > 0$ for all C .

A.9 Exposure Topology: Sign Reversal Counterexample

Proposition 7 (Exposure sign reversal). *For $K=3$ with A as the treatment variable, there exist a marginal distribution $\mathbf{p}$, treatment effects β , and a column-stochastic C with $C_{kk} > 0.5$ for all k , such that $\text{sign}(\text{plim } \hat{\beta}_1) \neq \text{sign}(\beta_1)$.*

Construction. Consider the exposure data-generating process (DGP) $Y = \beta_1 D_1 + \beta_2 D_2 + \varepsilon_Y$ with:

- Uniform prior: $\mathbf{p} = (1/3, 1/3, 1/3)$.
- True effects: $\beta_1 = 1, \beta_2 = -10$.

Choose the confusion matrix

$$C = \begin{pmatrix} 0.60 & 0.10 & 0.10 \\ 0.10 & 0.60 & 0.39 \\ 0.30 & 0.30 & 0.51 \end{pmatrix},$$

with $C_{kk} \in \{0.60, 0.60, 0.51\}$, all > 0.5 . Column sums are 1 (verifiable by inspection).

The OLS estimator regresses Y on $(1, D_1^*, D_2^*)$. The plim of $\hat{\beta}_j$ equals $\mathbb{E}[Y \mid A^*=j] - \mathbb{E}[Y \mid A^*=0]$, where

$$\begin{aligned} \mathbb{E}[Y \mid A^*=j] &= \sum_{k=0}^2 P(A=k \mid A^*=j) \beta_k \\ &= \sum_{k=1}^2 \frac{C_{jk} p_k}{q_j} \beta_k. \end{aligned}$$

Marginal distribution of A^ :*

$$\begin{aligned} q_0 &= \frac{1}{3}(0.60 + 0.10 + 0.10) = 0.267, \\ q_1 &= \frac{1}{3}(0.10 + 0.60 + 0.39) = 0.363, \\ q_2 &= \frac{1}{3}(0.30 + 0.30 + 0.51) = 0.370. \end{aligned}$$

Conditional means:

$$\begin{aligned} \mathbb{E}[Y \mid A^*=0] &= \frac{0.10 \cdot \frac{1}{3} \cdot 1 + 0.10 \cdot \frac{1}{3} \cdot (-10)}{0.267} \\ &= \frac{0.0333 - 0.3333}{0.267} = -1.125, \\ \mathbb{E}[Y \mid A^*=1] &= \frac{0.60 \cdot \frac{1}{3} \cdot 1 + 0.39 \cdot \frac{1}{3} \cdot (-10)}{0.363} \\ &= \frac{0.200 - 1.300}{0.363} = -3.028, \\ \mathbb{E}[Y \mid A^*=2] &= \frac{0.30 \cdot \frac{1}{3} \cdot 1 + 0.51 \cdot \frac{1}{3} \cdot (-10)}{0.370} \\ &= \frac{0.100 - 1.700}{0.370} = -4.324. \end{aligned}$$

Estimated coefficient:

$$\begin{aligned} \text{plim } \hat{\beta}_1 &= \mathbb{E}[Y \mid A^*=1] - \mathbb{E}[Y \mid A^*=0] \\ &= -3.028 - (-1.125) = -1.903. \end{aligned}$$

The true $\beta_1 = +1 > 0$ but $\text{plim } \hat{\beta}_1 = -1.903 < 0$: **sign reversal**.

Mechanism. The large off-diagonal entry $C_{12} = 0.39$ (probability of classifying true category 2 as $A^*=1$) routes substantial probability mass from category 2 ($\beta_2 = -10$) into the $A^*=1$ bin. Despite $C_{11} = 0.60 > 0.5$, the contamination from the strongly negative β_2 overwhelms the positive β_1 , reversing the sign of the estimated effect for category 1. This failure mode is structurally impossible at $K=2$: a 2×2 column-stochastic matrix with $C_{kk} > 0.5$ has the form $C = \begin{pmatrix} a & 1-b \\ 1-a & b \end{pmatrix}$ with $a, b > 0.5$, and the single coefficient $\hat{\beta}_1$ is a

positive-weighted average of β_1 , preserving sign. For $K \geq 3$, the additional degrees of freedom in C permit cross-contamination between categories with opposite-sign effects. $\square$

A.10 Summary of Conditions

“Proved” = algebraic proof via Cauchy–Schwarz (confounding) or Berkson attenuation (collider). “Empirical” = verified across 100,000 random column-stochastic matrices with zero violations; analytic proof for general K remains open. All conditions are on DGP parameters (not on C): once satisfied, sign preservation holds for all column-stochastic C with $C_{kk} > 0$.

Topology	Sufficient condition	Status	Paper DGP
Confounding	$\text{sign}(\tau) = \text{sign}(\tau + B_{OV})$	Proved	✓ (2.0 vs 2.3)
Collider	$ \tau > B_{col}$ (Prop. 6)	Proved	✓ (1.0 vs 0.15)
Mediation	$\text{sign}(\tau_{dir}) = \text{sign}(\tau_{total})$	Empirical	✓ (0.5 vs 0.85)
M-bias	Bias is attenuating (toward τ)	Empirical	✓
Front-door	$\text{sign}(\lambda_{\alpha_U}) = \text{sign}(\lambda_{\alpha_U} + \text{Cov}(T, f))$	Empirical	✓ (0.50 vs 0.80)
Exposure IV	Not guaranteed for $K \geq 3$ Not guaranteed (denominator instability)	Counterex.	See §A.9

Table 3: Sufficient conditions for sign preservation under non-differential misclassification.

Structural argument for M-bias topology. *M-bias:* The M-structure creates a backdoor path through two latent common causes. Misclassification of A weakens the statistical association along this backdoor path, reducing the spurious correlation rather than amplifying it. The bias therefore interpolates between the full M-bias (perfect classification, $\delta=0$) and zero (no classification information, $\delta \rightarrow \delta_{\max}$), monotonically non-increasing in δ .

M-bias sign-flip ASV $= \infty$. For the four covariate topologies (confounding, mediation, collider, M-bias), FWL invariance (Appendix A.2) implies that the coefficient on T depends on A^* only through the residualized relationship $\tilde{T} = T - \mathbb{E}[T \mid A^*]$. For M-bias, the spurious path through A is the sole source of bias, and misclassification attenuates this path for all column-stochastic C . Since $\text{bias}(C(\delta))$ is monotonically non-increasing in δ and starts at a value with the same sign as τ , no $\delta \in [0, \delta_{\max}]$ can produce a sign reversal. Therefore $\delta_{\text{sign}}^* = \infty$, which is a direct consequence of FWL invariance combined with the protective monotonicity of M-bias misclassification.

Remark. The sign preservation conditions for confounding, mediation, collider, and front-door

are conditions on the DGP parameters (the relative strength of the causal effect versus the confounding, indirect, or collider-induced bias), not on the confusion matrix C . Once the condition is satisfied, sign preservation holds for all column-stochastic C with $C_{kk} > 0$; no diagonal-dominance restriction is needed. Conversely, if the condition fails (i.e., the omitted-variable or total-effect estimate has the opposite sign from τ), then there exist confusion matrices that produce sign reversal. For the DGP parameters used throughout our experiments, the conditions are comfortably satisfied. The sign-alignment condition is violated when the confounder has opposing effects on treatment and outcome (e.g., $\gamma_A > 0$ on T but $\gamma_A < 0$ on Y), producing $\text{sign}(\tau) \neq \text{sign}(\tau_{OV})$. In applied settings, such opposing effects are uncommon when the analyst adjusts for A precisely because they believe it confounds the T - Y relationship in a consistent direction. When the condition fails, sign reversal becomes possible under sufficient misclassification, though this represents a pathological DGP where adjustment itself may be questionable.

A.11 Adversarial Monotonicity Verification

We test Conjecture 1 under four adversarial confusion matrix structures designed to challenge envelope monotonicity: near-permutation ($C_0 = (1-\epsilon)P + \epsilon \cdot \mathbf{1}\mathbf{1}^\top/K$, P a random permutation), block-diagonal (within-block confusion only), sparse (two nonzero entries per column), and near-uniform ($C_{kk} \rightarrow 1/K$). For each structure, we sweep $\delta \in [0, 1]$ along $C(\delta) = (1-\delta)I + \delta C_0$ and measure the maximum single-step monotonicity violation ε_{abs} .

$N=1,000$ matrices per cell, 545,945 total curves were tested.

Table 4: Adversarial monotonicity test across $K=3, 5, 7$ and seven topologies.

Structure	Non-monotone %	Max ε_{abs}
Near-permutation	35.4%	0.016
Block-diagonal	0%	0
Sparse	2.3%	0.003
Near-uniform	0%	0
All structures	11.9%	0.016

The four families span the extremes of the column-stochastic polytope: diagonal-dominant (near-identity), maximally reordered (near-permutation), degenerate (rank-deficient

sparse), and block-structured heterogeneity (block-diagonal). The global maximum violation $\varepsilon_{\text{abs}} = 0.016 < 0.02$ confirms that Conjecture 1 holds under adversarial matrix structures. Near-permutation matrices with alternating-sign β coefficients represent the most challenging case; even here, violations remain below the conjectured bound. Block-diagonal and near-uniform structures show exact monotonicity.

B Supplementary Tables, Figures, and Algorithm

B.1 Binarization Collapse

Collapsing K -class annotations to binary before analysis discards multi-class error structure. We quantify “binarization collapse” as $\Delta = |\text{bias}_K - \text{bias}_{K \rightarrow 2}|$.

For $K=3$, binarization introduces additional bias varying $>145\times$ across topologies: from $\bar{\Delta}=0.006$ (M-bias) to $\bar{\Delta}=0.926$ (IV). Exposure and IV also introduce sign flips during binarization (14% and 17.5%), a phenomenon absent in the other topologies (Figure 6).

Magnitude ASV is reported at $\tau=0.5$ with a 10% threshold. Power is evaluated at $\tau=0.05$ significance filter. “—” indicates not applicable.

Table 5: ASV type applicability and sensitivity across DAG topologies.

Topology	Mag. ASV	Signif.	Sign-flip	Rank
IV	0.101–0.607	0%	1.1–6.7	1 (fragile)
Front-door	0.064–0.387	100%	—	2
Confounding	0.081–0.487	100%	—	3
Mediation	0.076–0.454	0%	—	4
Collider	0.338–2.028	0%	—	5
M-bias	∞	96–100%	—	6 (immune)
Exposure	—	—	83.4%*	—

*Fraction of C_0 with $\tau_{\text{sign_flip}} < 1$. Exposure uses sign-flip ASV exclusively.

Results are averaged across 6 error structures with $N_{\text{sim}}=1,000$. Positive values indicate reduced RMSE; negative values indicate increased RMSE. Naive/MLA show 0.0% for FWL-invariant topologies because C^{-1} cannot change the T coefficient. **Bold**: best method per topology; underline: best gold-free method.




Figure 7: ASV magnitude ranges across topologies for HumAID ($K=10$). Left: magnitude ASV (δ^*) range across 8 LLM configurations; lower values indicate greater fragility. “Safe” topologies never exceeded the 10% bias threshold. Right: percentage of configurations flagged as ASV-unsafe.

Table 6: RMSE reduction (%) by correction method and DAG topology.

Topology	Naive	MLA	DSL	SIMEX	PPI++
Confounding	0.0	0.0	48.0	35.1	83.2
Mediation	0.0	0.0	52.9	39.1	54.1
Collider	0.0	0.0	-92.6	31.8	50.5
M-bias	0.0	0.0	-267	0.1	-88.0
Front-door	-52.7	-111.5	-29.0	-62.4	74.1
Exposure	-44.8	-89.0	-59.5	-9.8	75.5
IV	-40.4	-58.7	-206.1	59.4	79.4
Grand mean	-19.7	-37.0	-79.0	13.3	47.0

Error structures used in correction experiments. The six confusion matrix structures used for Table 6 are: (1) *Diagonal-dominant*: $C_{kk}=0.7$, off-diagonal uniform; (2) *Moderate noise*: $C_{kk}=0.5$, off-diagonal uniform; (3) *Near-permutation*: $C_{kk}=0.4$, largest off-diagonal 0.5; (4) *Asymmetric*: C_{kk} varies (0.8, 0.5, 0.3); (5) *Block confusion*: two classes confused ($C_{12}=C_{21}=0.4$), third clean; (6) *Heavy noise*: $C_{kk}=0.4$, off-diagonal uniform. All are column-stochastic with $K=3$. Results are averaged across all six structures.

This topology-based classification is robust to ε -approximate deviations; practitioners computing quantitative ASV thresholds should note the ε -approximate bounds (Conjecture 1, Section 3.3). Assumes linear structural equations; FWL-based routing may not apply under nonlinear models (see Appendix B.7). Thresholds are calibrated for $K=3$ linear SEM; recalibrate for different K using topology-specific bias functions. The full decision

procedure is given in Algorithm 1 (Section 4.5).

B.2 Proof of FWL Invariance

For the four topologies where A enters as a covariate (confounding, mediation, collider, M-bias), the regression is $Y = \tau T + \gamma^\top \mathbf{D}^* + \varepsilon$, where $\mathbf{D}^*$ is the $N \times (K-1)$ dummy matrix for A^* . Any matrix-based correction replaces $\mathbf{D}^*$ with $\tilde{\mathbf{D}} = \mathbf{D}^* M$ for some invertible $(K-1) \times (K-1)$ matrix M (e.g., $M = C^{-1}$ for plug-in, or $M = (C^\top C)^{-1} C^\top$ for maximum likelihood estimation (MLE)). By the Frisch–Waugh–Lovell theorem (Lovell, 1963), $\hat{\tau}$ depends only on the residual of T projected off $\text{col}(\mathbf{D}^*)$. Since $\text{col}(\mathbf{D}^* M) = \text{col}(\mathbf{D}^*)$ for invertible M , the projection is unchanged and $\hat{\tau}(\tilde{\mathbf{D}}) = \hat{\tau}(\mathbf{D}^*)$ exactly. SIMEX adds noise $\mathbf{D}_{\text{sim}}^* = \mathbf{D}^* + \lambda \mathbf{E}$; when $\mathbf{E}$ is generated from the same column space, the extrapolation to $\lambda = -1$ converges to a column-space-equivalent transformation, yielding 0% bias reduction empirically.

FWL breakdown under logistic regression.

The column-space argument applies only to estimators whose treatment coefficient is a function of T residualized off $\text{col}(\mathbf{D}^*)$, a property of linear projection. The logistic MLE solves $\sum_i (Y_i - \sigma(\mathbf{x}_i^\top \beta)) \mathbf{x}_i = \mathbf{0}$; the implicit weighting by $\sigma'(\cdot)$ depends on the actual values of the design columns, not only on their span, so replacing $\mathbf{D}^*$ by $\mathbf{D}^* M$ generally changes the fitted $\hat{\tau}$. We verify this on the confounding topology ($K=3$, $\tau=0.5$, $N=10,000 \times 500$ MC reps) under both linear ($Y = \tau T + \mathbf{D}\beta + \varepsilon$) and logis-

tic ($Y \sim \text{Bern}(\sigma(\tau T + \mathbf{D}\beta))$) DGPs, comparing the uncorrected estimator (fit on $\mathbf{D}^*$) against Naive plug-in (fit on per-sample $(C^{-1}\mathbf{e}_{A^*})_{1:K}$). Per-rep estimates are paired (same A^* , same outcome).

$|\Delta|_{\max} = \max_m |\hat{\tau}_{\text{corr}}^{(m)} - \hat{\tau}_{\text{un}}^{(m)}|$ over 500 MC reps; bias reduction = $1 - |\hat{\tau}_{\text{corr}} - \tau_*|/|\hat{\tau}_{\text{un}} - \tau_*|$. Confounding DAG, $K=3$, $N=10,000$.

Table 7: FWL invariance under OLS vs. breakdown under logistic regression.

CM (VAST K=3)	OLS		Logistic	
	$ \Delta _{\max}$	bias red.	$ \Delta _{\max}$	bias red.
deepseek_zero-shot	2.3×10^{-15}	+0.00%	8.1×10^{-4}	-0.09%
llama4_few-shot-3	5.2×10^{-15}	+0.00%	6.8×10^{-4}	-0.07%
llama4_few-shot-5	4.3×10^{-15}	+0.00%	6.8×10^{-4}	-0.06%
llama4_zero-shot	1.6×10^{-15}	+0.00%	5.9×10^{-4}	-0.05%
qwen2.5-7b_few-shot-3	7.2×10^{-15}	-0.00%	6.8×10^{-4}	-0.01%
qwen2.5-7b_few-shot-5	8.2×10^{-15}	+0.00%	5.6×10^{-4}	+0.05%
Mean	—	0.0%	—	-0.04%

Two findings: (i) Under OLS, $\hat{\tau}_{\text{corr}}^{(m)} = \hat{\tau}_{\text{un}}^{(m)}$ to machine precision in every replication, confirming Proposition B.2 as an exact algebraic identity. (ii) Under logistic regression, per-replication estimates differ by $|\Delta|_{\max} \sim 10^{-3}$, thirteen orders of magnitude larger than the OLS gap, so the invariance fails formally. However, the directional effect on bias is negligible: mean bias reduction is -0.04% across the six confusion matrices (range $[-0.09\%, +0.05\%]$, not even consistent in sign). Matrix-based corrections remain practically useless under logistic regression, not by an exact invariance but because the deviation is too small to systematically offset the misclassification-induced bias (which itself is ~ 0.10 in absolute units, two orders of magnitude larger than $|\Delta|$). The routing rule of Algorithm 1 (skip Naive/MLA for these four topologies) is therefore conservative-but-correct under logistic regression as well. Reproduction code and numerical results are included in the supplementary material.

B.3 Applied Illustration: ClinicalBERT Case Study

To demonstrate the framework’s end-to-end utility, we apply it retrospectively to Lee and Wood-Doughty (2024), who use ClinicalBERT to classify patient smoking status ($K=4$: never, past, current, unknown) from clinical notes and control for this predicted confounder when estimating the causal effect of echocardiography on mortality in 4,735 MIMIC-III sepsis patients. Our taxonomy identifies this as a confounding topology: the annotated

variable enters as a confounder between treatment and outcome. Three conclusions follow directly. First, the zero sign-flip guarantee for confounding (Table 1) ensures that the estimated protective effect (OR = 0.89) cannot reverse direction due to annotation error, a reassurance unavailable without topology-specific analysis. Second, the FWL invariance result implies that any correction reducing to C^{-1} applied to the dummy-coded confounder achieves exactly 0% bias reduction for the treatment coefficient; naive plug-in is provably futile here. Third, our correction guidance ranks PPI++ (83.2% RMSE reduction for confounding) above MC-SIMEX (35.1%) when gold labels are available, a concrete recommendation the original study lacked. This example illustrates the intended workflow: DAG identification $\rightarrow$ topology-specific guarantees $\rightarrow$ correction method selection.

B.4 IV Wald Estimator: Diagonal-Dominance Subspaces

Section 4.2 reports a 52.6% Wald sign-flip rate over the full Dirichlet(1, 1, 1) prior. Because realistic LLM confusion matrices are diagonal-dominant, we replicate the exposure-DAG stratification for the IV topology, restricting the sample space to $\{C : \min_k C_{kk} \geq d\}$ for $d \in \{0, 0.5, 0.7, 0.9\}$ (100,000 random matrices per threshold).

Rates are computed over 100,000 Dirichlet-truncated random C per threshold, with $K=3$ and $\tau_{\text{true}}=1.41$.

Table 8: IV Wald estimator bias by minimum-diagonal subspace.

Subspace	Violation	Sign Flip	Amplif.	Worst
Full	100.0%	51.7%	48.3%	$10^5 \times$
$d \geq 0.5$	100.0%	0.0%	100.0%	$28.6 \times$
$d \geq 0.7$	100.0%	0.0%	100.0%	$2.4 \times$
$d \geq 0.9$	100.0%	0.0%	100.0%	$1.2 \times$

The contrast with exposure is instructive: at $d \geq 0.7$, exposure violation drops to 13.7% (sign reversal 2.1%, amplification 11.7%), while IV maintains 100% violation with pure amplification. IV sign reversal is entirely driven by pathological (near-random) confusion matrices where the Wald denominator crosses zero; diagonal dominance reduces this crossing to near-zero under empirical accuracy conditions but cannot prevent amplification. Even at $d \geq 0.9$, the Wald estimator inflates $|\tau|$ by up to $1.24 \times$ with probability 1.

B.5 Non-Differential Misclassification Verification

The non-differential assumption $P(A^*=j \mid A=k, Y, \mathbf{Z}) = P(A^*=j \mid A=k)$ holds by design in standard LLM annotation pipelines: the model receives only source text as input and classifies without observing the outcome Y or covariates $\mathbf{Z}$. The annotation API call is structurally independent of downstream variables, satisfying the conditional independence requirement.

To provide empirical support beyond the design argument, we examine topic-stratified confusion matrices using the VAST dataset ($K=3$, stance detection), where the `topic` field (e.g., “gun control,” “college”) serves as a proxy for latent factors correlated with Y . Across 10 model configurations (3 LLMs $\times$ 2–4 prompt types) and 19 qualifying topics (minimum 10 samples per topic), we compute per-topic confusion matrices $C_{\text{topic}}(t) = P(A^*=j \mid A=k, \text{topic}=t)$ and compare against the overall $C = P(A^*=j \mid A=k)$.

Per-topic sample sizes range from 10 to 42, making individual topic CMs inherently noisy. A bootstrap permutation test ($B=10,000$) under the non-differential null shows that the expected grand maximum deviation $\max_{t,j,k} |C_{\text{topic}}(j, k) - C(j, k)|$ at the 95th percentile ranges from 0.90 to 0.98 across model configurations, driven by the smallest topics. For the 5 largest topics ($n \geq 28$), expected per-topic max deviation at the 95th percentile is 0.38–0.49, consistent with finite-sample variation rather than systematic topic-dependence. The observed per-topic variation is fully consistent with the non-differential null.

This proxy-based test is necessarily limited: topic is not the actual outcome Y , and small per-topic samples reduce statistical power. The design argument (LLM input excludes Y) remains the primary justification; the stress test in Section B.6 below quantifies bias degradation under hypothetical violations.

B.6 Non-Differential Assumption Stress Test

Our main analysis assumes non-differential misclassification: the confusion matrix C does not depend on the outcome Y . To assess sensitivity to violations of this assumption, we conduct a stress test on the exposure topology ($K=3$, $\min_k C_{kk} \geq 0.7$, 10,000 random matrices per ε). For each base matrix C_0 , observations with $Y > \text{median}(Y)$ are misclassified using a per-

turbed matrix $\tilde{C} = \text{renormalize}(C_0 + \varepsilon \cdot \Delta)$ where $\Delta_{jk} \sim \text{Uniform}(-1, 1)$ for $j \neq k$, while observations with $Y \leq \text{median}(Y)$ retain C_0 .

Results use the exposure DAG with $K=3$, $\min_k C_{kk} \geq 0.7$, and $N=50,000$ MC samples per trial.

Table 9: Effect of differential misclassification on bias violation rates.

ε	Sign-Flip (%)	Amplification (%)	Any Violation (%)
0.00	2.0	11.5	13.3
0.01	2.4	12.4	14.7
0.05	2.8	13.0	15.7
0.10	4.5	16.7	20.9
0.20	9.7	23.7	32.6
0.30	15.5	28.7	42.8
0.50	22.9	34.9	54.6

The $\varepsilon=0$ baseline (13.3% violation) closely matches the analytical result (13.7%) from Section 4.2. Small departures ($\varepsilon \leq 0.05$) change rates by fewer than 3 percentage points. At $\varepsilon=0.10$, the violation rate increases by 7.6 pp; at $\varepsilon=0.50$, it rises to 54.6% (+41.2 pp), indicating substantial sensitivity to large differential error. Our main conclusions hold under mild violations but should be interpreted cautiously when outcome-conditional annotation is suspected.

Extended stress test. An extended stress test across 96 configurations (4 topologies $\times$ 2 $K \times$ 3 CM qualities $\times$ 4 ε , 100 MC replications each; Table 10) confirms these patterns broadly: at $\varepsilon=0$, sign-flip rates are zero for all configurations; at $\varepsilon=0.10$, the maximum sign-flip rate is 7% (IV, low-quality $C_{kk}=0.4$, $K=3$). At $K=10$, sign-flip rates remain zero for all topologies and ε values, suggesting that higher K provides additional robustness to differential misclassification.

Table 10: Non-differential misclassification stress test: sign flip rate (SF) and mean bias ratio (BR = $|\hat{\tau}/\tau|$) under increasing differential departure ε .

Topology	K	Q	$\varepsilon=0.0$		$\varepsilon=0.01$		$\varepsilon=0.05$		$\varepsilon=0.1$	
			SF	BR	SF	BR	SF	BR	SF	BR
Confounding	3	H	0.00	1.10	0.00	1.10	0.00	1.11	0.00	1.12
		M	0.00	1.14	0.00	1.14	0.00	1.14	0.00	1.14
		L	0.00	1.15	0.00	1.15	0.00	1.15	0.00	1.15
	10	H	0.00	1.06	0.00	1.06	0.00	1.08	0.00	1.09
		M	0.00	1.09	0.00	1.09	0.00	1.10	0.00	1.10
		L	0.00	1.11	0.00	1.11	0.00	1.11	0.00	1.11
Mediation	3	H	0.00	1.35	0.00	1.36	0.00	1.42	0.00	1.47
		M	0.00	1.52	0.00	1.53	0.00	1.56	0.00	1.58
		L	0.00	1.61	0.00	1.61	0.00	1.62	0.00	1.61
	10	H	0.00	1.14	0.00	1.17	0.00	1.25	0.00	1.28
		M	0.00	1.24	0.00	1.26	0.00	1.30	0.00	1.32
		L	0.00	1.30	0.00	1.32	0.00	1.33	0.00	1.34
Exposure	3	H	0.00	0.72	0.00	0.71	0.00	0.67	0.00	0.64
		M	0.00	0.42	0.00	0.41	0.00	0.38	0.00	0.35
		L	0.00	0.11	0.00	0.10	0.01	0.08	0.05	0.06
	10	H	0.00	0.87	0.00	0.86	0.00	0.84	0.00	0.80
		M	0.00	0.71	0.00	0.70	0.00	0.67	0.00	0.64
		L	0.00	0.49	0.00	0.47	0.00	0.46	0.00	0.43
IV	3	H	0.00	1.42	0.00	1.45	0.00	1.58	0.00	1.69
		M	0.00	2.56	0.00	2.52	0.00	2.87	0.00	3.12
		L	0.00	12.79	0.00	12.21	0.02	19.97	0.07	43.44
	10	H	0.00	1.30	0.00	1.44	0.00	1.71	0.00	2.03
		M	0.00	1.83	0.00	1.96	0.00	2.43	0.00	2.84
		L	0.00	3.15	0.00	3.54	0.00	4.39	0.00	9.69

SF=sign flip rate; BR=mean $|\hat{\tau}/\tau|$ (1.00=unbiased). $\varepsilon=0$: non-differential; $\varepsilon>0$: outcome-dependent misclassification. Q: H/M/L=CM diagonal 0.8/0.6/0.4 with uniform off-diagonal. Each cell: $N=10,000 \times 100$ MC replications.

Structured violation models. Beyond uniform perturbations, we compare two structured violation models across all seven topologies ($K=3$, 1,000 matrices per ε , $\min_k C_{kk} \geq 0.7$): (1) *outcome-conditional*, where $\hat{C}$ applies to $Y > \text{median}(Y)$ while C_0 applies below; (2) *label-specific*, where off-diagonal perturbation magnitude scales with $|\text{corr}(A=k, Y)|$, modeling accuracy degradation on Y -predictive categories. For exposure at $\varepsilon=0.20$, label-specific violations are substantially milder than uniform (sign-flip 0.1% vs. 10.1%), while outcome-conditional violations are more severe (sign-flip 13.9%, mean relative bias 0.50 vs. 0.40). All five sign-preserving topologies maintain 0% sign-flip under both structured models at all ε tested; IV likewise shows 0% sign-flip (its vulnerability requires pathological diagonal structure, not differential error). These results indicate that the uniform perturbation model is neither uniformly conservative nor uniformly liberal: practitioners should assess which violation pattern is most plausible for their annotation pipeline.

B.7 Robustness to Nonlinear DGP

We replace the linear structural equations with $Y \sim \text{Bernoulli}(\sigma(\mathbf{x}^\top \beta))$ and estimate via logis-

tic regression (control function for IV) for all seven topologies, using $K=3$, six real VAST confusion matrices, 500 MC replications, and $N=10,000$. Mean sign-flip rates across the six CMs:

Topology	Linear SF_τ	Logistic SF_τ	Linear SF_{any}	Logistic SF_{any}
Confounding	0.000	0.000	0.573	0.559
Mediation	0.000	0.000	0.703	0.695
Collider	0.000	0.000	0.060	0.060
Exposure	—	—	0.577	0.574
M-bias	0.000	0.000	0.095	0.099
Front-door	0.000	0.000	0.590	0.579
IV	0.086	0.074	0.086	0.560

Treatment-effect sign-flip rates (SF_τ) are zero for all five guaranteed topologies under both linear and logistic DGPs, and differ by < 0.02 for IV. The SF_{any} discrepancy for IV reflects estimator dimensionality: the control function includes a residual coefficient prone to sign noise; the core comparison is SF_τ , where both estimators agree. Linear and logistic DGPs produce virtually identical sign-flip patterns across all seven topologies ($|\Delta \text{SF}_\tau| < 0.02$). Under logistic DGP, bias is measured in log-odds units; directional conclusions (sign preservation, topology ranking) transfer directly, while magnitudes are not directly comparable to linear plim expressions. For practitioners requiring finite-sample inference under logistic DGP, bias-corrected bootstrap provides a nonparametric alternative that bypasses the linear plim framework.

B.8 Bootstrap Uncertainty of ASV

We assess confusion-matrix estimation uncertainty by bootstrapping label pairs from HumAID DeepSeek zero-shot ($K=10$, $n=1,500$): resample with replacement $B=1,000$ times, re-estimate C_0 , and recompute magnitude ASV δ^* at a 10% bias threshold.

Topology	δ^* (point)	95% CI	Valid / 1000
Confounding	—	—	0
Exposure	0.937	[0.602, 1.363]	907
IV	0.754	[0.435, 1.290]	902

Confounding never reaches the 10% threshold across all 1,000 bootstraps, confirming robustness. For IV, the 95% CI spans a $3\times$ range (0.435 to 1.290); practitioners should bootstrap C_0 rather than rely on a point estimate when ASV is near the decision boundary. CI width scales as $O(1/\sqrt{n})$ where n is the gold-label validation set size; at $n=1,500$, IV CIs are $\sim 3\times$ wider than exposure CIs, reflecting the Wald estimator’s denominator instability. A convergence study ($K=3$, $B=1,000$) confirms this scaling: exposure CI width narrows from 0.60 ($n=500$) to 0.05 ($n=50,000$), and IV

from 0.52 to 0.04, with the product width $\times \sqrt{n}$ approximately constant ($CV < 0.20$).

B.9 Cross- K Robustness

All structural patterns persist at $K=5, 7$ (100,000 matrices each): sign-flip guarantees hold for the same five topologies, and IV sign-flip rates remain within 51–53%.

B.10 ASV Diagnostic Results

For HumAID ($K=10$, 56 scenarios), 55.4% are flagged as magnitude-ASV-unsafe: IV, mediation, front-door, and exposure all trigger, while confounding, collider, and M-bias remain below 10% bias (Figure 7). For CivilComments ($K=2$), 100% of scenarios are ASV-unsafe despite F_1 ranging from 0.69 (7B) to 0.82 (frontier).

Overlaying all 23 empirical matrices onto the theoretical Dirichlet violation-rate curves (Figure 4) confirms placement: VAST in the danger zone, HumAID mid-range, CivilComments approaching safe for exposure but still in IV amplification.

B.11 ASV Stability Across DGP Parameters

We test whether ASV rankings are stable across different data-generating process (DGP) configurations by varying the treatment-effect vector β across three settings: alternating-large ($\beta = [0, 1.0, -0.5]$), same-sign ($\beta = [0, 0.5, 0.3]$), and alternating-small ($\beta = [0, 0.2, -0.1]$), with $K=3$, $\tau=0.5$, and five C_0 diagonal values $\in \{0.5, 0.6, 0.7, 0.8, 0.9\}$.

Table 11: ASV stability across DGP configurations. Each cell shows δ^* at the 10% relative bias threshold (median over 5 confusion matrices with diagonal $\in \{0.5, 0.6, 0.7, 0.8, 0.9\}$). Lower δ^* indicates greater fragility. Confounding is consistently most robust; for alternating-sign β , the ranking Exposure $<$ IV $<$ Confounding holds. “ ∞ ” = threshold never reached within δ range.

DGP Configuration	Magnitude ASV $\delta_{10\%}^*$		
	Confounding	Exposure	IV
Alt. large: $\beta = [0, 1.0, -0.5]$	0.348	0.090	0.203
Same-sign: $\beta = [0, 0.5, 0.3]$	∞	0.223	0.203
Alt. small: $\beta = [0, 0.2, -0.1]$	∞	0.090	0.203

Three patterns emerge (Table 11). First, IV magnitude ASV is identical across all three DGPs ($\delta^*=0.203$). This is a mathematical identity, not an empirical finding: the Wald estimator $\hat{\tau}_{IV} = RF/FS$ has β in both numerator and denominator, and the ratio depends only on C and the first-stage

structure, making ASV invariant to β by construction. Second, exposure ASV depends on the coefficient ratio structure: alternating-large ($\beta_2/\beta_1 = -0.5$) and alternating-small ($\beta_2/\beta_1 = -0.5$) share the same ratio and yield identical $\delta^*=0.090$, while same-sign ($\beta_2/\beta_1 = 0.6$) gives $\delta^*=0.223$; thus the experiment provides two distinct sign-structure conditions, not three. Third, confounding ASV depends on effect magnitudes: only alternating-large triggers the 10% threshold ($\delta^*=0.348$); smaller effects never reach 10% bias.

Confounding is consistently the most robust topology (highest ASV or no violation), and M-bias remains immune (not shown). However, the relative fragility of exposure vs. IV inverts with coefficient sign structure: under alternating signs, exposure is most fragile ($0.090 < 0.203$); under same-sign effects, IV is most fragile ($0.203 < 0.223$). Practitioners should compute ASV with their specific DGP rather than relying on generic rankings.

B.12 Extended Conjecture Verification

We test envelope monotonicity (Conjecture 1) by sampling 10,000 random C_0 matrices from Dirichlet(1, ..., 1) per topology for $K \in \{3, 5, 7\}$ (210,000 total), sweeping δ in steps of 0.01, and counting matrices with at least one non-monotonic step exceeding tolerance 10^{-12} . Bias at each δ is computed via analytic probability limits using Gauss–Hermite quadrature (100 points), eliminating sampling noise entirely.

The δ step size is 0.01. “Worst mag.” is the largest single-step deviation $|B(\delta_{i+1})| - |B(\delta_i)|$ in the non-monotonic direction. *100,000 matrices at step = 0.005 yield exposure $\varepsilon = 0.075$.

Table 12: Envelope monotonicity violations per 10,000 random C_0 matrices.

Topology	$K=3$	$K=5$	$K=7$	Worst mag.	Direction
Confounding	8467	8063	7892	0.011	non-dec
Mediation	8236	7905	7893	0.013	non-dec
Collider	8683	7669	7570	0.005	non-inc
Exposure	485	205	77	0.056*	non-dec
M-bias	9791	7783	7399	0.002	non-inc
IV	7433	6015	5711	$> 10^6$	non-dec
Front-door	8429	8079	7899	0.013	non-dec

Violation counts are high (60–97% of matrices), but magnitudes are small for non-IV topologies under the default parameterization (Table 12); DGP sensitivity analysis (Table 14) shows ε scales with effect magnitude $|\beta|$, remaining below 0.02 for $|\beta| \leq 1$ across non-exposure topologies. These re-

Bootstrap Distribution of ASV δ^* (B=1000)




Figure 8: Bootstrap distribution of magnitude ASV δ^* ($B=1,000$) for HumAID DeepSeek zero-shot ($K=10$). Dashed line: point estimate; shaded: 95% CI.

sults are computed from analytic probability limits with zero sampling noise, confirming the violations are structural properties of multi-class confusion matrices rather than numerical artifacts. IV deviations reflect Wald estimator singularities when the first-stage coefficient passes near zero; this is a property of the Wald estimator, not of the ASV parameterization. Adversarial search (500 restarts per configuration, $K \in \{3, 5\}$) finds genuine counterexamples for all topologies except IV $K=3$, confirming that strict monotonicity does not hold universally. The bias envelope is ε -approximately monotonic with $\varepsilon < 0.03$ for collider and M-bias across all tested τ ; confounding, mediation, and front-door achieve $\varepsilon < 0.06$ at default DGP parameters (see DGP sensitivity below); exposure satisfies $\varepsilon_{p99} < 0.008$ but extreme-tail matrices ($< 0.02\%$ of draws) reach $\varepsilon \approx 0.15$. Empirically, ε scales as $O(1/K^\alpha)$ with $\alpha \in [0.6, 1.5]$ (topology-dependent): for confounding, mediation, and front-door, $\alpha \approx 1.4$; for exposure, $\alpha \approx 1.0$; for collider and M-bias, $\alpha \approx 0.7$. A universal empirical upper bound is $\varepsilon \leq 0.137/\sqrt{K}$, verified across all tested configurations ($K \in \{3, 5, 7\}$). This scaling supports the practical validity of ASV as a sensitivity tool at any meaningful bias threshold.

DGP parameter sensitivity. We repeat the verification across $\tau \in \{0.1, 0.5, 1.0, 2.0, 5.0\}$, varying the direct $T \rightarrow Y$ coefficient for confounding ($\beta_T = \tau$), mediation, collider, and M-bias (DGP τ), and rescaling the $A \rightarrow Y$ coefficient vector to $\beta \cdot \tau$ for exposure and front-door. With 10,000 random C_0 per (K, τ) setting (Table 13), ε remains small (< 0.03) for $\tau \geq 0.5$ for confounding, mediation, collider, and M-bias, but grows at $\tau=0.1$: mediation produces 1,950 / 30,000 violations across

$K \in \{3, 5, 7\}$ (worst $\varepsilon=0.130$ at $K=3$) and confounding 200 / 30,000 (worst $\varepsilon=0.118$ at $K=3$). The mechanism is dimensional: additive terms not scaling with τ (e.g. the β -contribution to $\mathbb{E}[YD_k]$ in mediation and confounding) contribute a τ -independent offset to the bias numerator, so the ratio $|B|/|\tau|$ inflates as $\tau \rightarrow 0$. A single mediation violation at $\tau=2.0$ ($\varepsilon=0.112$, 1/30,000 matrices) arises from a specific confusion matrix structure where off-diagonal entries induce partial cancellation along the mediation pathway; this isolated case does not affect the p_{99} bound ($\varepsilon_{p99}=0.005$). Collider and M-bias remain robust at all τ ($\varepsilon < 0.022$ and $\varepsilon < 0.015$ respectively). Exposure shows scattered tail outliers but $\bar{\varepsilon} \approx 3 \times 10^{-4}$ at every τ , consistent with linear scaling: when β rescales by τ , both bias and τ_{true} scale by the same factor and ε is theoretically invariant. Front-door inverts the pattern (ε grows with τ , reaching 0.139 at $\tau=5$, $K=3$, with 664 violations) because $\tau_{\text{true}} = \text{plim}(\hat{\tau}_T \mid C=I)$ is pinned at 0.5 by the U -confounding path while bias magnitudes scale with β ; this is an artefact of how τ_{true} is normalised, not a genuine failure of envelope monotonicity. In practice, ASV's δ^* thresholds (typically ≥ 0.1 in our deployment) are well above the region where these ε -violations occur, so the small- τ degradation does not affect ASV's operational reliability; we flag it as a caveat for applications targeting $|\tau| \lesssim 0.1$ in standardised units.

$\varepsilon \equiv \max_{\delta} (|B(\delta+\Delta)| - |B(\delta)|) / |\tau_{\text{true}}|$ in the non-monotonic direction, taking the max over $K \in \{3, 5, 7\}$ with 10,000 random C_0 per cell. Bold marks $\varepsilon > 0.02$ (non-exposure) or $\varepsilon > 0.08$ (exposure). Counts in parentheses are violations / 30,000. Tables 12 and 13 use independent random

seeds; for topologies with low violation rates (exposure $<0.02\%$), the worst-case ε is sample-sensitive.

Table 13: DGP parameter sensitivity of envelope monotonicity violations.

Topology	$\tau=0.1$	$\tau=0.5$	$\tau=1.0$	$\tau=2.0$	$\tau=5.0$
Confounding	0.118 (200)	0.027 (0)	0.014 (0)	0.007 (0)	0.002 (0)
Mediation	0.130 (1950)	0.026 (0)	0.013 (0)	0.112 (1)	0.003 (0)
Collider	0.022 (0)	0.006 (0)	0.005 (0)	0.004 (0)	0.002 (0)
Exposure	0.111 (5)	0.122 (2)	0.152 (2)	0.089 (3)	0.148 (3)
M-bias	0.014 (0)	0.003 (0)	0.001 (0)	0.001 (0)	0.000 (0)
Front-door	0.003 (0)	0.013 (0)	0.027 (0)	0.051 (0)	0.139 (664)

Effect-magnitude sensitivity. We further vary the direct effect $\beta \in \{0.1, 0.5, 1.0, 2.0, 5.0\}$ and noise variance $\sigma^2 \in \{0.5, 1.0, 2.0\}$ across all seven topologies ($K=3$, 10,000 Dirichlet(1, ..., 1) matrices per cell). Table 14 reports the worst-case relative envelope deviation ε_{rel} . For non-exposure topologies, ε scales approximately linearly with $|\beta|$; the safe/fragile partition is stable across all 105 cells.

β	σ^2	CONFO	MEDIA	COLLI	EXPOS	MBIAS	IV	FRONT
0.1	0.5	0.0002	0.0002	$5.7e-5$	0.1146	$7.6e-7$	0	0.0024
0.1	1	$9.3e-5$	0.0002	$6.0e-5$	0.1027	$4.6e-7$	0	0.0030
0.1	2	$4.4e-5$	0.0002	$5.8e-5$	0.0724	$1.9e-7$	0	0.0044
0.5	0.5	0.0040	0.0057	0.0014	0.0481	0.0003	0	0.0104
0.5	1	0.0019	0.0060	0.0014	0.0953	0.0002	0	0.0151
0.5	2	0.0010	0.0057	0.0014	0.0991	$9.8e-5$	0	0.0214
1	0.5	0.0083	0.0272	0.0048	0.1886	0.0021	0	0.0202
1	1	0.0066	0.0262	0.0048	0.1100	0.0015	0	0.0304
1	2	0.0039	0.0257	0.0048	0.2153	0.0009	0	0.0388
2	0.5	0.0148	0.1326	0.0137	0.0551	0.0074	0	0.0315
2	1	0.0091	0.0938	0.0117	0.0993	0.0063	0	0.0435
2	2	0.0125	0.0989	0.0109	0.0998	0.0045	0	0.0616
5	0.5	0.0083	0.8983	0.0305	0.1728	0.0155	0	0.0459
5	1	0.0230	0.8870	0.0267	0.0575	0.0156	0	0.0653
5	2	0.0122	0.9082	0.0206	0.0870	0.0132	0	0.0945

Table 14: Conjecture 1 DGP sensitivity: worst-case $\varepsilon_{\text{rel}} = \max_{C_0} (\text{wrong-direction step}) / |\tau_{\text{true}}|$ across 5 effect magnitudes β , 3 noise variances σ^2 , and 7 topologies ($K=3$, 10,000 Dirichlet(1, ..., 1) matrices per cell). ε scales approximately linearly with $|\beta|$ for non-exposure topologies; the safe/fragile partition is stable across all cells.

B.13 Proportional Effects Breakdown

Theorem 1 requires proportional effects ($f = \lambda g$); here we quantify how $\lambda(C)$ out-of-bounds (OOB) rates depend on the degree of departure from proportionality. We simulate $K=3$ mediation with 100,000 random confusion matrices (diagonal ≥ 0.5 , Dirichlet off-diagonal) at five proportionality degrees: exact ($f_j - f_0 = \lambda(g_j - g_0)$), near-proportional ($f = \lambda g + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$, $\sigma \in \{0.1, 0.3, 0.5\}$), and fully independent ($f \perp g$). Table 15 reports the $\lambda(C) \notin [0, 1]$ rate.

Results use $K=3$ with 100K random confusion matrices per row.

Table 15: $\lambda(C)$ out-of-bounds rate and mediation sign failure rate.

Proportionality degree	$\lambda(C) \notin [0, 1]$	Sign failure
Exact ($\sigma=0$)	0.0%	0.85%
Near ($\sigma=0.1$)	10.3%	0.87%
Near ($\sigma=0.3$)	19.0%	0.96%
Near ($\sigma=0.5$)	22.4%	0.93%
Independent ($f \perp g$)	22.9%	0.17%
NLP-realistic DGP	0.07%	0.0%

The λ OOB rate rises steeply from 0% to $\sim 10\%$ at $\sigma=0.1$ and saturates near 23% for independent (f, g). Sign failure rates remain below 1% across all conditions because λ excursions outside $[0, 1]$ are typically small and do not reverse the dominant effect direction. The NLP-realistic DGP ($\delta=[0, 0.8, -0.6]$, $\beta=[0, 1.0, -0.5]$) exhibits near-zero OOB (0.07%) and zero sign failures, confirming that proportional effects hold approximately in practice.

B.14 FWL Invariance Under Nonlinear Models

FWL invariance guarantees exactly 0% bias reduction from matrix-based corrections (MLA, naive plug-in) under linear OLS for covariate topologies (confounding, mediation, collider, M-bias). Here we test whether this invariance persists under logistic regression across sample sizes $N \in \{1000, 5000, 10,000, 50,000\}$ and three error levels ($C_{kk} \in \{0.9, 0.7, 0.5\}$), with 200 Monte Carlo replicates per cell. Table 16 reports the mean logistic bias reduction (positive = correction helps).

Results are averaged over $N \in \{1K, 5K, 10K, 50K\}$ with 200 MC reps each. Under linear OLS, all covariate topologies yield exactly 0%. Exposure is shown for comparison.

Table 16: Mean logistic bias reduction (%) from matrix-based correction.

Topology	Low ($C_{kk}=0.9$)	Med ($C_{kk}=0.7$)	High ($C_{kk}=0.5$)
Confounding	-0.003%	$<0.001\%$	$<0.001\%$
Mediation	-0.004%	$<0.001\%$	$<0.001\%$
Collider	-0.033%	-0.020%	$<0.001\%$
M-bias	-0.049%	-0.010%	$<0.001\%$
Front-door	0.001%	-0.001%	$<0.001\%$
Exposure	-68.5%	-48.7%	-22.5%

For all five covariate topologies, logistic bias reduction is negligible ($|\text{mean}| < 0.05\%$) across all

N and error levels, with no systematic N -scaling: medians cluster within $\pm 0.002\%$ at $N=50,000$. The large standard deviations observed at certain (N, error) cells for collider and M-bias (up to 2.6% std at $N=10,000$) reflect rare outlier replicates; medians remain near zero. This confirms that FWL invariance is a robust structural property that extends beyond linear models for covariate topologies.

In contrast, exposure shows substantial FWL breakdown under logistic regression: -22% to -78% bias reduction (negative = correction worsens bias), stabilising at $N \geq 5,000$. The breakdown magnitude increases with accuracy (C_{kk}), consistent with the exposure topology’s sensitivity to the nonlinear link function transforming category weights.

B.15 PPI++ Gold Label Ratio Sensitivity

We test PPI++ correction performance as a function of gold label fraction ($K=3$, $N=10,000$, 100 MC replications, 5 Dirichlet confusion matrices with $\min_k C_{kk} \geq 0.6$).

Negative values indicate PPI++ increases RMSE relative to the naive estimator.

Table 17: PPI++ RMSE reduction (%) vs. gold label fraction.

Topology	5%	10%	20%	50%	100%
Confounding	+84.8	+88.6	+91.8	+94.2	+95.2
Exposure	+75.2	+83.1	+87.1	+91.2	+93.8
Mediation	+56.7	+70.4	+77.2	+85.3	+88.7
IV	+1.4	+22.1	+50.1	+66.7	+76.4
M-bias	-125.3	-43.6	-19.6	+6.0	+18.5

For confounding and exposure, PPI++ achieves $>75\%$ RMSE reduction even at 5% gold labels; IV requires $\geq 20\%$ gold to reach 50% reduction, reflecting the Wald estimator’s higher variance. M-bias exhibits harmful correction at $\leq 20\%$ gold (-125% at 5%), because PPI++ re-exposes the M-bias pathway that annotation error had attenuated; only at $\geq 50\%$ gold does the estimator become weakly beneficial ($+6\%$).

B.16 Realistic Differential Misclassification Thresholds

Our non-differential assumption (Definition 1) requires $P(A^*|A, Y) = P(A^*|A)$. We estimate how much Y -stratified differential error each of our 23 empirical LLM confusion matrices can tolerate before exposure sign reversal exceeds 5%. For each matrix C , we construct $C_{\text{high}} = \text{renorm}(C + \varepsilon\Delta)$

and $C_{\text{low}} = \text{renorm}(C - \varepsilon\Delta)$ with random off-diagonal perturbation Δ (500 MC trials per ε), sweep $\varepsilon \in [0, 0.50]$, and report the safe threshold $\varepsilon_{\text{safe}}^* = \max\{\varepsilon : \text{reversal rate} < 5\%\}$ (Table 18).

For $K=2$ (binary), no configuration produces sign reversal up to $\varepsilon=0.50$. For $K \geq 3$, the majority of configurations are fragile ($\varepsilon_{\text{safe}}^*=0$), confirming that the non-differential assumption is load-bearing for multi-class exposure.

Table 18: Safe differential misclassification thresholds $\varepsilon_{\text{safe}}^*$ by dataset/ K .

Dataset	K	n	Mean ε^*	Median ε^*	$\varepsilon^*=0$
Civil Comments	2	7	0.50	0.50	0/7
VAST	3	8	0.003	0.00	7/8
HuMAID	10	8	0.045	0.00	7/8

The distribution is sharply bimodal: all $K=2$ configurations are immune to differential misclassification within the tested range ($\varepsilon_{\text{safe}}^*=0.50$), consistent with the binary sign-preservation guarantee (Proposition 1(c)). For $K \geq 3$, 14 of 16 configurations have $\varepsilon_{\text{safe}}^*=0$: sign reversal occurs at the smallest tested $\varepsilon=0.01$. Two outliers (VAST: $\varepsilon^*=0.02$; HuMAID: $\varepsilon^*=0.36$) have atypically structured confusion matrices. This confirms that even small violations of non-differential misclassification are consequential for multi-class exposure, underscoring the practical importance of validating $P(A^*|A, Y) \approx P(A^*|A)$ on gold-labeled subsets when $K \geq 3$.

B.17 Proportional Effects with Empirical Confusion Matrices

Theorem 1 assumes proportional effects ($f = \lambda g$). We evaluate this assumption using all 23 empirical LLM confusion matrices under four DGP groups in the mediation topology (23 matrices $\times$ 50 seeds $\times$ 100,000 precomputed samples per seed).

Evaluation uses the mediation DGP. $\lambda(C) \in [0, 1]$: proportion satisfying the Theorem 1 bound. Sign pres.: sign preservation rate of the direct effect estimate.

Table 19: Proportional effects validation with 23 real LLM confusion matrices.

DGP Group	$\lambda(C)$ mean	$\lambda(C) \in [0, 1]$	Sign pres.	CV
Proportional ($f=\lambda g$)	0.271	99.8%	100%	0.00
Near-prop. ($\sigma=0.1$)	0.285	99.2%	100%	1.06
Non-proportional	0.240	85.0%	97.3%	3.60
Sparse	0.296	75.8%	98.6%	6.59

Under proportional and near-proportional ef-

fects, $\lambda(C)$ remains within $[0, 1]$ for $\geq 99\%$ of configurations with 100% sign preservation: the Theorem 1 conditions are satisfied by real LLM confusion matrices. Even under non-proportional effects (f independent of g), sign preservation remains at 97.3%, with $\lambda(C)$ exceeding $[0, 1]$ in 15% of cases. The per- K breakdown shows that $\lambda(C)$ is dimension-dependent: $K=2$ (mean 0.249), $K=3$ (mean 0.082), $K=10$ (mean 0.490), with 100% in $[0, 1]$ across all K under proportional effects. Sparse DGP (zero-effect categories) produces the highest violation rate (24.2%), consistent with the breakdown mechanism: zero-effect categories make f/g ratios undefined, maximally departing from proportionality.

B.18 Case Study: ASV Diagnosis of Political Ad Tone Classification

We illustrate the ASV diagnostic workflow on a concrete CSS application. Egami et al. (2023) use LLM classifiers to annotate political advertising tone ($K=3$: promote/contrast/attack) from the Fowler & Ridout database, then estimate the causal effect of ad tone on voter outcomes. The annotated variable is the **exposure** (treatment). Using a representative $K=3$ confusion matrix from our empirical set (VAST DeepSeek zero-shot, mean diagonal 0.734) as a proxy:

Coefficient	δ_{mag}^* (10%)	δ_{sf}^*	Relative bias at $\delta=1$
β_1 (primary)	0.234	> 1.5	39.7%
β_2 (secondary)	0.154	> 1.5	60.4%

Both coefficients are magnitude-unsafe ($\delta^* < 1$), with the secondary coefficient more fragile ($\delta^*=0.154$). If the same annotator were instead classifying a **confounder** (e.g., media framing as a control variable), magnitude ASV rises to $\delta^*=0.690$, a $3\times$ increase in error tolerance from the same confusion matrix. This demonstrates the central insight: topology, not annotator quality, determines vulnerability. Algorithm 1 routes this exposure topology to PPI++ or DSL correction; matrix-based methods (naive plug-in, MLA) are ineffective due to FWL non-applicability.

B.19 End-to-End Demonstration

Topology	Naive Bias (%)	ASV δ_{mag}^*	Routed Method	Corrected Bias (%)	Improvement (%)
Confounding (naive plug-in)	$+22.71 \pm 2.29$	0.352	no correction (FWL)	$+22.71 \pm 2.29$	+0.0%
(MLA)	–	–	–	$+22.71 \pm 2.29$	+0.0%
Exposure	-42.87 ± 2.42	0.239	PPI++	$+22.71 \pm 2.29$	+0.0%
IV	$+82.40 \pm 12.22$	0.204	MC-SIMEX (gold-free)	-0.22 ± 4.94	+90.9%
				-197.00 ± 3.79	-139.1%

Table 20: End-to-end demo over 100 replicates ($N=10,000$, $K=3$, 20% gold). Confounding: plug-in corrections leave the treatment coefficient unchanged (FWL invariance), so Algorithm 1 routes to “no correction”. Exposure: PPI++ recovers the target coefficient. IV: Wald estimator instability produces large bias; PPI++ (OLS-based) is limited, routing to MC-SIMEX.

The demonstration uses a symmetric C_0 ($d_{\text{rel}}=0$); empirical $K=3$ matrices deviate substantially (median $d_{\text{rel}}=0.33$, Remark 5), but core ASV ranking and correction routing remain valid under Conjecture 1.

B.20 Semi-Real Case Study with Production LLM Confusion Matrices

We repeat the end-to-end pipeline using 10 confusion matrices estimated from production LLM annotation of VAST stance detection ($K=3$; DeepSeek, Llama4, Qwen2.5-7B, Qwen3 under zero-shot and few-shot prompting). These matrices exhibit median $d_{\text{rel}}=0.33$ (Remark 5), far from the symmetric assumption of Theorem 2. Results confirm both topology-dependent patterns: PPI++ achieves 93.6% median exposure bias reduction while plug-in corrections yield 0% improvement for confounding (FWL invariance, machine precision; Table 21).

LLM Config	d_{rel}	Exposure		Confounding	
		Naive (%)	PPI++ (%)	Naive (%)	Plugin (%)
DeepSeek FS-3	0.30	-61.3	+0.1	+28.5	+28.5
DeepSeek FS-5	0.30	-61.5	-0.5	+28.7	+28.7
DeepSeek ZS	0.31	-61.1	-0.5	+28.2	+28.2
Llama4 FS-3	0.37	-64.3	+0.0	+28.4	+28.4
Llama4 FS-5	0.32	-61.3	+0.0	+26.8	+26.8
Llama4 ZS	0.36	-59.3	-0.2	+27.9	+27.9
Qwen2.5-7B FS-3	0.38	-68.9	-0.5	+29.3	+29.3
Qwen2.5-7B FS-5	0.34	-68.7	+0.3	+29.0	+29.0
Qwen2.5-7B ZS	0.38	-65.5	+0.3	+28.7	+28.7
Qwen3 ZS	0.32	-60.3	-0.5	+28.2	+28.2
Median improvement		+93.6%		-0.0%	

Table 21: Semi-real case study with VAST $K=3$ confusion matrices from 10 LLM configurations. Bias is relative to true τ . Exposure: PPI++ substantially reduces bias. Confounding: plug-in correction is ineffective (FWL invariance).

B.21 Correction Method Sensitivity to K

We test whether the correction method ranking from Table 6 ($K=3$) is stable at $K=5$ and $K=10$. For each K , we sample 5 random Dirichlet

confusion matrices (diagonal ≥ 0.6), run 100 MC replications per topology–matrix combination ($N=5,000$, 10% gold fraction), and compute RMSE reduction for all five methods across all seven topologies.

Grand ranking stability. Spearman rank correlations across all topology–method pairs: $\rho_{3,5} = 0.80$, $\rho_{3,10} = 0.84$, $\rho_{5,10} = 0.93$ (all $p < 10^{-8}$). Per-topology correlations are even stronger: 5 of 7 topologies (confounding, collider, M-bias, IV, and exposure at $K=10$) achieve $\rho = 1.0$ between $K=3$ and $K=10$.

Method-specific patterns. PPI++ effectiveness degrades with K : confounding RMSE reduction drops from 85.6% ($K=3$) to 75.6% ($K=10$); mediation drops from 57.1% to 9.0%. At $K=10$, PPI++ becomes harmful for IV (-124%) and exposure (-11%). MC-SIMEX is the most K -robust method, maintaining the top rank for 5/7 topologies at $K=10$ (vs. 3/7 at $K=3$). FWL invariance (naive plug-in and MLA achieving 0% bias reduction on the treatment coefficient) holds at all K tested, confirming the structural nature of this result.

Bold indicates the top method for that topology– K combination.

Table 22: Top correction method and RMSE reduction (%) by topology across K .

Topology	$K=3$	$K=5$	$K=10$
Confounding	PPI++ 85.6	PPI++ 76.4	PPI++ 75.6
Mediation	PPI++ 57.1	SIMEX 59.7	SIMEX 59.0
Collider	SIMEX 55.4	SIMEX 63.4	SIMEX 57.3
Exposure	SIMEX 72.5	SIMEX 65.0	SIMEX 47.3
M-bias	SIMEX 3.2	SIMEX 47.0	SIMEX 43.7
IV	SIMEX 30.9	SIMEX 30.4	SIMEX 27.1
Front-door	PPI++ 82.9	PPI++ 69.0	SIMEX 70.1

B.22 Comparison with Related Sensitivity Frameworks

Table 23: Comparison of sensitivity frameworks.

	ASV	E-value	Partial R^2
Parameter space	$K \times K$ confusion $C(\delta)$	Scalar confounder	R^2 of omitted var
Sensitivity path	$(1-\delta)I + \delta C_0$	Monotone scalar	Monotone scalar
Topology-specific	✓	✗	✗
Sign-flip diagnostic	✓ (δ_{sign}^*)	✗	✗
Topology dependence	✓ (7 topologies ranked)	✗	✗
Monotonicity	Proved $K=2$; conj. $K \geq 3$	Guaranteed	Guaranteed
Target	Annotation error	Unmeasured confounding	Omitted variable

B.23 Per-Topology F_1 –Bias Correlation

The cross-topology correlation between macro F_1 and $|\text{bias}|$ is weak ($\rho = -0.10$, $p = 0.45$; §4.4), but decomposing by topology reveals a structured pattern. Table 24 reports Spearman correlations

computed over 23 LLM configurations (HumAID $K=10$, VAST $K=3$, CivilComments $K=2$; 161 total scenarios).

Correlations are computed over 23 empirical LLM configurations $\times$ 3 datasets. Negative ρ : higher $F_1 \rightarrow$ lower bias (expected). Positive ρ : higher $F_1 \rightarrow$ higher bias (protective attenuation removed).

Table 24: Within-topology Spearman $\rho(F_1, |\text{bias}|)$ across LLM configurations.

Topology	n	ρ	p
Confounding	23	-0.557	0.006
Mediation	23	-0.565	0.005
Exposure	23	-0.531	0.009
Front-door	23	-0.568	0.005
IV	23	-0.700	<0.001
Collider	23	$+0.310$	0.150
M-bias	23	$+0.867$	<0.001

Five topologies show the expected negative correlation: better annotation quality ($F_1 \uparrow$) reduces causal bias. IV shows the strongest effect ($\rho = -0.70$), consistent with the Wald estimator’s sensitivity to misclassification in the first stage. M-bias uniquely reverses: $\rho = +0.87$ ($p < 0.001$). The mechanism is that misclassification attenuates the spurious M-bias path, so the naive estimator under noisy annotations is closer to the true causal effect; improving annotation quality (higher F_1) removes this protective attenuation and increases bias. This is consistent with the correction results (§4.5), where PPI++ harms M-bias (-88%) by removing the same protective noise. Collider shows a positive but non-significant trend ($\rho = +0.31$, $p = 0.15$), consistent with Berkson bias attenuation. Per-dataset breakdowns (CivilComments $K=2$: $\rho = -1.0$ for all topologies with $n=7$) should be interpreted with caution given the small sample and limited F_1 variation within the binary task.

B.24 Confusion Matrix Estimation Robustness

The ASV framework requires the confusion matrix C , which in practice must be estimated from a held-out gold-labeled subset. We assess how estimation noise in $\hat{C}$ propagates to the safe/fragile partition and ASV ranking. For each of 8 empirical confusion matrices (VAST, 4 LLMs $\times$ 2 prompting strategies), we draw $n_{\text{gold}} \in \{50, 100, 200, 500\}$ samples from a Multinomial($n_{\text{gold}}, C_{\cdot k}$) distribution per column, compute $\hat{C}$, and evaluate the re-

sulting ASV ranking across 200 bootstrap replicates.

Spearman ρ is computed vs. true- C ranking under multinomial bootstrap from finite gold sets (8 empirical matrices $\times$ 200 bootstrap replicates per n_{gold}).

Table 25: $\hat{C}$ estimation robustness: partition stability and ASV ranking correlation.

n_{gold}	Partition stability	ρ (median)	ρ (q_{25})
50	50.4%	0.991	0.789
100	82.0%	1.000	0.991
200	85.6%	1.000	0.991
500	88.5%	1.000	0.991

The ASV ranking is highly robust: median Spearman $\rho \geq 0.99$ even at $n_{\text{gold}}=50$. The safe/fragile partition (5 safe vs. 2 fragile topologies) recovers reliably from $n_{\text{gold}} \geq 100$ (82% stability), increasing to 89% at $n_{\text{gold}}=500$. The remaining instability concentrates in source matrices with low diagonal dominance (e.g., VAST DeepSeek few-shot: $\min C_{kk} = 0.16$), where bootstrap noise can shift borderline ASV values across the partition boundary. All 8 truth partitions match the canonical safe/fragile partition, confirming that the theoretical dichotomy holds empirically.

B.25 Sign Reversal Determinants

We investigate which confusion matrix features predict exposure sign reversal across the 23 empirical LLM configurations. For each matrix, we compute 10 structural features: $\min_k C_{kk}$ (minimum diagonal), $\overline{C_{kk}}$ (mean diagonal), $\max_{j \neq k} C_{jk}$ (maximum off-diagonal), $\max_{j \neq k} C_{jk}/C_{kk}$ (off-diagonal ratio), $\log(\kappa(C))$ (log condition number), mean and max column entropy, spectral gap, $\det(C)$, and minimum singular value. Sign reversal is defined as any exposure coefficient reversing sign under misclassification.

Best single-feature rule: $\det(C) < 0.014 \rightarrow$ sign reversal (accuracy 91.3%, sensitivity 85.7%, specificity 100%). An AND rule ($\min_k C_{kk} < 0.55$ and mean entropy > 0.40) achieves 95.7% accuracy with 100% sensitivity.

Table 26: Sign reversal rates by K and best single-feature decision rule.

K	n	Sign reversal rate
2	7	0/7 (0%)
3	8	7/8 (87.5%)
10	8	7/8 (87.5%)

Binary matrices ($K=2$) never produce exposure sign reversal, consistent with classical attenuation theory. For $K \geq 3$, the determinant of C is the strongest single predictor: $\det(C) < 0.014$ identifies sign-reversing matrices with 100% specificity (zero false positives) and 85.7% sensitivity. An AND rule combining $\min_k C_{kk} < 0.55$ with mean column entropy > 0.40 achieves 100% sensitivity at 88.9% specificity. These rules are robust to β -scaling (accuracy unchanged at $0.3\times$ and $3\times$ effect magnitude) but sensitive to β -sign structure: under uniform-positive β (all coefficients same sign), no matrices produce sign reversal regardless of C , as cancellation across differently-signed coefficients is the underlying mechanism.

B.26 Y-Stratified Non-Differential Verification

We directly test the non-differential assumption (Definition 1) by stratifying confusion matrices by outcome-proxy variables and testing $C(Y=\text{high}) \approx C(Y=\text{low})$. Using 12 LLM configurations on VAST ($K=3$; 4 models $\times$ 3 prompting strategies, 1,500 samples each), we compute per-stratum confusion matrices under three Y-proxies and apply permutation tests ($B=10,000$, test statistic: $\max_{j,k} |C_{jk}^{(0)} - C_{jk}^{(1)}|$).

Content-independent Y (text length) shows no significant violation (1/12 rejection $\approx$ expected false positive rate at $\alpha=0.05$). Content-dependent proxies (topic, sentiment) show expected violations since LLMs observe text content.

Table 27: Y-stratified non-differential test.

Y proxy	Type	Reject ($\alpha=0.05$)	Median p
Text length	Content-indep.	1/12	0.490
Topic	Content-dep.	11/12	0.002
Sentiment	Content-dep.	10/12	0.017

The single text-length rejection (Llama-4 few-shot-3, $p=0.027$) is consistent with the expected false positive rate ($0.05 \times 12 = 0.6$). Content-dependent proxies (topic, sentiment) show systematic rejection because LLMs process text content from which these variables can be inferred, violating $Y \perp \text{text} \mid A$. In causal inference applications,

2382 Y is typically an external outcome (e.g., health
2383 status, economic indicator) that cannot be inferred
2384 from the annotated text, satisfying the indepen-
2385 dence condition. The content-dependent rejections
2386 confirm statistical power: the test can detect viola-
2387 tions when they exist.

2388 **C Notation**

Table 28: Key notation used throughout the paper.

Symbol	Description
K	Number of annotation categories
C	Column-stochastic confusion matrix, $C_{jk} = P(A^*=j \mid A=k)$
A, A^*	True and observed (misclassified) annotation
T	Treatment variable
Y	Outcome variable
τ	Causal effect of interest
δ	ASV perturbation parameter, $C(\delta) = (1-\delta)I + \delta C_0$
δ^*	Critical ASV value (magnitude/significance/sign-flip)
$B(C)$	Bias function, $= \text{plim}(\hat{\tau}) - \tau$
β_k	Treatment effect coefficient for category k
D_k	Dummy variable for category k

2389 **D Use of AI Assistance**

2390 We used Claude for coding, code drafting, debug-
2391 ging, and language polishing.